

TOWARDS NEURAL ENCODING OF
HIGHER-ORDER COGNITION WITH LARGE
LANGUAGE MODELS

SEAN MATA

ADVISOR: PROFESSOR PETER RAMADGE, PROFESSOR URI HASSON

SUBMITTED IN PARTIAL FULFILLMENT
OF THE REQUIREMENTS FOR THE DEGREE OF
BACHELOR OF SCIENCE IN ENGINEERING
DEPARTMENT OF ELECTRICAL AND COMPUTER ENGINEERING
PRINCETON UNIVERSITY

MAY 2025

I hereby declare that I am the sole author of this thesis.

I authorize Princeton University to lend this thesis to other institutions or individuals for the purpose of scholarly research.

Sean Mata

I further authorize Princeton University to reproduce this thesis by photocopying or by other means, in total or in part, at the request of other institutions or individuals for the purpose of scholarly research.

Sean Mata

Abstract

Large Language Models (LLMs) have demonstrated remarkable language understanding capabilities, significantly advancing neural encoding models. While recent progress has improved LLMs’ reasoning abilities, they still face limitations in social contexts that require nuanced understanding of human cognition. This thesis investigates whether inference-time techniques like chain-of-thought reasoning—which externalize an LLM’s reasoning process—can be leveraged for neural encoding of human reasoning and higher-order cognitive processing. Specifically, we focus on theory of mind and social reasoning within conversational contexts, examining the detection of subtext (unspoken beliefs, intentions, and desires) in communication.

Towards this we present two complementary projects: First, we develop a novel Theory of Mind-inspired Tree of Thought approach for dialogue generation that explicitly models subtextual reasoning. Our method produces diverse, contextually appropriate outputs, though automatic selection mechanisms currently struggle to outperform single-step generation. Through perturbation analysis and evaluator consistency tests, we gain insights into model confidence and the distribution of plausible responses. Second, we apply LLM-derived embeddings to neural encoding tasks, successfully modeling brain activity associated with both word- and sentence-level language processing using banded-ridge regression. Our findings suggest the potential for modifying sentence-level inputs to target the encoding of higher-order thinking in high-frequency Electrocorticography (ECoG) data. Together, these projects highlight both the promise and current limitations of LLMs in modeling social reasoning.

We conclude with proposals for future work, including improving training data, context structuring, evaluation metrics, and thought representation. Our findings point to new directions in aligning language models with human-like communicative competence and mental state inference.

Acknowledgements

An uncountable number of positive influences have shaped my journey to this point. For me, this thesis marks not just the culmination of a project, but also the nearing end of my undergraduate years—years that have defined much of my young adult life.

Within the context of this project, I am deeply grateful to Itamar Jalon for his mentorship and unwavering support. Without his guidance, I would have been lost in how to even begin tackling this work. Our regular meetings were often a highlight of my week, filled with discussions about neuroscience, as well as general life and the future. Thank you to Peter Ramadge for your invaluable support and for your guidance in managing both the research and the writing process. I am also especially grateful towards Professor Uri Hasson for welcoming me into his lab and facilitating the research opportunities that made this thesis possible. To all the members of the Hasson Lab, thank you for making me feel welcome, answering my many questions, and creating an environment in which I could learn and grow.

Beyond this project, there are many more people I wish to thank. At Princeton University, I am deeply grateful to everyone who helped me build a sense of home within this “orange bubble.” I’m proud of the community we created—rich with stories, memories, and shared experiences. From travels across states and countries—New York, Chicago, Puerto Rico, Mexico, and more—to countless campus events spanning sports, music, and community and awareness-building, to the joy of watching and performing in dance shows, each moment added depth and color to my time here. I look forward to the friendships that will continue to grow long after graduation, and I hope to cross paths with many of these friends again in the years to come.

To my family: thank you for your constant love and support. I miss you every time I find myself living in a different city or country. To my parents, no matter where

I go, I will always be your baby. To my siblings, I will always be there to support you.

Finally, to my friends from Oxford, secondary school, and throughout the years: thank you for the countless moments of growth, learning, and laughter we shared. We navigated so many phases of life together, and I look forward to continuing to grow and experience new chapters with you and all my dear friends.

Contents

Abstract	iii
Acknowledgements	iv
List of Figures	x
List of Symbols	xv
1 Introduction and Background	1
1.1 Reasoning versus language abilities in LLMs	2
1.2 Inference time generation methods	2
1.3 Theory of mind in LLMs	3
1.4 Use of Language Models to study processing in the brain	5
1.4.1 Sentence embedding encoding	7
1.5 Report Overview	7
2 Methodological Background	8
2.1 Natural Language Processing using Deep Neural Networks	8
2.2 Text similarity metrics	12
2.2.1 N-gram matching	14
2.2.2 Edit distance	15
2.2.3 Trained metrics	15
2.2.4 Perplexity	15
2.2.5 Embedding based metrics	17

3 Methodology	20
3.1 Problem setup	20
3.2 Generation strategies	21
3.2.1 Single step utterance generation	21
3.2.2 Tree of thought	23
3.3 Ranking and State Evaluation Metrics	28
3.4 Majority Voting	31
4 Experimental Results	32
4.1 Experimental Setup	32
4.1.1 Dataset curation	32
4.1.2 Large Language Model Configuration	34
4.2 One step utterance generation	35
4.2.1 Sentence generation improves with model size and context length- Up to a Limit	35
4.3 Thought generation for theory of mind reasoning	37
4.3.1 The best generated responses based on subtext thoughts im- prove upon baseline	38
4.3.2 Selecting Responses Based on a Single Subtext Underperforms Baseline	39
4.3.3 Self-Reflective and Goal-Oriented Thoughts Rank Higher . . .	39
4.3.4 Sampling multiple times from a varied distribution can still lead to improvement in BERTScore	40
4.3.5 Using random ideas in subtext does not have clear trends in scores	43
4.4 State Evaluator	44
4.4.1 Choosing the best ranked thought improves BERTScore . . .	44

4.4.2	Both LLM-based state evaluators are not effective in selecting the best subtext	45
4.5	Evaluation over multiple shuffles	47
4.5.1	Evaluation by Scoring and by Voting are biased towards certain numbers	47
4.5.2	Evaluation by Scoring is more consistent than by Choice . . .	49
4.6	Majority Voting	49
4.6.1	Majority voting does not improve performance	50
4.6.2	High consistency on evenness index leads to lower variation of BERTScore of output	50
4.7	Overall findings	52
5	Discussion	54
5.1	Future work	54
5.1.1	Language Model Capabilities	54
5.1.2	Context Representation	55
5.1.3	Evaluation Metrics.	55
5.1.4	Thought Representation and Tree of Thought.	56
5.1.5	Filtering to Target Turns that Require Reasoning	56
5.2	Limitations	57
6	Related Explorations-Neural Encoding using Sentence and Word Embeddings	58
6.1	Data collection and pre-processing	59
6.2	Method: Ridge regression	60
6.3	Method: Banded Ridge regression	61
6.4	Results	62

6.4.1	Encoding with context and sentence embeddings outperforms static embeddings, but not contextual embeddings	62
6.4.2	Sentence Embeddings take up variance in Joint Context+Sentence encoding models	63
6.4.3	Encoding with static and sentence embeddings outperforms static embeddings, but not contextual embeddings”	64
6.4.4	Joint Sentence+static embeddings enable encoding for best electrodes over a greater period of time compared to only static embeddings	65
6.4.5	Overall findings	67
7	Future Directions	69
8	Conclusion	71
A	Code Implementation	73
A.1	System prompt used	73
A.2	Prompt Generation	73
A.3	Dataset filtering and setup	79
A.4	Inference generation	89
A.5	Language model hyperparameters	126
A.6	Banded Ridge Regression Code	127
B	Engineering Standards	146

List of Figures

2.1 Transformer Architecture from Attention is All You Need (Vaswani et al., 2017)[1]	9
3.1 An illustration of the single-step generation procedure, where the model generates a response conditioned using a context length of 1 (only on the previous utterance in the conversation.)	22
3.2 A diagram representing the Tree of Thought and evaluation methodology for utterance generation. Orange highlighted cells represent the actual states searched through using the LLM. Yellow highlighted cells represent the set of potential model responses. Red arrows represent LLM generation while the dashed arrows represent comparison using BERTScore.	23
4.1 A histogram of the percentage of sentences spoken by the dominant speaker in each conversation	34
4.2 Average sentence similarity scores between model generated and human reference turns of conversation for differing amounts of context and different model sizes	37
4.3 Box plot showing the BERTScores for all ranked ToM subtext generations (1st, 2nd, 3rd best, etc.) for each sample.	38

4.4	Box plot showing the BERTScore distribution for each different subtext over all samples	39
4.5	Bar chart showing the Average Reciprocal Rank for each theory of mind subtext. The possible range is between 0 and 1, with a perfect score of 1 being achieved by having a rank of 1 over all samples	40
4.6	BERTScore outcomes for different generation strategies with various perturbation conditions. Dark blue bars represent the average best score per sample, light blue bars the worst and thick black bars the score from randomly sampling a subtext each sample. The yellow bar shows the average score with a simple label. The black dashed line is the model baseline.	41
4.7	Bar chart showing the Average Reciprocal Rank for each random idea subtext. The possible range is between 0 and 1, with a perfect score of 1 being achieved by having a rank of 1 over all samples	43
4.8	Box plot showing the BERTScore distribution for each different subtext over all samples	44
4.9	Box plot showing the BERTScores for all ranked random idea subtexts (1st, 2nd, 3rd best, etc.) for each sample.	45
4.10	Scatter plots of the BERTScore based on model response from the best ranked thought (left) and a randomly selected thought (right) against baseline response	46
4.11	Scatter plots of the BERTScore based on model response from the voting over all possible subtexts (left) or scoring each subtext individually (right) against the baseline response.	46

4.12 Violin Plot of the BERTScore for different State Evaluators. The dark blue line represents the mean of the distribution while the black lines represent the datapoints 1.5 times the interquartile range (IQR) from the first and third quartile.	47
4.13 The frequency each numerical index in the list is chosen when the state evaluator votes over states	48
4.14 The frequency specific digits appear when the state evaluator scores each option. Calculated over 10 shuffles.	48
4.15 The continuous density plot of evenness index of all samples for different shuffles. An evenness index of 0 would indicate perfect consistency, while an index of 1 would indicate a perfectly random/uniform evaluation process.	49
4.16 The histogram of maximum number of agreements on the best subtext over different shuffles for all samples. A score of 10 would indicate perfect consistency while a uniform/random evaluator would score 1.67 on average.	50
4.17 A scatter plot of the BERTScore of the majority vote meta-state evaluator voting by voting and by scoring. General points are in grey while coloured points refer to high consistency examples with number of agreements ≥ 3	51
4.18 A scatter plot of the BERTScore of the majority vote meta-state evaluator voting by voting and by scoring. General points are in grey while coloured points refer to high consistency examples with evenness index > 0.55	51

4.19 Violin Plot of the BERTScore for the majority vote meta-State Evaluators based on scoring and voting. Also plotted are the distributions of high consistency examples highlighted in the scatter plots in Figures 4.18 and 4.17	52
6.1 Scatter plots of encoding performance (correlation) for joint (sentence+contextual word) embeddings against static and contextual embeddings	63
6.2 The encoding performance using either the contextual embedding or sentence embedding feature space plotted against each other	64
6.3 Correlation scores per electrode mapped onto the brain for joint (contextual word + sentence), contextual word and sentence embeddings	64
6.4 Scatter plots of encoding performance (correlation) for joint (sentence+static word) embeddings against static and contextual embeddings	65
6.5 The encoding performance using either the contextual embedding or sentence embedding feature space plotted against each other	66
6.6 Correlation scores per electrode mapped onto the brain for joint (static word + sentence), static word and sentence embeddings	66
6.7 The correlation of model and actual neural signal plotted over time and averaged over best performing electrodes. Black lines show the encoding done using a split feature space after banded ridge regression for static embeddings (left) and sentence embeddings (right). The blue line represents the joint embedding performance. Electrodes where the joint (static+sentence) correlation had a performance of more than 0.15 were selected.	67

6.8 The correlation of model and actual neural signal plotted over time and averaged over best performing electrodes. Black lines show the encoding done using a split feature space after banded ridge regression for contextual embeddings (left) and sentence embeddings (right). The blue line represents the joint embedding performance. Electrodes where the joint (static+sentence) correlation had a performance of more than 0.15 were selected	68
6.9 The correlation of model and actual neural signal plotted over time and averaged over best performing electrodes. Embeddings used in the model in order of envelop size are contextual, joint (contextual+sentence), joint (static+sentence) (plotted twice), static, and split static.	68

List of Symbols

$\mathcal{D} = \{D_1, D_2, \dots, D_J\}$: Collection of dialogues, where each dialogue D_j is a sequence of utterances between two speakers.

$D = \{u_1, u_2, \dots, u_N\}$: A dialogue between two speakers, represented as a sequence of N utterances.

$c_i = \{u_{i-w}, \dots, u_{i-1}\}$: Context for utterance u_i , consisting of the w preceding utterances, or $\min(i - 1, w)$ if $i - 1 < w$.

w : The number of preceding utterances included in the context for generating each response. This can be referred to as the context length or width.

$C = \{c_{w+1}, c_{w+2}, \dots, c_N\}$: Set of all contexts used for response generation, corresponding to turns $i = w + 1$ to N .

$R = \{r_{w+1}, r_{w+2}, \dots, r_N\}$: Model-generated responses for turns $i = w + 1$ to N , conditioned on context c_i .

$L = \{l_{w+1}, l_{w+2}, \dots, l_N\}$: Human-written reference responses for turns $i = w + 1$ to N .

Note: For turns $i \leq w$, the model does not generate responses due to insufficient context.

S : Set of scores produced by the evaluation metric, $\{s_1, s_2, \dots, s_N\}$.

Q : Set of human quality annotations, $\{q_1, q_2, \dots, q_N\}$.

p_θ : A pre-trained language model (LM) with parameters θ .

x, y, z, s, \dots : Language sequences, where $x = (x[1], \dots, x[n])$ represents a collection of tokens, with each $x[i]$ being a token.

$p_\theta(x)$: The probability of sequence x under the pre-trained model p_θ , given by:

$$p_\theta(x) = \prod_{i=1}^n p_\theta(x[i] \mid x[1 \dots i]).$$

$\bar{X}, \bar{Y}, \bar{Z}, \bar{S}, \dots$: Collections of language sequences, represented in uppercase with a bar.

$p_\theta(y \mid \text{prompt}_{IO}(x))$: The probability of generating output y from input x wrapped in a prompt with task instructions and/or few-shot input-output examples.

$p_\theta^{\text{prompt}}(\text{output} \mid \text{input})$: Simplified notation for the probability of output given the input with a prompt.

prompt: A prompt that contains task instructions and/or few-shot input-output examples used to generate the next response.

$T = \{\tau_0^{(m)}\}_{m=1}^M$: The collection of subtexts $\tau_0^{(m)}$, indexed by m from 1 to M , where each $\tau_0^{(m)}$ represents a potential subtext idea based on a theory of mind taxonomy.

$G(p_\theta, s, k)$: Thought generator that generates k candidate thoughts based on the pre-trained language model p_θ and tree state s .

$s = [x, z_1, \dots, z_i]$: A tree state, where x is the initial input, and z_1, \dots, z_i are the previous thoughts generated.

LLM- Large Language Model

ToM -Theory of Mind

Chapter 1

Introduction and Background

Large Language Models (LLMs) have recently exploded in popularity due to their ability to model language and their emerging, but imperfect, capacity for reasoning [2]. One application of LLMs is in human neuroscience, where they have been used to model neural signals from the literal content of words spoken [3]. This thesis aims to leverage recent advances in LLM inference-time reasoning to not only model the literal content of words during language production and comprehension but also subtext, meaning possible unspoken thoughts and inferences that arise during communication. Towards investigating this, this thesis presents two projects. The first applies an in-context learning technique, tree of thought, to a novel domain by aiming to improve dialogue generation over a naturalistic dataset by reasoning over potential subtextual thoughts — unspoken beliefs, desires, or intentions — inspired by theory of mind. The second builds on prior neural encoding projects in the Hasson lab, developing models that leverage both word-level and sentence-level embeddings to encode neural activity. The rest of this section will give a summary of related work and an outline of this report.

1.1 Reasoning versus language abilities in LLMs

It is under debate whether LLMs can truly reason in an abstract way to solve problems or whether they learn narrow-non-transferable skills for task-solving [4]. Nevertheless, they have had increasingly strong performance on reasoning benchmarks such as math, logic and reasoning [5] as well as language [6, 7]. This leads to the question of whether intelligence can be gained solely through mastery of language. This is related to a common fallacy conflating language and thought ability, where the ability to use language is neither a necessary nor sufficient condition for reasoning (Mahowald, Ivanova et al. coined the difference between knowledge of linguistic rules and understanding of language as formal linguistic competence and functional linguistic competence [8]). In humans, language and reasoning have been suggested to use different brain networks [9, 10], and individuals with aphasia who exhibit severe linguistic deficits still have other intact cognitive abilities such as social and mathematical reasoning [11]. Nevertheless, language can still bootstrap cognition, facilitating processes such as inner speech[12] and enabling faster development of social reasoning[13]. However, it seems unlikely that language models can acquire robust reasoning ability solely through next-token prediction during pre-training -motivating the development of various inference-time techniques aimed at enhancing LLM reasoning.

1.2 Inference time generation methods

Inference time generation methods are a method to improve the performance of LLMs without the large amounts compute and data needed for pre-training. This can be used alongside other methods of post-training improvement such as fine-tuning, reinforcement learning from human feedback and pure reinforcement learning in models such as o1 by OpenAI [14], R1 by DeepSeek [15], and QwQ by Qwen Team[16].

Inference-time methods come in many flavours, they can explore a single reasoning path often by optimizing a single prompt, such as in Chain of Thought [17, 18], multiple reasoning paths with or without multiple agents as in multi-agent debate [19] or boosting of thoughts [20] or can include an iterative or memory mechanism as in Reflexion [21] or buffer of thoughts[22]. However, there is no technique that consistently outperforms others in all tasks, so the implemented method must be chosen according to the specific task [23].

In this project, the strategy implemented is tree of thought [24]. This has the benefit of having multiple reasoning chains to track different ideas, as well as the flexibility to modify the tree’s complexity and structure as needed for deeper or shallow levels of reasoning. Following an analysis of bias, a majority voting mechanism was also implemented [25].

1.3 Theory of mind in LLMs

Theory of Mind (ToM) refers to the cognitive capacity to infer the mental states—such as beliefs, desires, and intentions—of others, allowing us to predict and interpret their behavior. Traditionally, ToM has been modeled as a predictive model of others’ actions based on their mental states. However, this approach often struggles with flexible generalization to novel situations and is limited in its use for complex planning tasks that involve many possible actions. A more powerful alternative frames ToM as a causal model—an abstract, structured representation of mental states that supports efficient reasoning, planning, and decision-making over a wide space of possibilities [26].

One domain where this richer model of ToM becomes essential is pragmatic language use. For instance, speakers frequently tailor their words to achieve specific emotional or social outcomes, such as using indirect or softened language to avoid

offending someone. This type of communication requires balancing honesty, social norms, and an estimation of the listener’s mental state. The Rational Speech Act (RSA) framework captures this complexity by modeling communication as a Bayesian reasoning process. Within RSA, speakers choose utterances by anticipating how listeners will interpret them, and listeners interpret utterances under the assumption that speakers are cooperative and informative. ToM is also central in interpersonal affect regulation, where individuals deliberately attempt to influence others’ emotional states—calming them down, cheering them up, or helping them regain control. This kind of planning involves both empathy and strategic reasoning, highlighting the real-world complexity of applied ToM.

The extent to which large language models (LLMs) can perform ToM reasoning is the subject of ongoing debate [27, 28]. To address this, researchers have introduced a number of benchmarks to systematically evaluate ToM abilities in LLMs [29, 30, 31, 32]. While LLMs show strong results on certain tasks, their overall ToM performance remains inferior to that of humans, and is often accompanied by signs of overfitting [33].

Current datasets for evaluating Theory of Mind (ToM) in large language models (LLMs) suffer from notable limitations. Rich, human-authored narratives that naturally involve ToM reasoning are both rare and costly to curate. Consequently, researchers often rely on synthetic, template-based datasets that are overly simplistic and lack ecological validity. These datasets frequently include explicit mental state language (e.g., “Amy thinks that...”), which makes the reasoning task unrealistically easy and reduces the diagnostic value of the evaluation. Moreover, they tend to focus narrowly on action prediction, failing to capture the broader and more nuanced reasoning processes involved in real-world ToM. One dataset that attempts to address these shortcomings is SimpleToM [34], which includes both types of reasoning: the recognition of mental states (explicit ToM) and the application of that knowledge

to predict actions or responses (applied ToM). They found that while LLMs tend to perform well on explicit ToM tasks, their performance on applied ToM remains weak—highlighting a critical gap and a central challenge relevant to our domain. Additionally, ToM performance in LLMs can vary significantly depending on the scenario, highlighting the need for methods that enhance LLMs’ capacity for robust, context-sensitive reasoning rather than relying solely on narrow or superficial cues.

Several methods have been proposed to enhance Theory of Mind (ToM) performance at inference time [33, 35], but many of these approaches rely on few-shot examples or rigid assumptions that hinder scalability. While techniques such as custom chain-of-thought prompting can guide models toward more accurate ToM reasoning, they are often fragile and require extensive manual tuning [34]. Structured frameworks have also been introduced to support ToM reasoning. For instance, Wu et al. [36] developed COKE, a system that reasons over a knowledge graph by instantiating ToM as a collection of manually verified cognitive chains. These chains represent human mental activities in specific social contexts, including their behavioral and affective responses. More recently, Kim et al. [37] proposed a similar approach to the one in this project, propagating a series of hypothetical thoughts to explain and infer human actions with a method inspired by the sequential Monte Carlo algorithm.

1.4 Use of Language Models to study processing in the brain

A computational model’s neural plausibility can be evaluated by assessing whether its internal representations align with patterns of brain activity elicited by corresponding stimuli. If a model consistently predicts neural responses to novel inputs, this suggests it captures key aspects of the brain’s representational structure. This approach has been used in decoding brain activity in the visual pathway, using techniques

such as motion energy models, Bayesian decoders, and deep learning frameworks [38]. Similarly, large language models (LLMs), which have achieved remarkable success in modeling human language, raise the question: how closely do their internal representations mirror those of the human brain?

To quantify representational similarity between LLMs and the brain, researchers typically employ two types of models: encoding models, which assess how well LLMs predict neural activity in response to linguistic stimuli, and decoding models, which infer the stimulus that produced a given neural response [39]. Empirical studies have revealed overlaps in representational structure—particularly between contextual embeddings from LLMs and decontextualized word embeddings. For example, fMRI studies [40] and intracranial recordings [41, 42] indicate that contextual embeddings better align with brain representations than static embeddings.

To deepen our understanding of this shared structure, Tuckute et al. [39] propose two key strategies for leveraging LLMs in the study of brain function. The first involves systematically probing LLMs by varying their architecture, training data, and behavioral outputs to isolate factors that align with brain activity. The second involves using LLMs as tools for *in silico* neuroscience experiments, allowing researchers to simulate and test hypotheses derived from empirical data.

These approaches have been extended to higher-level cognition. For instance, one study used fMRI to measure brain responses to 1,000 diverse sentences and showed that an encoding model based on GPT could predict the magnitude of the neural response to each sentence [43]. Interestingly, for sentences describing others’ mental states, the model did not explain significantly more variance (beyond surprisal), suggesting that language-related brain regions are not modulated by social content. This finding supports the view that the brain’s language network is functionally distinct from the theory of mind network.

One intriguing direction for future research is whether in-context learning—a

method known to improve reasoning in LLMs—could be used to simulate or even enhance models of human reasoning. Such approaches may offer novel tools to investigate the mechanisms of cognitive processing in the brain.

1.4.1 Sentence embedding encoding

When working with fMRI signals, sentence-based encoding and decoding is more common due to the longer time duration of the signals. Early research demonstrated the ability to encode sentences using embeddings [44], and further analysis has shown that syntactic structure, rather than lexical-semantic content, is the primary contributor to the similarity between artificial neural network representations and brain responses in the language network [45]. Sentence embeddings are typically generated by averaging the word embeddings in the sentence or using the representation of the last sentence token (often the final period token ".") as a sequence summary.

1.5 Report Overview

In the following chapters, we will review the relevant literature and methodological background for natural language processing using transformers and dialogue modeling metrics, go over the methodology and discuss experimental results on the naturalistic 24/7 Conversations dataset, including future work and current limitations. Finally, we go over the method and results for the sentence encoding project and explore strategies for integrating it with the dialogue modeling work presented in this report.

Chapter 2

Methodological Background

2.1 Natural Language Processing using Deep Neural Networks

Natural language processing involves many different tasks which includes classification, text generation and sentence pair matching. A variety of deep neural network models can be used, including recurrent neural networks, long short-term memory networks, and most commonly now transformers. Modern transformer-based models based on the original architecture by Vaswani et al., 2017[1] give us the best performance on the relevant tasks for this project which requires text generation and sentence similarity scoring.

An overview of the transformer architecture is given in Figure 2.1. While a detailed account is covered in textbooks [46], the following section will give an overview of the key aspects of the models that we will repeatedly refer to.

Tokenisation

In NLP, a token is a single unit of text, such as a word, subword, or character, depending on the tokenization method used. For example, in the sentence "I love




Figure 1: The Transformer - model architecture.

Figure 2.1: Transformer Architecture from Attention is All You Need (Vaswani et al., 2017)[1]

NLP!”, the tokens could be [”I”, ”love”, ”NLP”, ”!”] (word-level tokenization) or [”I”, ” lo”, ”ve”, ” NLP”, ”!”] (subword tokenization like in Byte Pair Encoding). Tokenization is a fundamental step in processing text for machine learning models. Different models may have different tokenization schemes with different vocabulary sizes and elements.

Static Embeddings

Static embeddings map each unique word in a vocabulary to a dense d-dimensional vector that capture meaning and similarity. This can be thought of as an improvement

over one-hot encoding, where each word is represented as a large sparse vector with no meaningful relationship between each vector. In static embeddings, words with similar meanings have similar vectors and tend to cluster together in the embedding space. Furthermore, certain directions within this space can even capture specific relational meanings, for example, the vector difference between king and queen is similar to that between man and woman. The key attribute of static embeddings is that each word has a single fixed vector, regardless of its surrounding context. Common methods of creating these embeddings include Word2Vec and GloVe.

Contextual Embeddings

By contrast, with contextual embeddings, such as those learned by masked language models like BERT, each word w will be represented by a different vector each time it appears in a different context. The embeddings are made contextual through the self-attention mechanism.

The attention mechanism then measures how much focus each token should give to every other token. This is done by computing the similarity between the Query (Q) of a token and the Keys (K) of all tokens through the dot product QK^T . In the decoder only part this computation is masked, so can only look at the past tokens. The result is then scaled by $\frac{1}{\sqrt{d_K}}$ to normalize the values, helping stabilize gradients during training. Then the softmax function normalises the attention scores to ensure they sum to 1, converting raw similarity scores into a probability distribution. The result of this determines how much weight each token should assign to others in the sequence, and finally these computed attention weights are applied to the Value (V) vectors. Tokens with higher attention scores contribute more to the final representation of a given token. The weighted sum of values produces a new representation for each token, capturing contextual dependencies dynamically. This mechanism can be applied multiple times sequentially and for multiple heads. So when we refer to

”contextual embeddings” this refers to weighted value vector after attention is applied using the Q and K matrices.

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_K}} \right) V \quad (2.1)$$

Next token prediction models

For generation of dialog I will use a decoder-only transformer e.g. DeepSeek, Mistral, GPT. In next-token prediction models like decoder-only transformers, text is generated one token at a time based on previous context. These models are also trained autoregressively. After processing input tokens, the hidden state of the last token in the final layer $h_N \in \mathbb{R}^{1 \times d}$ is passed through a linear layer using an unembedding matrix $W \in \mathbb{R}^{d \times V}$ (called the language modeling head) to produce *logits* $z_i \in \mathbb{R}^{1 \times V}$ — raw scores for each token in the vocabulary.

$$z_i = h_N W$$

These logits are converted to probabilities using the softmax function, producing $P(\text{next token} | \text{context})$:

$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_{j=1}^V e^{z_j}}$$

The model can then either pick the most likely token (greedy decoding) or sample from this distribution (temperature sampling/ to generate the next token.

Masked word prediction

As opposed to autoregressive decoder models, masked language models like BERT operate by randomly masking words in sentences and training the model to predict these missing words using surrounding context. This approach falls under denoising

training objectives, where inputs are deliberately corrupted through masking, substitution, reordering, deletion, or insertion, and the model learns to reconstruct the original input by minimizing cross-entropy loss between its predictions and the true values of the missing tokens.

This bidirectional approach enables learning rich contextual word representations applicable to various NLP tasks such as text classification, named entity recognition, and sentence similarity assessment. Often called encoder-only models, they excel at producing contextualized embeddings for interpretative tasks rather than generating text, making them particularly valuable for methods like BERTScore.

2.2 Text similarity metrics

For this project, we need a reference-based metric(s) to compare the similarity of the model generated response to the human reference response. In practise, human evaluation (typically from crowd-workers over experts for dialog) is used a gold standard for text and dialog evaluation with various strategies used for data collection and different questions asked to evaluators [47]. However, collecting this data from humans is time-intensive and expensive, so it is typically only used for a final evaluation, while automatic evaluation metrics are used during development to optimise model design and the choice of hyperparameters.

The ideal automatic metric maximally correlates with human judgement, however most metrics struggle to achieve this goal [48] hence metric-creation is still an open problem itself[49][50]. To improve accuracy or cost a hybrid approach may also be used, potentially combining multiple automatic metrics [51][52] or human and automatic metrics [53][54]. There have been several reviews on the several text similarity metrics [51] [55], and here we summarize their findings.

Formally, we can describe dialog evaluation metrics as follows: Given a dialog con-

text $c = \{c_1, \dots, c_m\}$, model response $r = \{r_1, \dots, r_n\}$, and human-written reference response $l = \{l_1, \dots, l_k\}$, the goal is to learn a function:

$$f : (c, r, l) \rightarrow s \quad (2.2)$$

that evaluates the generated response. (A metric is said to be reference-free if it can be expressed as a function of only c and r .)

These metrics are assessed by comparing them to human judgment. Concretely, a human annotator (or several annotators) scores the quality of a given response conditioned on the dialog context:

$$(c, r) \rightarrow q. \quad (2.3)$$

Given the scores produced by a particular metric, $S = \{s_1, s_2, \dots, s_n\}$, and the corresponding human quality annotations, $Q = \{q_1, q_2, \dots, q_n\}$, we can measure the performance of the metric by calculating the correlation between S and Q .

There are some important features and challenges to consider in the selection of this metric. The first is that the surface form representation is less important than the semantic meaning of the text, for example ””The cat sat on the mat.”, ”A mat was where the cat chose to sit.”, ””The cat settled on the mat.” , ”The mat became the cat’s resting place.” should all receive similar similarity scores. Furthermore, the metric should be robust without changing greatly word small changes in word order or synonyms. Additionally, metrics can typically be evaluated at the utterance level as well as the system level by computing the average utterance-level score across many utterances [56], and having a strong correlation at the system-level is still useful even if the metric per utterance is noisy or has a lower correlation with human judgement.

Achieving these goals can be complicated, especially because of the long-tailed nature of language. You can not train a model on every possible combination of

sentences because there are too many, which is why supervised trained metrics tend to overfit the dataset they were trained on [57]. Additionally, it is important to consider all metrics will have limitations and optimising a particular metric can lead to unintended effects [58][59].

In this section we go over the various options for evaluation of natural language text generation for dialogue, and explain the workings of BERTScore -the final metric selected. In the following sections, existing metrics can be broadly categorized into using n-gram matching, edit distance, embedding matching, or learned functions.

2.2.1 N-gram matching

N-gram matching metrics are based on word or character based overlap between the prediction and reference. The most popular example is BLEU, with other metrics including METEOR AND NIST. While originally used in translation to detect the textual similarity when changing language, they can also be used for evaluating textual similarity within the same language.

As an example, the BLEU score for a corpus of candidate translation sentences is a function of the n-gram word precision over all the sentences combined as well as a brevity penalty computed over the whole corpus [46]. Here n-gram refers to a sequence of token-length n. N-gram precision would be the percentage of n-gram tokens in the candidate translation that also occur in the reference translation. The harmonic mean of the unigrams, bigrams, up to 4-grams is used in the calculation of BLEUscore. Since it is a word-based metric, it is very sensitive to word tokenization, making it impossible to compare different systems if they rely on different tokenization standards. It also cannot detect semantic equivalence, so two sentences with the same meaning but mostly different characters can receive an incorrect low score. Because of this, it is not generally used for sentence similarity evaluation [60] and it is not a good fit for our uses.

2.2.2 Edit distance

Minimum edit distance/word error rate is another common metric with its origin in automatic speech recognition. However, since it is commonly used in speech recognition this equally penalises the completely wrong word in a sentence as well as a close synonym. Therefore, it is also misses capturing some of the idea of the semantic meaning of sentences. This lack of semantic meaning makes it unsuitable for this project.

2.2.3 Trained metrics

Various metrics are trained to optimise correlation with human judgement. The input to these models can include other different metrics such as character n-grams in BEER which used a regression model or ADEM which uses the context and the reference response as inputs to a recurrent neural network trained to predict appropriateness ratings by human judges. These methods usually have the highest correlation to human judgement [61]. However, these methods require expensive human judgments as supervision for each dataset and also generally have poor generalisation to new domains and data. These metrics have been shown to not be robust [57]. For example with ADEM, in 48.66% of cases the predicted score increased when the generated response was reversed. In 86.93% of cases the predicted score increased when the generated response was replaced with a dull dummy response.

2.2.4 Perplexity

Perplexity is a metric which was first introduced by Jelinek et al. in 1977 while working on automatic speech recognition [62], which was its original use case [63]. However, now it is most commonly now it is used to decide how to sample from a language model for text generation [64][65].

Given the probability assigned by the language model to a sequence \mathbf{x} as:

$$p_{\theta}(\mathbf{x}) = \prod_{i=1}^n p_{\theta}(x[i] \mid x[1 \dots i-1])$$

Perplexity (PPL) can be derived from this probability as;

$$\text{PPL}(\mathbf{x}) = p_{\theta}(\mathbf{x})^{-\frac{1}{n}} = \exp \left(-\frac{1}{n} \sum_{i=1}^n \log p_{\theta}(x[i] \mid x[1 \dots i-1]) \right)$$

where n is the sequence length. Perplexity is highly interpretable, because it quantifies how "surprised" the model is by the given sequence—lower perplexity means the model assigns higher probability to the observed sequence. Perplexity can range from a theoretical value of 1, where the language model is completely certain of its prediction to infinity, where the language model is completely uncertain.

One example of its use in sampling is the maximum mutual information (MMI) scoring function [66] [67] has been used to predict such "boring" responses. MMI employs a pre-trained backward model to predict source sentences from given responses, i.e., $P(\text{Source}|\text{target})$. The probability of $P(\text{Source}|\text{Hypothesis})$ can be used to rank the quality of hypotheses. Intuitively, maximizing backward model likelihood penalizes the bland hypotheses such as "that's interesting" or "ok", since the ability to discern the source sentence from such a target sentence is low, because these hypotheses can be associated with many possible queries. Another use case has been to evaluate the coherence of two consecutive sentences [68] for .

Perplexity is not as commonly used for evaluation of dialogue models. One reason for this is that this metric must be model-dependent, and it is better to avoid situations where the model generating the response also evaluates its quality because this would lead to bias [48]. Perplexity is also not comparable between language models with different vocabularies [63]. Moreover, perplexity is slightly different to the other metrics because it does not directly measure textual or semantic similarity between

two samples. Even if two texts have similar perplexity values, they could be semantically very different. It could still be used to evaluate which subtext or thought (or no thought) most likely results in a given human reference, however it does not give us a direct measure of how well a model response matches the reference response. Hence, I did not prioritize implementing this metric.

2.2.5 Embedding based metrics

As discussed in Section 2.1, token (word) embeddings are learned dense token (word) representations that can be represented in a continuous vector space where semantic similarity is encoded so that tokens (words) with similar meanings are closer in the vector space. Embedding based metrics tend to have the highest correlation to human judgment based on their ability to match semantic meaning between candidate and reference sentences. Hence, this is what we chose to use in the rest of this report.

The general method of computing similarity is between two tokens is using the cosine similarity between two embedding vectors v and w as:

$$\text{cosine}(v, w) = \frac{v \cdot w}{\|v\| \|w\|}$$

As mentioned in previous section, these embeddings can either be static and learned from large text corpora such as Word2Vec or GloVe, or they can be from contextual models like BERT. This is important for measuring semantic similarity of texts, since contextual embeddings allow the same word to have a different representation depending on its surrounding words, allowing the embeddings to capture nuanced meanings, polysemy, and relationships within the sentence context. For example, the word "bank" will have different embeddings in "river bank" versus "financial bank," reflecting its specific usage in each sentence. Static embeddings will measure the similarity of the word "bank" in these two sentences as exactly the same. Hence, using

contextual embeddings for this task tends to do better[69].

The specific metric chosen for this project is BERTScore based on its superior performance for measuring textual similarity[70][69]. This metric is based on BERT, which is an encoder that creates contextual embeddings. Encoder-only models like BERT are known for producing high-quality, robust embeddings, making BERTScore an ideal choice for this task.

To explain how BERTScore works in detail, the BERTSCORE algorithm first passes the reference sentence x and the candidate sentence \tilde{x} through BERT, computing a BERT embedding for each token x_i and \tilde{x}_j .

Given this sequence of tokens in a reference sentence $x = \langle x_1, \dots, x_k \rangle$ and a candidate sentence $\tilde{x} = \langle \tilde{x}_1, \dots, \tilde{x}_l \rangle$, the similarity is calculated using a greedy matching algorithm.

Each token in x is matched to a token in \tilde{x} to compute recall, and each token in \tilde{x} is matched to a token in x to compute precision (with each token greedily matched to the most similar token in the corresponding sentence). BERTSCORE provides precision and recall (and hence F_1):

$$R_{\text{BERT}} = \frac{1}{|x|} \sum_{x_i \in x} \max_{\tilde{x}_j \in \tilde{x}} (x_i \cdot \tilde{x}_j)$$

$$P_{\text{BERT}} = \frac{1}{|\tilde{x}|} \sum_{\tilde{x}_j \in \tilde{x}} \max_{x_i \in x} (x_i \cdot \tilde{x}_j)$$

$$F_{\text{BERT}} = \frac{2 \cdot P_{\text{BERT}} \cdot R_{\text{BERT}}}{P_{\text{BERT}} + R_{\text{BERT}}}$$

Additionally, BERTScore includes baseline scaling. From the formula mentioned above, scores would have the numerical range of cosine similarity between -1 and 1, however scores are observed in a smaller range in practice. To improve readability, BERTScore rescales its output with respect to an empirical lower bound b . This is

computed by finding the average score between 1M candidate-reference pairs made by grouping two random sentences in the Common Crawl monolingual datasets. This lower bound is then used to rescale BERTScore linearly. The rescaled BERTScore is given as:

$$\hat{R}_{\text{BERT}} = \frac{R_{\text{BERT}} - b}{1 - b}$$

Chapter 3

Methodology

In this chapter, we will formalize the dialog generation task and describe the methods implemented to improve on this task. Then I will outline the generation strategies used to create a baseline score and the tree-of-thought and majority voting strategies implemented to aim to improve upon this score.

3.1 Problem setup

Given a dataset of conversational turns, we aim to construct a set of dialogues between two speakers. Each sample consists of a context, represented as a sequence of previous dialog turns c_1, \dots, c_m , and a label, which is the true next utterance written by a human, denoted as l_1, \dots, l_k .

The objective is to improve the ability of a pre-trained language model p_θ to generate a model response r_1, \dots, r_n that closely matches the human label response, l_1, \dots, l_k . This similarity is measured using a metric such as BERTScore, which produces a set of scores $S = s_1, s_2, \dots, s_n$ comparing the generated response to the human-written label.

Our approach is to enhance the model’s generation process by conditioning its response on reasoning paths which are guided by prompts that consider a set of po-

tential subtexts, $\tau_0^{(i)M}_{c=1}$, which should help the model generate output that reflects the thought process the human speaker likely followed. By incorporating these reasoning paths into the model’s response generation, we aim to produce more coherent and contextually appropriate responses that better match the human label. If successful, this would not only improves the quality of the model’s output in natural dialogue settings but also demonstrates the model’s ability to understand and reflect the underlying mental states involved in conversation.

3.2 Generation strategies

This work explores two distinct next-utterance generation approaches. The first is a straightforward baseline that conditions only based on previous context. The second seeks to improve upon this baseline by adopting a tree of thought reasoning framework, in which the model explicitly reflects on the next speaker’s possible mental states and intentions before generating a response.

3.2.1 Single step utterance generation

This section describes our procedure for generating single-step utterances using large language models, and introduces the notation we use throughout. This notation based on the original Tree of Thought paper [24], with adaptations to suit the specific context of our work.

We begin by formalizing the general language generation process, before describing how we apply this framework to obtain a baseline generation score without any in-context learning. Let p_θ denote a pre-trained LM with parameters θ , and lowercase letters x, y, z, s, \dots to denote a language sequence which is a collection of tokens, i.e., $x = (x[1], \dots, x[n])$ where each $x[i]$ is a token. Then $p_\theta(x) = \prod_{i=1}^n p_\theta(x[i] \mid x[1 \dots i])$. We also use uppercase letters with a bar $\bar{X}, \bar{Y}, \bar{Z}, \bar{S} \dots$ to denote a collection of

language sequences.

We can turn a problem input x into an output y using a language model by wrapping input x with a prompt containing task instructions and/or few-shot input-output examples. We can write this as $y \sim p_\theta(y | \text{prompt}_{IO}(x))$. We can simplify this notation further by saying that $p_\theta^{\text{prompt}}(\text{output} | \text{input}) = p_\theta(\text{output} | \text{prompt}(\text{input}))$

Now we can describe the specific implementation approach. The relevant variables here are the context c and the prompt prompt . We can tie this together as:

$$r_1^{(i)} \sim p_\theta^{\text{prompt}}(r_1^{(i)} | c) \quad (3.1)$$

Figure 3.1 depicts what this generation process would look like where the context c consists of a single preceding utterance. We will create our baseline based on the generation procedure outlined in Equation 3.1. The prompt will be manually iterated to produce the best performance while the effect of changing context length will be included in the results.




Figure 3.1: An illustration of the single-step generation procedure, where the model generates a response conditioned using a context length of 1 (only on the previous utterance in the conversation.)

3.2.2 Tree of thought




Figure 3.2: A diagram representing the Tree of Thought and evaluation methodology for utterance generation. Orange highlighted cells represent the actual states searched through using the LLM. Yellow highlighted cells represent the set of potential model responses. Red arrows represent LLM generation while the dashed arrows represent comparison using BERTScore.

In this work, we adapt the original Tree of Thought (ToT) framework to the domain of dialogue generation. To structure our method clearly, we follow the organization of the original ToT paper, which frames a specific instantiation of ToT around four key design questions: (1) How should the reasoning process be decomposed into intermediate thought steps? (2) How should potential thoughts be generated from each state? (3) How should states be heuristically evaluated? (4) What search algorithm should be used to navigate the tree? In the next section, we describe how we adapt and extend each of these components to address the challenges specific to dialogue generation.

At a high level, our method generates intermediate thoughts based on a fixed

collection of subtexts of size M , which guide the generation of the final dialogue response. If all branches were explored, this would result in a tree of height 2 and width N as depicted in Figure 3.2, with the final layer corresponding to M candidate responses. However, in practice, the state evaluator $V(p_\theta, S)$ and search algorithm will be used to identify and follow the most promising branch without exploring the whole tree. All the explored states are highlighted in orange in Figure 3.2.

1. Thought Decomposition

While Chain of Thought samples thoughts without explicit decomposition, Tree of Thought designs and decomposes intermediate thought steps using the problem properties. In the following section, I will describe how the model uses a collection of subtexts $T = \{\tau_0^m\}_{m=1}^M$ to guide the reasoning process along specific paths. However, I will not focus on the decomposition of these thoughts, as we use a simple, low-depth tree structure, making further decomposition less necessary. This is because there is only one intermediate thought step before the model generates its response. Future work could explore increasing the depth of the tree to capture higher-order theory of mind or, alternatively, simplifying the thought process, as proposed in Chain of Draft [71].

2. Thought Generation

We require a thought generator $G(p_\theta, s, k)$ that generates k candidate thoughts based on the pre-trained language model p_θ and tree state s . Given a tree state $s = [x, z_1 \dots z_i]$, there are two strategies in the original paper to create a generator to create k candidates for the next thought step:

- (a) **Propose thoughts sequentially using a "propose prompt":**

$$[z^{(1)}, \dots, z^{(k)}] \sim p_\theta^{\text{propose}}(z_{i+1}^{(1\dots k)} | s).$$

(b) **Sample i.i.d. thoughts from a CoT prompt:**

$$z^{(j)} \sim p_{\theta}^{\text{CoT}}(z_{i+1}|s) = p_{\theta}^{\text{CoT}}(z_{i+1}|x, z_1 \dots z_i), \quad (j = 1 \dots k).$$

Proposing thoughts sequentially using a propose prompt works better when the thought space is more constrained. However, dialogue presents an inherently open thought space, suggesting that sampling i.i.d. thoughts from a CoT prompt would be more appropriate. The standard Tree of Thought approach would sample diverse reasoning paths independently like this, although sampling thoughts in this way may not lead to results that are interpretable from a theory of mind perspective.

To address these challenges, we instead initialize the first layer of tree states $S_0 = s_0^{(1\dots k)}$ by adding to the context c with a collection of subtexts $T = \{\tau_0^{(i)}\}_{i=1}^M$, which reflect different mental states or intentions the human speaker might have had. Specifically, each initial state is defined as $s_0^{(i)} = (c, \tau_0^{(i)})$. This can be interpreted as creating a manual thought generator $G_0(p_{\theta}, s_0 = c, k = M)$ which generates M new nodes given a context c based on the set of potential subtexts T , however this distinction is a matter of perspective rather than mechanism.

With the issue of creating thought diversity handled, the primary thought generator then does not need to calculate k new candidates given a tree state, instead it effectively sets $k = 1$ and is primarily concerned with propagating a node forward to produce intermediate thought(s) and then the final model response.

The final tree structure uses one intermediate thought which results in a tree structure of width M and height 2, which is summarized in Equation 3.2 and 3.3 below. Note that the prompts, and therefore the thought generators, for producing the intermediate thought and final model response are distinct because they create fundamentally different types of content. This contrasts with the original Tree of Thought framework, which does not explicitly distinguish between the generation of

intermediate thoughts and final output, instead treating all nodes as being of the same type.

$$z_1^{(i)} \sim p_{\theta}^{\text{prompt1}} \left(z_1^{(i)} \mid \tau^{(i)}, c \right) = G_1(p_{\theta}, s_0^{(i)}, 1) \quad (3.2)$$

$$r^{(i)} = z_2^{(i)} \sim p_{\theta}^{\text{prompt2}} \left(r_2^{(i)} \mid z_1^{(i)}, \tau^{(i)}, c \right) = G_2(p_{\theta}, s_1, 1) \quad (3.3)$$

In this implementation discussed in the results, a set of six subtexts are used, based on a framework of self versus other for feelings, goals and information seeking within a larger taxonomy of theory of mind reasoning [72]. The selection of these subtexts are discussed in more detail in Section 1. The exact prompts are specified in Appendix XYZ.

3. State Evaluator

Given a series of different thoughts we also need a state evaluator $V(p_{\theta}, S)$ to evaluate the progress they make towards solving the problem. Traditionally, this is used as a heuristic for a search algorithm to determine which states to keep exploring and in which order. We use the language model to reason over these states/thoughts to choose which one to generate a response for. I implemented both strategies that the original paper provides to evaluate states either independently or together:

- (a) **Value each state independently:**

$$V(p_{\theta}, S)(s) \sim p_{\theta}^{\text{value}}(v|s), \quad \forall s \in S,$$

where a value prompt reasons about the state s to generate a scalar value v (e.g., 1-10) or a classification (e.g., sure/likely/impossible) that could be heuristically turned into a value.

(b) **Vote across states:**

$$V(p_\theta, S)(s) = \mathbb{1}[s = s^*],$$

where the "best" state $s^* \sim p_\theta^{\text{vote}}(s^*|S)$ is voted for based on deliberately comparing different states in S using a vote prompt. This is useful when problem success is harder to directly value.

We additionally define two reference state evaluators for comparison. The first is the random evaluator which can act as a baseline and the second is a perfect (oracle) evaluator which would represent an ideal upper bound.

(c) **Random state evaluator:**

$$V_{\text{random}}(S)(s) = \frac{1}{|S|}, \quad \forall s \in S,$$

where each state $s \in S$ is selected uniformly at random. This gives an equal value to each state, so that one is selected randomly at each level of the tree.

(d) **Perfect state evaluator:**

$$V_{\text{perfect}}(S)(s) = \mathbb{1}[s = s^*],$$

where:

$$s^* = \arg \max_{s \in S} \text{BERTScore}(\text{complete}(s))$$

and $\text{complete}(s)$ denotes the best possible full solution reachable from state s .

This always chooses the state that will end up maximising the BERTScore. While defining this function is very hard (and using LLMs to approximate this is a focus of this thesis), by propagating each branch of the tree forward to completion and backtracking we can deduce what the ideal state evaluator would have chosen.

4. Search Algorithm

Finally, within the Tree of Thoughts (ToT) framework, one can use different search algorithms depending on the tree structure. The original paper mentions the depth first and breadth first algorithms as starting points. Given the simple tree structure here, one of the main constraints is the accuracy of the state evaluator $V(p_\theta, S)$. In particular, it is likely easier for the language model to evaluate the likelihood of a particular mental state directly using a subtext instead of a thought generated by that subtext, because it will potentially match what it has seen in its training distribution better.

Hence, the final algorithm we implement can be thought of a breadth first search over the first layer, including (context+subtext) with a greedy rollout of the most promising state to produce a thought based on the subtext and then the final model response as in Equations 3.2 and 3.3.

The algorithmic implementation of this for this tree structure is included below.

Algorithm 1 Greedy Rollout After Breadth-First Search

Require: Input x , LM p_θ , subtexts $T = \{\tau_0^{(i)}\}_{i=1}^N$, thought generator $G()$, state evaluator $V()$.

- 1: $S'_0 \leftarrow \{[x, \tau_0^{(i)}] \mid \tau_0^{(i)} \in T\}$
 - 2: $V_0 \leftarrow V(p_\theta, S'_0)$
 - 3: $s_0^* \leftarrow \arg \max_{s \in S'_0} V_0(s)$
 - 4: $s_1 \leftarrow G_1(p_\theta, s_0^*, 1)$
 - 5: **return** $G_2(p_\theta, s_1, 1)$
-

3.3 Ranking and State Evaluation Metrics

For being able to characterise the ranking of the thoughts as well as the consistency of the voting process, there are various metrics we will use.

Ranking thoughts

Average Reciprocal Rank

To be able to rank the thoughts we use the Average Reciprocal Rank (ARR) which is very similar to the Mean Reciprocal Rank used to evaluate the effectiveness of ranking systems such as search engines and recommendation systems. The main difference is that ARR is calculated per thought instead of per-query. It is given as:

$$ARR_i = \frac{1}{N} \sum_{k=1}^N \frac{1}{r_{ik}} \quad (3.4)$$

Where ARR_i is the Average Reciprocal Rank of thought i and r_{ik} is the rank of thought i in trial k . This creates a score that varies from worst at $1/N$ to a best score of 1. While doing a sum of ranks would treat 1st vs 2nd vs 3rd as linear differences, MRR or calculating the number of times a thought is ranked first would only consider the rank of the 1st ranked thought. Hence, ARR provides a middle ground where being 1st is prioritized, but having a good ranking such as always ranking second will still have a good score.

Measuring consistency in voting

When prompting the LLM with the subtext thoughts in different orders, I observed that the voted thought depended on the order of thoughts in the prompt. Ideally, the selection mechanism would be perfectly consistent, however there must be some bias or uncertainty in the voting process that decreases this consistency. We can quantify the consistency of the voting process using two complementary measures:

1. Maximum Agreement Score

For each sample we define $V_{i,j} = \{v_{i,j}^{(1)}, v_{i,j}^{(2)}, \dots, v_{i,j}^{(L)}\}$ as the set of votes received by thought j in sample i across L shuffles, where: $v_{i,j}^{(\ell)} = \begin{cases} 1 & \text{if thought } j \text{ was selected in shuffle } \ell, \\ 0 & \text{otherwise.} \end{cases}$

Then the count matrix $C_{i,j}$ represents the total number of times thought j was selected for sample i across all shuffles:

$$C_{i,j} = \sum_{\ell=1}^L v_{i,j}^{(\ell)}$$

(Note that the number of shuffles can be recovered by calculating $L = \sum_{i=1}^k c_i$).

The maximum agreement score for sample i is then:

$$\max_j C_{i,j}.$$

This score ranges from T/L (minimal consistency) to L (perfect consistency). Higher values indicate stronger robustness to ordering effects.

2. Evenness Index

The evenness index provides a normalized measure of vote distribution uniformity for sample i . It ranges from 0 to 1, where a value of 1 indicates the distribution is uniform and highly inconsistent, while a value of 0 represents perfect consistency. It is defined as:

$$E_i = \frac{H_i}{\log_2(L)} = \frac{-\sum_{i=1}^k p_i \log_2 p_i}{\log_2(L)} \quad (3.5)$$

Where E_i is the evenness index for sample i , H_i is the Shannon entropy of the vote distribution at sample i and $p_j = \frac{C_{i,j}}{L}$ is the probability that thought j was selected at sample i . The denominator $\log_2(L)$ normalizes the entropy by the maximum possible entropy over L shuffles, providing a measure of how evenly the votes are distributed across the thoughts.

3.4 Majority Voting

Using the measure of consensus described in the previous section, we can create a quasi-mixture of experts voting approach by aggregating scores/choices over different order of subtexts and selecting the most popular vote.

Effectively, this creates a meta-state evaluator, that takes in the output of the state evaluator $V(p_\theta, S)(s)$ over L shuffles. The voted subtext thought for sample i using count matrix $C_{i,j}$ by seeing what the most popular thought is would be:

$$\arg \max_j C_{i,j}.$$

Chapter 4

Experimental Results

4.1 Experimental Setup

4.1.1 Dataset curation

The first step of the experimental setup is setting up the dataset. For a general dialogue dataset with a series of utterances between two speakers, the first step is to segment it into a collection of dialogues, where each dialogue $D = \{u_1, u_2, \dots, u_N\}$ is represented as a sequence of N utterances between two speakers. After each dialogue is created, we can extract a series of contexts $C = \{c_{w+1}, c_{w+2}, \dots, c_N\}$ which are used to generate model responses $R = \{r_{w+1}, r_{w+2}, \dots, r_N\}$ to compare to human reference responses $L = \{l_{w+1}, l_{w+2}, \dots, l_N\}$.

To construct the sets C and L , a sliding window of size w is applied over each dialogue. For each turn i from $w + 1$ to N , we collect the w utterances immediately preceding u_i as the context c_i , and define the corresponding label l_i as the ground-truth utterance u_i . Since at least w previous utterances are required to form a complete context, the model only generates responses for turns where $i > w$. This ensures that predictions are all conditioned on the same amount of prior context for fairness.

24/7 Conversations

To implement this setup for the "24/7 Conversations" dataset, a series of preprocessing steps are required. To give some context, the "24/7 Conversations" dataset contains 750 hrs of continuous electrocorticography (ECoG) data from epilepsy patients engaging in free daily life conversations alongside audio recordings using a microphone to capture all the speakers in the room. These recordings have already been preprocessed by members of the Hasson lab using an Automatic Speech Recognition system to generate a time-stamped transcript. For one of the patients, patient 798, we have this data segmented by sentence with speaker labels.

To prepare the dataset as a collection of dialogues \mathcal{D} , we first transform the transcript into a sequence of utterances. Multiple consecutive sentences spoken by the same speaker are merged into a single utterance. After constructing the collection of dialogues $\mathcal{D} = \{D_1, D_2, \dots, D_J\}$, we apply steps which aim to filter out interactions that lack meaningful turn-taking and exchange or are not useful for modeling conversational dynamics.

The first filtering step targets one-sided conversations, which are identified based on the ratio of speaker turns. Specifically, we exclude any dialogue in which a single speaker contributes more than 75% of the utterances, as these are less likely to have meaningful two-way exchanges. This step is visualized in Figure 4.1. In this dataset, we observe several instances where one participant—typically a doctor—dominates the conversation, often explaining something while the other speaker listens passively. These are some of the interactions we specifically seek to exclude.

Finally, we apply a length-based filter to the reference responses. Many utterances in naturalistic dialogue are extremely short—often backchanneling cues such as "mhm," "yeah," or "okay"—which are challenging to model and offer limited value for evaluating response generation. On the other hand, extremely long utterances introduce their own modeling difficulties. Therefore, in our analysis, we restrict at-




Figure 4.1: A histogram of the percentage of sentences spoken by the dominant speaker in each conversation

tention to responses containing between 5 and 25 words. Then, for patient 798, the final dataset consists of 827 labels across 318 dialogues, the average number of turns per dialogue being 13.3. These statistics correspond to the case where the maximum number of context turns used in the analysis is $w = 15$. A smaller value of w would yield more datapoints, as fewer turns would be excluded for lacking sufficient preceding context.

(The code for this filtering is included in Appendix A.3)

4.1.2 Large Language Model Configuration

When selecting a language model for this project, we considered various models with different trade-offs, such as ease of use, performance on specific benchmarks, and flexibility for future fine-tuning. After careful evaluation, we chose the Llama 3.x series of models [73], which offers several key advantages. Notably, at the time of development it was one of the most powerful models with publicly available weights, providing the flexibility for potential future fine-tuning. Additionally, Llama 3.x is

available in multiple sizes—1B, 7B, and 80B parameters—trained on similar datasets, providing both adaptability across different use cases and the ability to assess the impact of model size on performance.

In particular, we use the instruction-tuned versions of Llama 3, which are specifically fine-tuned to improve the quality of generations by understanding more complex requests and producing more relevant responses. Generally, models are fine-tuned for specific tasks such as question-answering, summarization, and coding. Without this fine-tuning, the model would only predict the next token based on statistical patterns in the training data, leading to less rich and context-aware dialogue. The instruction-tuned models also enable us to incorporate a ”system prompt,” which helps shape the style and content of the model’s responses.

The hyperparameters, such as temperature, that control the sampling behavior for generating the next token, are detailed in Appendix A.5.

4.2 One step utterance generation

This section will describe the results of the single step utterance generation method described in Section 3.2.1. The system prompt for this generation is given in Appendix A.1. The code for generation is given in Appendix A.4.

4.2.1 Sentence generation improves with model size and context length-Up to a Limit

Figure 4.2 presents the average BERTScore across all samples, varying by model size and context length (measured as the number of preceding utterances). Broadly, we observe that BERTScore improves with both increasing model size and longer context. However, the gains from additional context diminish beyond approximately 10 preceding turns. This could suggest that speakers typically do not reference dialogue

more than 10 turns prior as often, or that excessive context may dilute the model’s attention to relevant parts of the conversation.

With regard to model size, both the 8B and 70B parameter models outperform the 1B model as context increases, with the 70B model slightly edging out the 8B.

This score is computed by averaging the BERTScore across randomly shuffled candidate-reference label pairs, which approximates the baseline generation method used in the original BERTScore paper [69] (the primary difference being that the original paper compares 1M pairs).

Two baseline scores are included as dashed lines in Figure 4.2. The red line represents the average BERTScore when comparing each human label l_a to another randomly selected label l_b from the dataset. Computing the score over randomly shuffled candidate- reference pairs of labels in this way approximates the method used to create a baseline in the original BERTScore paper [69] (the main difference being that the original paper compares 1M pairs). The blue line shows the average BERTScore between human labels and random generations from an 8B model without context. These baselines offer reference points: for instance, we expect models to perform no worse than the random generation baseline of 0.638.

Interestingly, the 8B and 70B models only begin to surpass the pairwise random-label baseline after a context length of 5, while the 1B model never exceeds it.

To evaluate the effectiveness of in-context prompting methods in later sections, we adopt the BERTScore of the 8B model with 10 turns of context as our model baseline because this configuration yielded the highest observed score.




Figure 4.2: Average sentence similarity scores between model generated and human reference turns of conversation for differing amounts of context and different model sizes

4.3 Thought generation for theory of mind reasoning

This section will go over the results of the multi-step sentence generation based on the thought generator by evaluating all possible branches in the tree of thought. By characterizing the performance of each branch, we can characterize the problem of choosing the best thought in more detail. Overall, I find that the model response based on the best subtext outperforms the baseline, however randomly selecting a subtext to generate based off of will not improve the baseline, hence being able to find the best ranked thought is important.

4.3.1 The best generated responses based on subtext thoughts improve upon baseline

Figure 4.3 presents the distribution of scores for each ranked thought, with thoughts ordered from highest to lowest score within each sample. The box plot illustrates how the top-ranked, second-ranked, third-ranked, and subsequent thoughts perform across all samples. Notably, selecting the model response associated with the first or second highest-ranked subtext consistently outperforms the baseline. In contrast, the third-ranked thought performs comparably to the baseline on average, while lower-ranked thoughts generally underperform. These results highlight a clear opportunity to improve upon the baseline by selecting higher-ranked generations.




Figure 4.3: Box plot showing the BERTScores for all ranked ToM subtext generations (1st, 2nd, 3rd best, etc.) for each sample.




Figure 4.4: Box plot showing the BERTScore distribution for each different subtext over all samples

4.3.2 Selecting Responses Based on a Single Subtext Underperforms Baseline

While choosing a highly ranked generation can outperform the baseline, consistently selecting responses based on a single subtext tends to slightly underperform. This trend is evident in Figure 4.4, which presents the average BERTScore for model responses categorized by subtext. From a Theory of Mind perspective, this result is logical: applying the same inferred mental state across all interactions overlooks the nuances of individual interactions and lacks conversational specificity to the context. As a result, biasing the selection mechanism toward any single subtext would not necessarily lead to any improved performance.

4.3.3 Self-Reflective and Goal-Oriented Thoughts Rank Higher

We use the Average Reciprocal Rank (ARR), defined in Equation 3.4, to evaluate the relative ranking performance of different subtexts. As shown in Figure ??, there is some variation across subtexts, though none deviate drastically from the expected

ARR for a uniform distribution of ranks, given by $\frac{1}{M} \sum_{k=1}^M \frac{1}{k} = 0.408$. Notably, subtexts that prompt the model to think about *goals* tend to achieve higher ARR scores, while those focused on *decision-making* or *information processing* perform slightly worse. Across all subtexts, prompting the LLM to consider the *upcoming* speaker’s state of mind from their own perspective generally yields better rankings than prompting it to reflect on the *previous* speaker’s state.




Figure 4.5: Bar chart showing the Average Reciprocal Rank for each theory of mind subtext. The possible range is between 0 and 1, with a perfect score of 1 being achieved by having a rank of 1 over all samples

4.3.4 Sampling multiple times from a varied distribution can still lead to improvement in BERTScore

Having shown that a thought generation approach based on a set of subtexts can improve text similarity scores, we next explore how perturbing this method affects performance.

Figure 4.6 summarises this analysis. To produce this graph, we perturb several key components involved in dialogue generation—the context c , the list of subtext




Figure 4.6: BERTScore outcomes for different generation strategies with various perturbation conditions. Dark blue bars represent the average best score per sample, light blue bars the worst and thick black bars the score from randomly sampling a subtext each sample. The yellow bar shows the average score with a simple label. The black dashed line is the model baseline.

thoughts T , and the intermediate generated thought z . The only major variable that remains unchanged is the prompt.

To perturb the context c , we use the Random Generation (RandomGen) condition. As in Section 4.2, this involves removing the context from the LLM input and instead sampling random text generations. We sample $M = 6$ generations to match the multi-sample setup used when conditioning on different subtexts.

For the Random Subtext condition, we replace the theory-of-mind-based subtexts with a set of unrelated, casual thoughts that are not clearly tied to mental states or intentions. The set includes: “Hm. Where was I going?” “Mmm. I could have toast later.” “I need to remember to call my mom later.” “What’s the weather like today?” “I need to remember to pick up groceries later.” “Do fish ever get thirsty?”

In the Random Thought condition, we retain the original theory-of-mind subtexts, but instead of generating thoughts conditioned on the subtext and context, we substitute the intermediate thought z with a completely random text sample.

We also include a simple baseline label—“Yeah, that’s what I see too.”—as a final comparison. Since this case does not involve multiple generations, it is plotted as a single yellow bar in Figure 4.6.

The figure shows the following: the light blue bars represent the average worst score across samples for each condition, while the dark blue bars show the average best score. The thick black horizontal lines indicates the expected score if a sample were selected at random. Lastly, the dashed black line represents the baseline model response from Section 4.2.

First, note that sampling multiple times from a distribution with some variance and choosing the best outcome will, by definition, outperform the mean of the distribution. This section evaluates that intuition in practice by modifying the generation process to inject variety.

The results confirm that all generation strategies exhibit variance in output quality, meaning that sampling multiple times and choosing the best result consistently improves BERTScore. Interestingly, both the theory-of-mind and random subtext conditions yield similar distributions, with the random subtext condition exhibiting slightly more variance—resulting in a higher best-case score.

Surprisingly, the Random Thought condition outperforms both subtext-based strategies. One plausible explanation emerges from examining the performance of the Simple Label condition, which closely matches the model baseline despite its simplicity. This suggests that basic or generic statements such as “Yeah, that’s what I see too” can score well in similarity metrics. By introducing randomness in thought generation, the model may ”get distracted” and fall back to more simple, generic responses. These responses would cluster near the average score, thereby lowering the worst-case performance but potentially boosting the maximum due to the occasional strong match.

Given these findings, for this strategy to have any use, it is crucial to be able to




Figure 4.7: Bar chart showing the Average Reciprocal Rank for each random idea subtext. The possible range is between 0 and 1, with a perfect score of 1 being achieved by having a rank of 1 over all samples

interpret performance changes in relation to underlying phenomena, rather than the simple effect of sampling multiple times from a noisy distribution. In this context, further analysis is needed to determine whether improvements in generation quality can be meaningfully linked to real-world mental states.

4.3.5 Using random ideas in subtext does not have clear trends in scores

Given our perturbation analysis, we can examine the random subtext and theory of mind subtext conditions in more depth for interpretability.

Looking at Figure 4.7 and comparing it to 4.5, we can see that the random subtext condition has no clear trend in performance based on the random subtext ideas. There is also more deviation from the uniform ranking ARR value.

Additionally, when computing the average scores across all samples in Figure 4.8, the results are comparable to those for the theory of mind subtexts shown in




Figure 4.8: Box plot showing the BERTScore distribution for each different subtext over all samples

Figure 4.4. One likely reason for this is that the large amount of conversational context—specifically, the inclusion of 10 preceding turns—allowed the model to produce a sensible reply, even when the subtext and resulting thought were largely ignored. When inspecting the response often they were mainly related to the context and not directly to the subtext thought.

4.4 State Evaluator

This section will review all the selection mechanisms implemented. In summary, no selection mechanism was able to select subtexts based on the conversational context to improve the generation BERTScore.

4.4.1 Choosing the best ranked thought improves BERTScore

To be able to contextualize our LLM-based voting mechanisms, the results from the random and perfect state evaluators are included as a reference for comparison here. A scatter plot of the scores from both state evaluators against the baseline




Figure 4.9: Box plot showing the BERTScores for all ranked random idea subtexts (1st, 2nd, 3rd best, etc.) for each sample.

generation score is included in Figure 4.10. We can see that generating a model response based on the perfect state evaluator which chooses the subtext resulting in the best outcome has a BERTScore which is almost always better than baseline single-step generation. However, when selecting a generation based on a random subtext there is no improvement in performance. This is consistent with the data in Figure 4.3 which shows that either the first or the second ranked subtext must be used on average to improve on the baseline.

4.4.2 Both LLM-based state evaluators are not effective in selecting the best subtext

Figure 4.11 depicts the scatter plot of the scores from both LLM-based state evaluators against the baseline generation score. Visually, it seems that the voting mechanism has no noticeable change on the distribution of scores, with the distribution appearing much more similar to the random state evaluator compared to the perfect state evaluator in Figure 4.10.




Figure 4.10: Scatter plots of the BERTScore based on model response from the best ranked thought (left) and a randomly selected thought (right) against baseline response

The results in Figure 4.12 confirm these observations. All selection mechanisms apart from the perfect state evaluator perform marginally worse than the baseline. It can be seen that the perfect state evaluator also has a slightly less spread distribution.




Figure 4.11: Scatter plots of the BERTScore based on model response from the voting over all possible subtexts (left) or scoring each subtext individually (right) against the baseline response.




Figure 4.12: Violin Plot of the BERTScore for different State Evaluators. The dark blue line represents the mean of the distribution while the black lines represent the datapoints 1.5 times the interquartile range (IQR) from the first and third quartile.

4.5 Evaluation over multiple shuffles

The following results are obtained by applying the voting scoring method and voting choice method 10 times each with the order of subtexts to choose randomized each time.

4.5.1 Evaluation by Scoring and by Voting are biased towards certain numbers

One reason that the state evaluation mechanism is not working well could be because of bias, which can be an issue when using LLMs as evaluators [74][75]. An obvious way this appears is a preference for certain numbers.

In Figure 4.13 we can see that for example the LLM prefers not to choose the first item in the list regardless of which subtext it represents. Similarly, in Figure 4.14 we

can see that the scoring mechanism has a strong preference for the digits 2,4 and 8. With a couple other digits having frequencies one or multiple orders of magnitude less, since the graph is plotted using a logarithmic scale.




Figure 4.13: The frequency each numerical index in the list is chosen when the state evaluator votes over states




Figure 4.14: The frequency specific digits appear when the state evaluator scores each option. Calculated over 10 shuffles.

4.5.2 Evaluation by Scoring is more consistent than by Choice

We can also quantify the consistency of the voting mechanism by using the metrics defined in Section 3.3, evenness index, and maximum agreement score.

Figures 4.15 and 4.16 show the distribution of these metrics over all the samples. It is clear that both voting mechanisms have high levels of bias; for example, for most samples, the maximum agreement on the best subtext over the different shuffles did not reach a majority of more than 4. However, these plots do show that the evaluation method by scoring independently is more consistent than voting over all the possible options, since the average entropy is closer to 0 and the number of maximum agreements also increases.




Figure 4.15: The continuous density plot of evenness index of all samples for different shuffles. An evenness index of 0 would indicate perfect consistency, while an index of 1 would indicate a perfectly random/uniform evaluation process.

4.6 Majority Voting

Since we implemented a shuffling mechanism, we can also create a meta-state evaluator by selecting the most popular subtext after a certain amount of shuffles.




Figure 4.16: The histogram of maximum number of agreements on the best subtext over different shuffles for all samples. A score of 10 would indicate perfect consistency while a uniform/random evaluator would score 1.67 on average.

4.6.1 Majority voting does not improve performance

In the scatter plots in Figures 4.17 and 4.18 we can see that sampling the state evaluator multiple times does not have a noticeable improvement on the distribution of scores.

4.6.2 High consistency on evenness index leads to lower variation of BERTScore of output

Figure 4.19 shows the distribution of scores for the majority vote meta-state evaluators as well as the distributions of the samples evaluated with high consistency (which are also highlighted in the scatter plots in Figures 4.18 and 4.17). While neither the majority vote meta-State Evaluators or the high consistency examples have a greater similarity to the human reference compared to the baseline, the evenness index metric of the scores or voting is indicative of the variance of the final BERTScore distribution. That is, the distribution of high consistency examples plotted as judged by evenness




Figure 4.17: A scatter plot of the BERTScore of the majority vote meta-state evaluator voting by voting and by scoring. General points are in grey while coloured points refer to high consistency examples with number of agreements ≥ 3 .




Figure 4.18: A scatter plot of the BERTScore of the majority vote meta-state evaluator voting by voting and by scoring. General points are in grey while coloured points refer to high consistency examples with evenness index 0.55.

index has a lower variation of BERTScore compared to all other conditions.




Figure 4.19: Violin Plot of the BERTScore for the majority vote meta-State Evaluators based on scoring and voting. Also plotted are the distributions of high consistency examples highlighted in the scatter plots in Figures 4.18 and 4.17

4.7 Overall findings

In this project, we further demonstrated that large language models are capable of generating sensible and coherent dialogue. Our experiments showed that both model size and context length positively correlate with similarity to human-generated references.

We introduced a thought generator based on a theory-of-mind taxonomy, which produced diverse response sets—some outperforming and others underperforming the baseline. This motivated the implementation of a selector mechanism designed to choose the most promising subtext to guide response generation. To better understand the generator’s behavior, we conducted a perturbation analysis, exploring how changes to different components influenced the distribution of model outputs. While BERTScore appeared to be a reasonable metric for assessing textual similarity, further investigation is needed to determine whether it captures sufficient semantic nuance,

and whether the generated thoughts and responses genuinely align with the theory-of-mind taxonomy provided.

Finally, we developed three state evaluation methods to select the best subtext for response generation. However, none of these approaches outperformed the baseline model that generates responses in a single step. Interestingly, we observed inconsistency and bias in the voting and scoring evaluators, though the consistency of an evaluator across multiple data shuffles—measured using the evenness index—appeared to reflect the spread in BERTScores of the model outputs. This suggests that LLMs might have some capacity to internally assess or express confidence in their generated responses.

Chapter 5

Discussion

While the ultimate goal of automatically improving dialogue generation using a theory of mind-based in-context learning approach was not fully achieved, this project still provides valuable insights—both for future improvements and for our broader understanding of conversational modeling.

To better understand the performance and shortcomings of the system, we can break down the implementation into its key components. Given the complexity and number of moving parts, this component-wise reflection is essential for identifying areas for refinement.

5.1 Future work

5.1.1 Language Model Capabilities

First, the capabilities of the language model itself, denoted p_θ , may have been a limiting factor. It is possible that the model’s social reasoning abilities are insufficient for accurate inference in this domain. This challenge is compounded by the inherent noisiness and ambiguity of social reasoning tasks. If the LLM has not been trained extensively on such data—likely, given the difficulty of curating high-quality theory-

of-mind datasets—it may lack the inductive biases required for these tasks. Although the weights of the LLaMA 3.x models are publicly available, their training data is not, making it difficult to verify this hypothesis. Potential next steps include fine-tuning the model on curated social reasoning datasets, such as SocialIQA [76], or designing a new benchmark of simpler, well-scaffolded theory-of-mind tasks to evaluate and adapt the model progressively.

5.1.2 Context Representation

Second, the context c provided to the model may not have been rich or structured enough. While it included recent conversational history, many other types of information are relevant for generating socially coherent dialogue: background world knowledge, participant personalities, shared memories, and conversational goals. A more structured or explicitly abstracted representation of context may better support theory-of-mind reasoning, aligning with thoughts in other work [26].

5.1.3 Evaluation Metrics.

Third, the evaluation metric—BERTScore—may lack the granularity necessary to capture nuanced social reasoning or alignment with human intent. Although BERTScore correlates with textual similarity, it does not assess propositional content and does not have extremely high correlations with human judgment in other tasks. Alternative metrics could include: using a pipeline to convert responses and labels to propositional form, to remove any reliance on surface representation of text before using a text similarity like BERTScore. Additionally, encoding accuracy from the sentence embedding of the generated text could be used as a score of how good a generation is. Then BERTScore as a metric could also be validated by seeing how well it correlated with encoding accuracy. Perplexity could also be used as a metric to capture likelihood of a response given context. Finally, human evaluations can be used to validate

the use of these automatic scores.

5.1.4 Thought Representation and Tree of Thought.

Another area for exploration is the structure of the thought process itself. It remains unclear whether the intermediate thoughts generated by the LLM accurately reflect the model’s internal reasoning [77, 78, 79, 80]. Doing more perturbations of thoughts could help show this.

An especially compelling dataset to apply this method to is ESConv [81], which features emotional support conversations between a ”speaker” describing their problems and a crowdworker ”listener” providing support. Each response from the listener is annotated with a communicative strategy, selected from a set that includes Questions, Self-disclosure, Affirmation and Reassurance, Providing Suggestions, Other, Reflection of Feelings, Information, and Restatement or Paraphrasing. This makes ESConv a rich resource for evaluating whether social reasoning capabilities—such as those enabled by a tree-of-thought approach—can be leveraged to predict the underlying social strategies used by the listener.

Furthermore, one potential direction is to simplify the thought process, for instance by reducing the tree depth from two to one, and conditioning the final response more directly on the intermediate reasoning steps. Prompt engineering—particularly for generating or interpreting theory-of-mind subtext—could also be refined.

5.1.5 Filtering to Target Turns that Require Reasoning

An important open question in conversational analysis is determining when participants engage in more automatic, reflexive responses versus when they employ deliberate reasoning. Our current predictions may capture turns that elicit automatic replies, which are inherently more challenging to model due to their low cognitive load and variability. Incorporating more sophisticated filtering methods could help isolate

conversational turns that genuinely require thought or reasoning. For instance, the high BERTScore of the generic label “Yeah I see that too” in Section 4.3.4 suggests that such low-effort responses may be disproportionately influencing our evaluation metrics, indicating a need for refined filtering criteria.

5.2 Limitations

First, the project assumes that conversations can be segmented into rigid, alternating turns. While this simplifies implementation, it does not reflect the fluid nature of real-world dialogue, which often includes interruptions, overlaps, or simultaneous speech. Moreover, it overlooks how interlocutors build shared context over time. The common ground theory [82] describes how shared knowledge accumulates and simplifies communication—for example, initial explanations may be detailed, but future references become more abbreviated. While this project began with segmenting dialogues into discrete two-person turns, future work could explore continuous, overlapping, or multi-party dialogue representations.

Second, the model is fundamentally unimodal—it operates on language alone. However, in human conversation, responses are influenced not only by the preceding words but also by a range of visual, situational, and embodied cues. Listeners naturally integrate information such as visual scene understanding [83], the goals and perspectives of the speaker [84], the physical context or affordances of objects in the environment [85], and paralinguistic signals such as facial expressions or body language [86, 87]. While some of these aspects may be indirectly encoded in language, much of this contextual richness remains implicit and inaccessible to language-only models. As such, future modeling efforts could benefit significantly from incorporating multimodal inputs or grounding the model in real-world context.

Chapter 6

Related Explorations-Neural Encoding using Sentence and Word Embeddings

This chapter presents the second project of the thesis: exploring how to perform sentence-level encoding of intracranial neural recordings. Unlike fMRI, intracranial recordings offer much higher temporal resolution—often far exceeding the duration of a single sentence. This mismatch in timescales makes sentence-level alignment particularly challenging.

Previous work on this dataset has primarily focused on encoding neural activity using token-level embeddings. However, incorporating sentence embeddings could provide a window into higher-order cognitive processes by capturing more abstract semantic representations. In this section, we investigate the feasibility of combining word and sentence embeddings to model neural activity using a banded ridge regression framework.

6.1 Data collection and pre-processing

The original dataset and pre-processing method is provided by Zada et al. (2025) [88]. To summarise, a 30 minute audio story (podcast) was played to nine participants while undergoing electrocorticographic monitoring for epilepsy. A simple preprocessing pipeline was implemented to extract the high-gamma band power per electrode. Alongside the neural recording, the story was also manually transcribed with each word timestamped at high temporal resolution.

The stimulus presented to participants was a segment from the podcast “This American Life” entitled “So a Monkey and a Horse Walk Into a Bar: Act One, Monkey in the Middle” released on November 10, 2017. The original audio and transcript are freely available online (<https://www.thisamericanlife.org/631/transcript>). We manually transcribed the story and timestamped words at high temporal resolution.

Raw electrode data underwent the following preprocessing pipeline, which is quoted from the Zada et al [88]. The steps include ”removing bad electrodes, downsampling, despiking and interpolating high amplitude spikes, common average re-referencing, and notch filtering. First, we visualized the power spectrum density of each electrode per subject. From this, we were able to annotate unusual electrodes that did not conform to the expected $1/f$ pattern, had a consistent oscillatory pattern, or showed other unusual artifacts. We found 31 such electrodes and marked them as “bad” (identified in the accompanying metadata). The source of these artifacts may be due to several factors, including excessive noise, epileptic activity, no noise, or poor contact. For data acquired with a sampling rate greater than 512 Hz, we downsampled to 512 Hz to match the sampling rate across subjects. We then applied a despiking and interpolation procedure to remove time points that exceeded four quartiles above the median of the signal and refill it using pchip interpolation. For re-referencing, we subtracted the mean signal across all electrodes per subject from each of their

individual electrode time series. Finally, we used notch filters at 60, 120, 180, and 240 Hz to remove power line noise. This pipeline produces a “cleaned” version of the raw signal.”

6.2 Method: Ridge regression

Here, I will describe the original method in the paper for ridge regression to decode neural activity from token-level embeddings.

Let $X \in \mathbb{R}^{n \times p}$ be a feature matrix with n samples and p features, $y \in \mathbb{R}^n$ a target vector, and $\alpha > 0$ a fixed regularisation hyperparameter. Ridge regression defines the weight vector $b^* \in \mathbb{R}^p$ as:

$$b^* = \arg \min_b ||Xb - y||_2^2 + \alpha ||b||_2^2. \quad (6.1)$$

The simplest way of casting this problem would be to have n be the number of words and p be the number of electrodes, so that the target matrix Y has length of number of electrodes. Each entry in X would contain the word embedding for the word at each new word onset, while Y would contain the voltage signals recorded from the electrodes placed on the brain’s surface in Volts. ECoG electrodes measure the aggregate electrical activity of thousands to millions of neurons and the recorded will mostly reflect local field potentials (LFPs), which arise from the synchronous activity of many neurons. Additionally, if the word consists of more than one token, we average all of them to get the embedding per word.

However, this method is unideal because the electrode data (512Hz) is much higher frequency than the spoken words (2Hz). To solve this, we epoch the electrode data around the onset of each word, so that X now has shape (number of words, number of ECoG electrodes x number of lags) and Y has shape (number of ECoG electrodes x number of lags). The epoch period chosen is from -2 seconds to +2 seconds relative to

word onset. This is equivalent to fitting a separate regression model for each electrode and each epoch.

Before the model can be fit, there are two additional elements in the pipeline we define in Scikit-learn [89]. First, Y and the regressors in X were mean-centered using StandardScaler. Second, because the design matrix was wider than it is long, we used the kernel method to solve the ridge regression in its dual form [90]. Specifically, we used a linear kernel for each feature space separately before fitting the model.

The penalty term b is learned using an inner cross-validation setup within the training set using the himalaya python library [91]. Then an outer cross-validation loop is used to evaluate the encoding model with $k = 2$. We evaluated the performance of each encoding model on the held-out fold by correlating the model-predicted ECoG signal Y_{preds} with the actual ECoG signal Y_{test} in the test fold. Finally, we averaged the two correlations for each fold to obtain one correlation value denoting the encoding performance for each feature space, electrode, and lag.

6.3 Method: Banded Ridge regression

We use a banded ridge regression model to regress to get the ECoG signal in Y with two different feature spaces with their own regularisation parameters as follows.

$$b^* = \arg \min_b \left\| \sum_{i=1}^m X_i b_i - y \right\|_2^2 + \sum_{i=1}^m \alpha_i \|b_i\|_2^2 \quad (6.2)$$

The variables are similar to Equation 6.1, where X_i represents the feature matrix corresponding to the i -th group, b_i denotes the coefficient vector for the i -th group, α_i is the regularization parameter specific to the i -th group and m is the total number of feature groups.

In practice for constructing X , the shape of X_1 and X_2 are both (number of words,

number of ECoG electrodes x number of lags). Because a sentence contains multiple words, the sentence embedding will be the same for each word in that sentence.

The embeddings used are:

- **Static word embeddings:**

- Word2Vec embeddings (300-dimensional)

- **Contextual word embeddings:**

- GPT-2 (13 layers of embeddings, each word embedding is 768-dimensional)

- **Sentence embeddings:**

- SentenceTransformer intfloat/e5-mistral-7b-instruct (4096-dimensional)

6.4 Results

6.4.1 Encoding with context and sentence embeddings outperforms static embeddings, but not contextual embeddings

Figure 6.1 depicts the encoding performance of a banded ridge regression with joint contextual and static embeddings compared to a ridge regression with either only static or contextual embeddings. It is clear that the joint embedding model improves on the static embeddings however compared to the contextual embeddings the performance appears very similar.

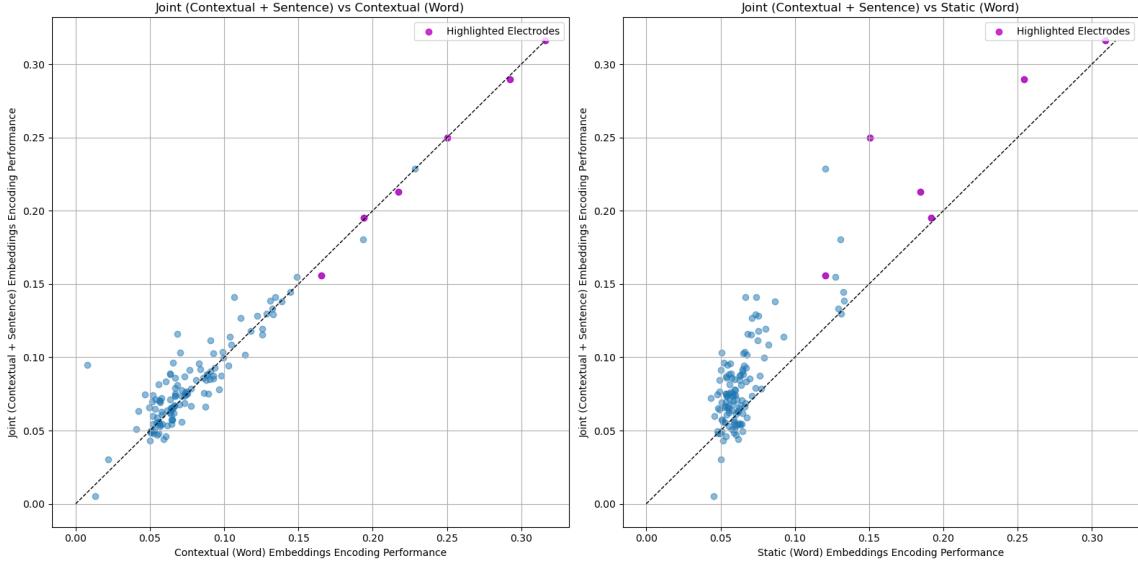


Figure 6.1: Scatter plots of encoding performance (correlation) for joint (sentence+contextual word) embeddings against static and contextual embeddings

6.4.2 Sentence Embeddings take up variance in Joint Context+Sentence encoding models

While using banded ridge regression with contextual and sentence embeddings does not improve upon the performance of ridge regression using only contextual embeddings, examining how the variance is split in the joint model can still be useful.

Figures 6.5 and 6.6 show that variance is split to a large extent between contextual and sentence embeddings. Figure 6.5 shows the encoding performance with either feature space separately, which shows that by themselves the overall encoding performance decreases. However, there are some electrodes where the sentence embeddings have a higher relative encoding performance and others where contextual embeddings have a much higher encoding performance on the right of the graph.

Plotting the encoding performance on the brain in Figure 6.6 we can see that the electrodes where the contextual embeddings do best are around XYZ brain areas. However, the correlation for sentence embeddings throughout the brain is generally more diffuse.

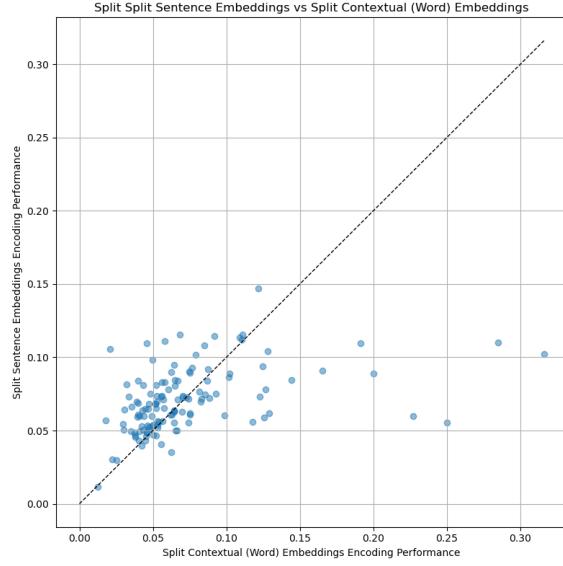


Figure 6.2: The encoding performance using either the contextual embedding or sentence embedding feature space plotted against each other

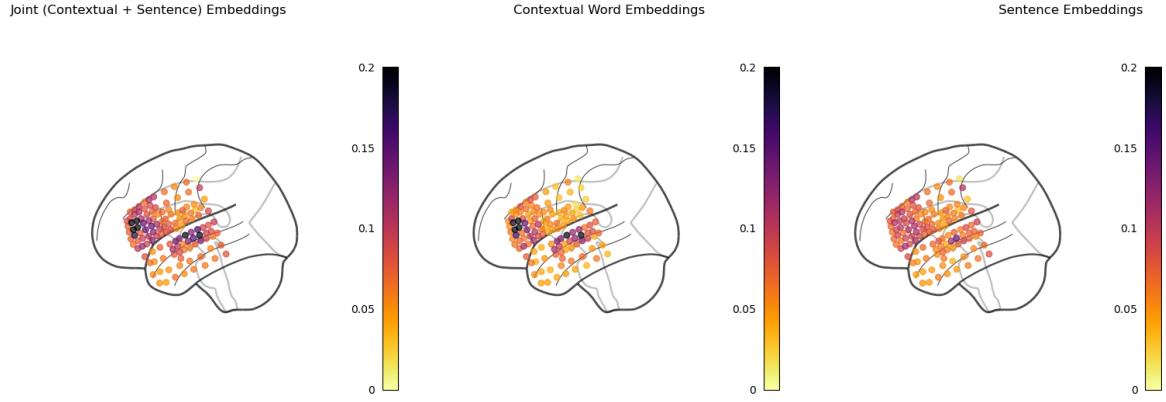


Figure 6.3: Correlation scores per electrode mapped onto the brain for joint (contextual word + sentence), contextual word and sentence embeddings

6.4.3 Encoding with static and sentence embeddings outperforms static embeddings, but not contextual embeddings”

In contrast to Figure 6.1 we compare the encoding performance of joint (static + sentence) embeddings compared to contextual and static embeddings in Figure 6.4

as opposed to joint (contextual+sentence) embeddings.

In Figure 6.4 we can see that the joint (static+sentence) model performs worse than model based on just contextual embeddings, however this joint model improves performance on models trained with just static embeddings. Most of this increase in performance happens in a range between a correlation of 0.05 to 0.09 for the static model, which increases to a range of 0.05 to 0.13. The performance for high correlation electrodes tends to be similar for both static and sentence embeddings.

6.4.




Figure 6.4: Scatter plots of encoding performance (correlation) for joint (sentence+static word) embeddings against static and contextual embeddings

6.4.4 Joint Sentence+static embeddings enable encoding for best electrodes over a greater period of time compared to only static embeddings

Another way of plotting our results is by collapsing the electrode dimension by selecting a subset and averaging, rather than collapsing the time dimension by choosing the maximum correlation over all time.

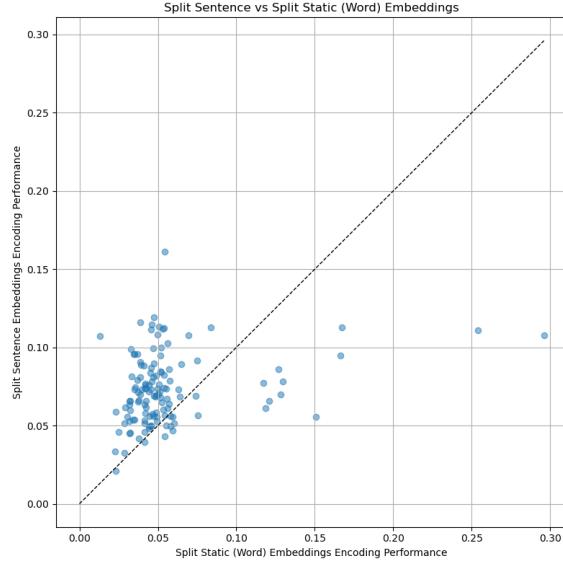


Figure 6.5: The encoding performance using either the contextual embedding or sentence embedding feature space plotted against each other




Figure 6.6: Correlation scores per electrode mapped onto the brain for joint (static word + sentence), static word and sentence embeddings

Figures 6.7 and 6.8 show the correlation over time for a subset of the best-performing electrodes. From Figure 6.7 is clear that encoding using sentence embedding in addition to static embeddings increases the encoding performance over a greater period of time. In Figure 6.8 we can see this effect as well, but it is less noticeable because the contextual embeddings already perform better over a greater period of time compared to the static embeddings. Nevertheless, by looking at the split embedding spaces we can see that the word embeddings generally peak a few

100ms after word onset, but the encoding accuracy using sentence embeddings is more spread out over time.

To be able to relate the size of the outer blue envelopes in Figures 6.7 and 6.8 we also plot them overlapping alongside the results from contextual embeddings in Figure 6.9(a). We can see that the outer envelopes of contextual, contextual+sentence and static+sentence have similar profiles. While in Figure 6.9(b) we can see that encoding using either static or split-static embeddings has a more narrow profile.

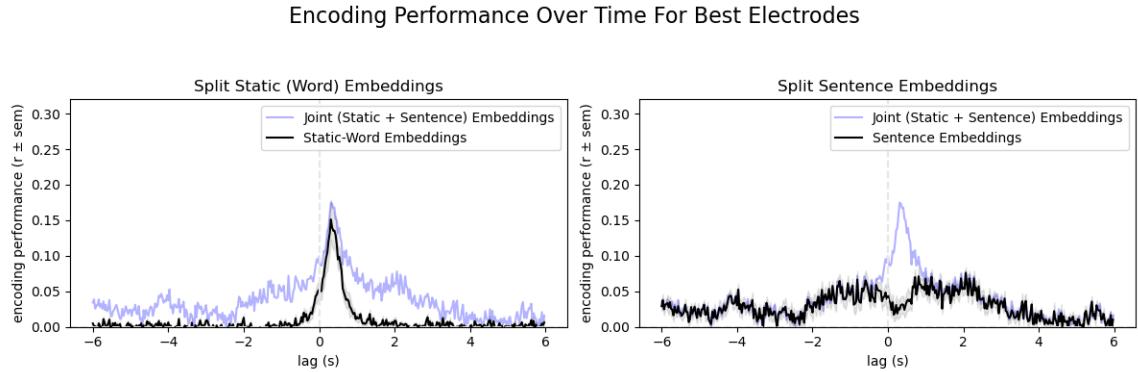


Figure 6.7: The correlation of model and actual neural signal plotted over time and averaged over best performing electrodes. Black lines show the encoding done using a split feature space after banded ridge regression for static embeddings (left) and sentence embeddings (right). The blue line represents the joint embedding performance. Electrodes where the joint (static+sentence) correlation had a performance of more than 0.15 were selected.

6.4.5 Overall findings

In this project, we successfully implemented a model for sentence encoding. We found that static and sentence embeddings when combined can get encoding performance that approaches that of contextual embeddings, with the added flexibility of being able to change the content of the sentence embedding easily compared to contextual word embeddings.

Encoding Performance Over Time For Best Electrodes

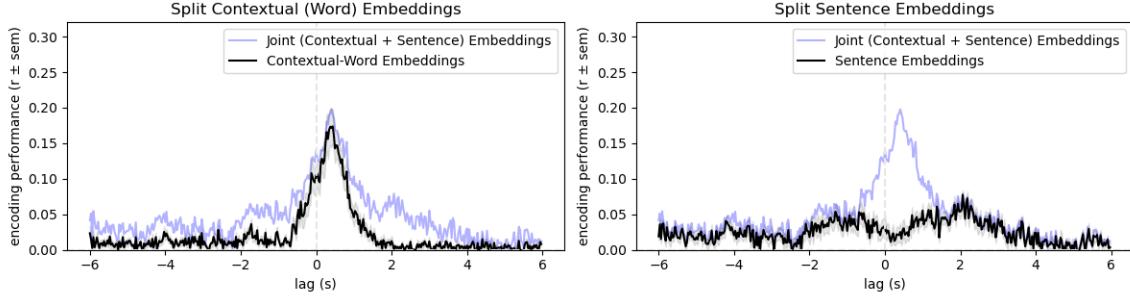


Figure 6.8: The correlation of model and actual neural signal plotted over time and averaged over best performing electrodes. Black lines show the encoding done using a split feature space after banded ridge regression for contextual embeddings (left) and sentence embeddings (right). The blue line represents the joint embedding performance. Electrodes where the joint (static+sentence) correlation had a performance of more than 0.15 were selected

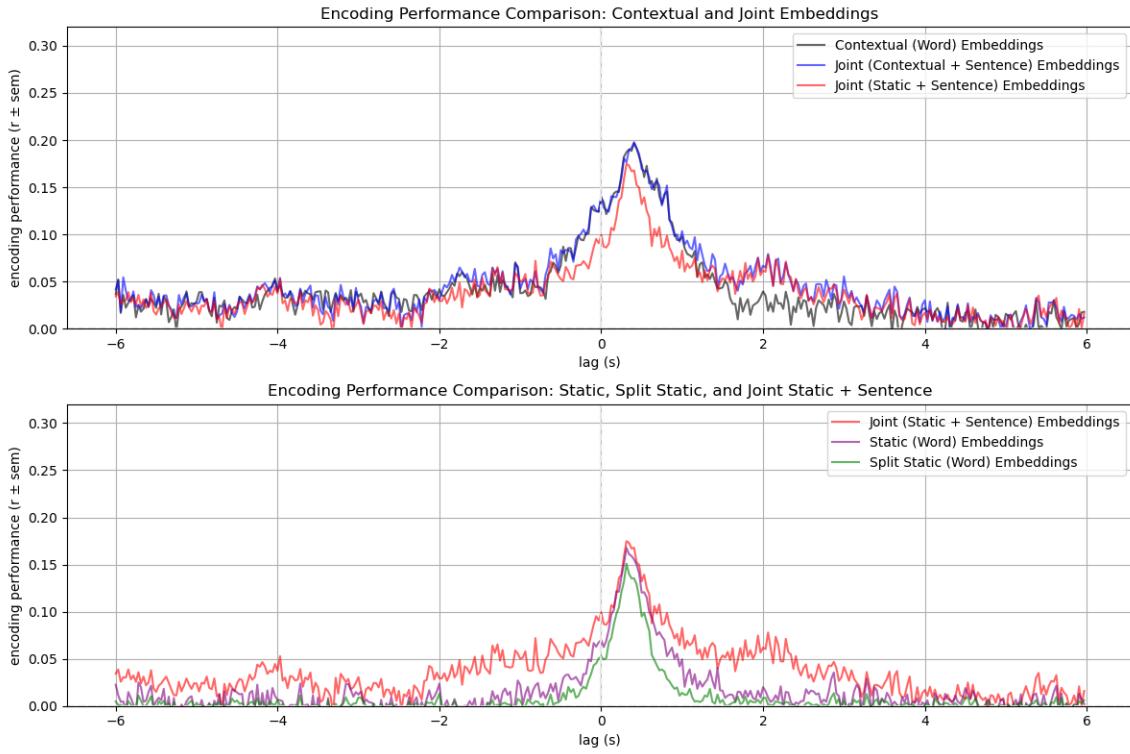


Figure 6.9: The correlation of model and actual neural signal plotted over time and averaged over best performing electrodes. Embeddings used in the model in order of envelop size are contextual, joint (contextual+sentence), joint (static+sentence) (plotted twice), static, and split static.

Chapter 7

Future Directions

Now that the two projects in this thesis have been established, it is worth exploring how they might be integrated to offer insight into the neural basis of theory of mind.

One direct way to bridge the conversational modeling project with neural data is by using the voted subtexts as labels for each utterance. These labels could serve as supervision signals in a neural classifier aimed at predicting the mental state most likely associated with each utterance. A crucial aspect of this task would involve distinguishing between moments when a person is engaged in mental state reasoning versus more routine conversational behavior—an open and intriguing question.

If the dialogue generation task proves challenging, an interesting alternative would be to adapt ideas from masked sentence modeling. Rather than predicting the final utterance in a conversation, the model could instead predict a missing utterance from the middle. This would allow the model to leverage both past and future context, potentially resulting better ability to reason over theory of mind.

In terms of evaluation, rather than relying solely on BERTScore, one could assess how strongly generated utterances drive activity in neural encoding models. After training on a subset of ground-truth data, such an approach could help identify "supernormal stimuli"—stimuli that elicit exaggerated responses from a given neural

system, as described by Barrett (2010) [92].

This project also lends itself naturally to an in-context learning framework, and its application need not be limited to social reasoning alone. Insights from psycholinguistics could enrich the modeling, particularly theories related to how people construct situation models through bridging inferences during discourse. Some accounts propose that literal and non-literal meanings are computed simultaneously within a parallel distributed framework—phenomena that could be explored by prompting large language models and examining their alignment with brain activity.

Another compelling area for exploration is inner speech. Inner speech serves diverse functions such as deliberation, planning, clarification, problem-solving, and even self-regulation or self-motivation (e.g., “You can do this,” “Don’t do that”) [93]. Modeling these internal dialogues might provide a novel pathway for connecting language-based AI models with neural data, offering a richer picture of cognition and self-directed thought.

Chapter 8

Conclusion

This thesis explored the extent to which Large Language Models can perform socially informed reasoning during dialogue, focusing on the inference of subtextual thoughts—unspoken beliefs, desires, and intentions—through a theory-of-mind-inspired approach. By applying the Tree of Thought framework to naturalistic dialogue generation, we demonstrated that LLMs are capable of generating plausible, socially coherent responses. However, our results also revealed that existing evaluation metrics and selection mechanisms struggle to consistently identify the most appropriate subtext or response, highlighting limitations in both model expressivity and current assessment tools.

In the second part of this thesis, we extended LLM applications to neural encoding, showing how both word- and sentence-level embeddings can be used to model brain activity during language comprehension. This builds on prior work in the Hasson lab and offers new directions for integrating computational representations of language with neuroscience.

Across both projects, we found that while LLMs exhibit emerging capacities for modeling literal and inferred content, their reasoning processes remain opaque and prone to inconsistency. Evaluator bias, limited context representation, and inade-

quate scoring methods present ongoing challenges. Nonetheless, the promising signals observed in perturbation and consistency analyses suggest that LLMs may possess latent capacities for internal assessment and social inference that are yet to be fully unlocked.

Future progress will likely depend on more structured representations of context, improved datasets for training and evaluation, and new methods for interpreting model reasoning. Ultimately, this work contributes to the growing effort to bridge the gap between surface-level language modeling and deeper social cognition in AI systems, with implications for both computational neuroscience and the development of more human-aligned dialogue agents.

Appendix A

Code Implementation

A.1 System prompt used

You will be given an excerpt of a conversation that includes the next speaker. It will be formatted as Speaker 1: message, Speaker 2: message. Reply as the next speaker in the conversation only. Format your reply as Speaker 1: message. Do not say you are a chatbot.

A.2 Prompt Generation

```
1 {
2
3 import json
4 import random
5
6 #final JSON file format:
7 #prompts_subtext: a list of length s (number of subtexts) containing
8     lists of length n (number of steps)
9 #subtexts: a list of length s containing the subtexts
```

```

9 #prompts_subtext_voting: a list of length 3 containing the prompts
  for the subtext generation (list of length s), voting (single
  string), and next speaker sentence generation (single string)

10
11
12
13 #number of subtexts=s, number of steps=n
14
15 def exportJSONFile(data,fileName):
16     json_object = json.dumps(data)
17
18     with open(fileName, "w") as outfile:
19         outfile.write(json_object)
20
21
22 subtextPrompt="You will be given an excerpt of a conversation
  between two speakers. Format your output as {Output:}. Your task
  is to generate the subtext thought that the next speaker is
  having, based on the following idea:"
23
24 #length s
25 subtextIdeaList=["Do I agree with this?", "What do they want from me?
  ", "How are they feeling right now?", "What are they avoiding
  mentioning?", "What's their body language saying?", "Should I
  clarify my earlier point?", "What do they really mean by this?"]
26 #subtextIdeaList2=["Why are they telling me this now?", "What
  triggered this topic in the conversation?", "How does this relate
  to what they said earlier?", "Does this connect to something
  they've shared about themselves before?", "How did they
  accomplish this, or how are they planning to?", "What resources or
  skills did they use to solve this problem?", "What are the
  implications of this for me or them?", "What are they not saying"

```

```

        that might be important?", "What are they trying to achieve by
        saying this?", "How does this align with their long-term goals or
        ambitions?"]

27 #subtextIdeaList3=["What are they about to say or ask?", "Are they
        feeling frustrated, happy, or anxious?", "Is this behavior typical
        for them, or is it out of character?", "Are they hinting at a
        deeper issue without explicitly saying it?", "Are they trying to
        persuade me?", "Their background might explain why they see things
        this way."]

28 #subtextIdeaList4=["How do I feel about this?", "How does the other
        person feel about this?", "What do I want to happen next?", "What
        do they want to happen next?", "What do I need to know to make a
        decision?", "What do they need to know to make a decision?"]

29

30

31 randomSubtextList=["Hm. Where was I going?", "Mmm. I could have toast
        later.", "I need to remember to call my mom later.", "Whats the
        weather like today?", "I need to remember to pick up groceries
        later.", "Do fish ever get thirsty?"]

32

33

34 subtextIdeaList_only4=["What do I want to happen next?", "What do
        they want to happen next?", "What do I need to know to make a
        decision?", "What do they need to know to make a decision?"]

35

36 prompt1=[]
37 prompt1Random=[]
38 selectedSubtexts=subtextIdeaList_only4
39 print(f"The number s of subtexts is {len(selectedSubtexts)}")
40 for i in range(len(selectedSubtexts)):
        #try this later (include info in the content) prompt1.append({
        "role": "system", "content": subtextPrompt + " " + s})

```

```

42     prompt1.append(subtextPrompt + " " + selectedSubtexts[i])
43     prompt1Random.append(subtextPrompt + " " + randomSubtextList[i])
44
45
46 #promptVoting="Based on the subtext generated, select the thought
   that you think is most likely to be the subtext of the next
   speaker. Output a single number corresponding to the thought you
   choose. Output this number as {Output:num}"
47 promptVoting="Based on the subtext generated and the conversation
   history provided, select the thought that is most likely to be
   the subtext of the next speaker. Output a single number
   corresponding to the thought you choose. Output this number as {
   Output:num}. "
48
49 promptVotingScoring="Based on the subtext generated and the
   conversation history provided, output a score corresponding to
   the likelihood that the subtext is the true subtext of the next
   speaker. Output this score as a number from 1 to 10 as {Output:
   num}. "
50
51
52 #or can include a newly formatted prompt that just contains relevant
   info (does not include the last prompt)
53 #prompt2A="You will be given an excerpt of a conversation that
   includes two speakers. It will be formatted as {Speaker 1:
   message, Speaker 2: message}, followed by {Thought:}. Reply as
   the next speaker in the conversation only based on the
   conversation history and the subtext thought. Format your reply
   as {Speaker X:message}. Do not say you are a chatbot. Do not
   output any further thoughts."
54 prompt2A="You will be given an excerpt of a conversation between two
   speakers. It will be formatted as {Speaker 1: message, Speaker

```

```

2: message}, followed by {Thought:}. Reply as the next speaker in
the conversation only based on the conversation history and the
subtext thought of the speaker replying. Format your reply as {
Speaker X:message}. Do not say you are a chatbot. Do not output
any further thoughts."
55
56 #systemPromptRigged="You will be given an excerpt of a conversation
between two speakers. Format your output as {Output:}. Your task
is to generate the subtext thought that the next speaker is
having, based on the following idea:"
57 promptVotingRiggedChoice="You are a cognitive scientist. Based on
the conversation history, which of these subtexts was the last
speaker most likely thinking before they spoke? Choose the one
that best fits the context. Output a single number corresponding
to the thought you choose. Output this number as {Output:num}. "
58 promptScoringRiggedScoring="You are a cognitive scientist. Based on
the conversation history and the given subtext, rate how likely
it is that this subtext reflects what the last speaker was
thinking before they spoke. Output a score from 1 (very unlikely)
to 10 (very likely) as {Output:num}"
59
60
61 all_subtexts_parallel_list=[]# a list of length s (number of
subtexts) containing lists of length n (number of steps)
62 all_subtexts_parallel_list=[prompt1,prompt2A]
63 all_subtexts_random_parallel_list=[prompt1Random,prompt2A]
64
65
66 #[[prompt1 (generate subtext)*10], prompt2 (voting), prompt3 (
generate next speaker sentence)]
67 prompts_subtext_voting1=[prompt1,promptVoting,prompt2A]#choice

```

```

68 prompts_subtext_voting2=[prompt1,promptVotingScoring,prompt2A]#
    scoring

69

70 prompts_subtext_voting3=[prompt1,promptVotingRiggedChoice,prompt2A]#
    voting_choice_rigged

71 prompts_subtext_voting4=[prompt1,promptScoringRiggedScoring,prompt2A
    ]#voting_scoring_rigged

72

73

74 myDict = {}

75 myDict["prompt_singleStep"] = "You will be given an excerpt of a
    conversation that includes the next speaker. It will be formatted
    as {Speaker 1: message, Speaker 2: message}. Reply as the next
    speaker in the conversation only. Format your reply as {Speaker
    1: message}. Do not say you are a chatbot."

76 myDict["prompts_subtext_multiStep"] = all_subtexts_parallel_list

77 myDict["prompts_subtext_multiStep_random"] =
    all_subtexts_random_parallel_list

78 myDict["subtexts"] = selectedSubtexts

79 myDict["prompts_subtext_voting1"]=prompts_subtext_voting1

80 myDict["prompts_subtext_voting2"]=prompts_subtext_voting2

81 myDict["randomSubtextList"] = randomSubtextList

82 myDict["prompts_subtext_voting3"] = prompts_subtext_voting3

83 myDict["prompts_subtext_voting4"]=prompts_subtext_voting4

84

85

86 exportJSONFile(myDict,"prompts/promptsC_only4.json")

87

88 promptOld="You will be given an excerpt of a conversation from the
    point of view of one speaker. Your task is to reply as the other
    speaker in the conversation. Reply as a human in the conversation
    not a chatbot."

```

```

89
90
91
92 #instead we want prompts to be structured like this:
93 #[[prompt1 (generate subtext)*10], prompt2 (voting), prompt3 (
94     generate next speaker sentence)]
95 }
```

A.3 Dataset filtering and setup

```

1 import pandas as pd
2 import numpy as np
3 import json
4 import matplotlib.pyplot as plt
5 from datasetUtilities import createDataset, exportJSONFile
6 import csv
7
8 df = pd.read_csv('decoding_df_798_final_notasks.csv')
9 print(df.columns)
10
11
12 sentence=df ["sentence"]
13 corrupted=df ["corrupted"]
14 speaker=df ["speaker"]
15 conv_name=df ["conversation_name"]
16
17
18
19 onsets=df ["onsets"]
20 offsets=df ["offsets"]
21
22 gap_time = np.zeros(len(onsets)) # Convert gap_time to a NumPy
```

```

        array

23 for i in range(len(onsets)):

24     if i == 0:

25         gap_time[i] = 0 # or some default value

26     else:

27         gap_time[i] = onsets[i] - offsets[i - 1]

28

29 print(gap_time[0:5])
30 print(len(gap_time))

31 N=len(speaker)#length of dataset

32 switchConvo = np.zeros(N, dtype=object)

33

34

35 #segmenting data into "two speaker segments", and also segmenting
    out corrupted segments

36 #get first two unique speakers

37 lastTwoSpeakers , index=np.unique(speaker,return_index=True)
38 lastTwoSpeakers=lastTwoSpeakers [np.argsort(index)]
39 lastTwoSpeakers=list(lastTwoSpeakers) [:2]

40

41 print(lastTwoSpeakers)

42

43 for i in range(len(sentence)):

44     curSpeaker=speaker[i]

45     #keep track of last two speakers, add new segment if one of them
        changes and update lastTwoSpeakers

46     if (curSpeaker not in lastTwoSpeakers):

47         prevSpeaker=speaker[i-1]

48         lastTwoSpeakers.clear()

49         lastTwoSpeakers.append(prevSpeaker)

50         lastTwoSpeakers.append(curSpeaker)

51         switchConvo[i]=1

```

```

52     #print("switch speaker" + str(i))
53
54     #print(curSpeaker)
55
56     #segment out corrupted data
57     if (corrupted[i]==1):
58
59         switchConvo[i]=1
60
61         if (i != 0):
62
63             if (corrupted[i-1]==1):
64
65                 switchConvo[i]=1
66
67                 if (conv_name[i]!=conv_name[i-1]):#add new segment if
68                     conversation as labelled in data changes
69
70                     switchConvo[i]=1
71
72
73 # all the indicies where a new segment starts
74 start_indices = np.where(switchConvo == 1)[0]
75 start_indices =np.insert(start_indices ,0,0)
76
77
78 # Create masks for each segment
79 masks = []
80 for i, start in enumerate(start_indices):
81
82     # Determine the end of the segment (either next start index or
83     # end of array)
84
85     end = start_indices[i+1] if i+1 < len(start_indices) else len(
86         switchConvo)
87
88     # Create a mask for the segment
89
90     mask = np.zeros_like(switchConvo, dtype=bool)
91
92     mask[start:end] = True
93
94
95     #remove corrupted segments

```

```

81     if (corrupted[start]==False):
82         masks.append(mask)
83
84
85 #segmenting data into segments based on masks
86 sentenceSegments = [sentence[mask] for mask in masks] # masked
87             sentences
88 speakerSegments = [speaker[mask] for mask in masks] # masked speaker
89             labels
90 gap_timeSegments = [gap_time[mask] for mask in masks] # masked gap
91             times
92 overallConversation=[]
93 overallGapTimes=[]
94
95 #checking for corrupted segments
96 print("Checking for corrupted segments")
97 corruptedSegments = [corrupted[mask] for mask in masks]
98 for i in range(len(corruptedSegments)):
99     if (corruptedSegments[i].any()):
100         print("missed corrupted segment!!!")
101
102 #finding the percentage of the sentences spoken by the speaker that
103             talks the most in each conversation
104 percentageSpeak=[]
105 for i in range(len(speakerSegments)):
106     percentageSpeak.append(speakerSegments[i].value_counts(normalize
107             =True).max())
108
109 plt.figure(figsize=(10, 6))
110 plt.hist(percentageSpeak, bins=30, edgecolor='black')
111 plt.xlabel('Percentage Spoken')
112 plt.ylabel('Frequency')

```

```

108 plt.title('Histogram of percentageSpeak -percentage of sentences
109 spoken by the speaker that talks the most in each conversation')
110 plt.show()
111
112 #joining sentences from each speaker into conversational turns , and
113 #formatting data to input to createDataset function
114 for i in range(len(sentenceSegments)):
115     #if conversation dominated by one speaker skip it
116     if (speakerSegments[i].value_counts(normalize=True).max() >
117         0.75):
118         continue
119
120     sentenceSegment=list(sentenceSegments[i])
121     speakerSegment=list(speakerSegments[i])
122     gap_timeSegment=list(gap_timeSegments[i])
123
124     prevSpeaker=speakerSegment[0] #get initial speaker
125     curUtterance=""
126
127     conversationHistory=[]
128     gapTimeHistory=[]
129     for j in range(len(sentenceSegment)):
130         #if speaker changes , add utterance so far from prev speaker
131         #to conversation history
132         if (prevSpeaker!=speakerSegment[j]):
133             conversationHistory.append(curUtterance)
134             curUtterance=""
135
136             curGapTime=gap_timeSegment[j]
137             gapTimeHistory.append(curGapTime)
138
139             #add next sentence to current uttranace (add a space if not
140             #the first sentence)

```

```

135     if j !=0:
136         curUtterance=curUtterance+ " " + sentenceSegment[j]
137     else:
138         curUtterance=sentenceSegment[j]
139         gapTimeHistory.append(0)#add 0 gap time for first
140         sentence
141
142         #updating the previous speaker
143         prevSpeaker=speakerSegment[j]
144         if gapTimeHistory:
145             gapTimeHistory.pop() # Remove the last entry of
146             gapTimeHistory
147             overallConversation.append(conversationHistory)
148
149
150
151
152
153 print(len(overallGapTimes[0]))
154 print(len(overallConversation[0]))
155
156 #shows how many conversations were filtered out
157 print("Num conversations before:")
158 print(len(sentenceSegments))
159 print("Num conversations after:")
160 print(len(overallConversation))
161
162 #plotting number of conversational turns in each conversation
163 conversationalTurns=[]
164 for i in overallConversation:

```

```

165     conversationalTurns.append(len(i))

166

167

168

169 plt.figure(figsize=(10, 6))
170 plt.hist(conversationalTurns, bins=30, edgecolor='black')
171 plt.xlabel('Number of conversational turns')
172 plt.ylabel('Frequency')
173 plt.title('Histogram of number of conversational turns in each
174 conversation')
175 plt.show()
176 plt.savefig('conversationLengths.png')

177 print("Average number of conversational turns:", np.mean(
178     conversationalTurns))

179

180 #creating correctly formatted datasets
181 references=list()
182 labels=list()
183 references2=list()
184 labels2=list()

185

186 max_num_turns=15

187

188 references,labels=createDataset(overallConversation,1,max_num_turns
189 -1)
190 references2,labels2=createDataset(overallConversation,2,
max_num_turns-2)
191 references3,labels3=createDataset(overallConversation,3,
max_num_turns-3)
192 references4,labels4=createDataset(overallConversation,4,
max_num_turns-4)
193 references5,labels5=createDataset(overallConversation,5,
max_num_turns-5)

```

```

        max_num_turns-5)

192 references10,labels10=createDataset(overallConversation,10,
    max_num_turns-10)

193 references15,labels15=createDataset(overallConversation,15,
    max_num_turns-15)

194

195 referencesGapTimes15,labelsGapTimes15=createDataset(overallGapTimes
    ,15,max_num_turns-15)

196

197

198 references3_extra,labels3_extra=createDataset(overallConversation
    ,3,3-3)

199 referencesGapTimes3_extra,labelsGapTimes3_extra=createDataset(
    overallGapTimes,3,3-3)

200

201 print("Number of label sentences: " + str(len(labels)))

202 print("labels identical:")

203 print(labels10==labels15)

204 print(labels5==labels10)

205 print(labels3==labels5)

206 print(labels2==labels3)

207 print(labels==labels2)

208

209 exportJSONFile(references,labels,"lab_data_no_tasks_one_turn.json")

210 exportJSONFile(references2,labels2,"lab_data_no_tasks_two_turn.json"
    )

211 exportJSONFile(references3,labels3,"lab_data_no_tasks_three_turn.
    json")

212 exportJSONFile(references5,labels5,"lab_data_no_tasks_five_turn.json"
    )

213 exportJSONFile(references10,labels10,"lab_data_no_tasks_ten_turn.
    json")

```

```

214 exportJSONFile(references15,labels15,"lab_data_no_tasks_fifteen_turn
    .json")
215
216 exportJSONFile(referencesGapTimes15,labelsGapTimes15, "
    lab_data_no_tasks_gap_times_fifteen_turn.json")
217
218 exportJSONFile(references3_extra,labels3_extra, "
    lab_data_no_tasks_three_turn_extra.json")
219 exportJSONFile(referencesGapTimes3_extra,labelsGapTimes3_extra, "
    lab_data_no_tasks_gap_times_three_turn_extra.json")
220
221 print(len(labelsGapTimes3_extra))

```

```

1     import json
2
3 #formats the dataset for inference
4 #input is structured by conversation and then conversational turns
5 #turnCount is num of prev. conversational turns to include
6 #startPoint is used to make sure different datasets have the same
    references e.g. if startPoint=0, one turn would have more data
    than two turn (one sentence for each conversation)
7 def createDataset(dataset,turnCount,startPoint=0):
8     _prompts=list()
9     _references=list()
10
11     #for each conversation
12     for i in range (len(dataset)):
13         excerpt=dataset[i]
14         #for all sentence sequences we can use
15         if len(excerpt)<turnCount+1:
16             continue
17
18         for j in range (startPoint,len(excerpt)-turnCount):

```

```

19
20     prompt=""
21
22     #add current speaker to prompt string (assumes
23     #alternating speakers)
24
25     for k in range(0,turnCount):
26
27         speakerNum=k %2 + 1
28
29         prompt = prompt + "Speaker " + str(speakerNum) + ":" + str(excerpt[j+k])
30
31         if k!=turnCount-1:
32             prompt=prompt + " "
33
34     #use next utterance as the reference
35     reference=excerpt[j+turnCount]
36
37     _prompts.append(prompt)
38     _references.append(reference)
39
40     return _prompts, _references
41
42
43
44
45     def exportJSONFile(_prompts,_references,fileNamed):
46
47         myDict = {}
48
49         myDict["references"] = _prompts
50
51         myDict["labels"] = _references
52
53
54         json_object = json.dumps(myDict)
55
56
57         with open(fileName, "w") as outfile:
58
59             outfile.write(json_object)

```

A.4 Inference generation

```
1
2 import torch
3 import time, subprocess, argparse, json, re, math
4 import accelerate.utils
5 import numpy as np
6 import pandas as pd
7 from transformers import AutoModelForCausalLM, AutoTokenizer,
8     AutoModel
9
10 from evaluate import load
11 from pathlib import Path
12 from datetime import datetime
13 import os
14 import math
15 import random
16
17 from inference import generateInference
18 from distutils.util import strtobool
19
20
21 numSamplesTotal=200
22
23
24
25 def getSimilarityResults(_references,_labels, sys_prompt, model,
26     tokenizer, bertscore, modelType="Llama8B", bertscore_model="distilbert-base-uncased", _temperature=1.0):
27     #perform inference and get similarity results through BERTscore
```

```
28 #SET UP PROMPTS
29
30 #Other prompt: You will be given an excerpt of a conversation
31 from the point of view of one speaker. Your task is to reply as
32 the other speaker in the conversation. Reply as a human in the
33 conversation not a chatbot.
34
35 myinput=[]
36
37 for i in range (0,len(_references)):
38     myinput.append([{"role": "system", "content": sys_prompt}, {
39         "role": "user", "content": _references[i]}])
40
41 #getting next sentence and outputting similarity results
42
43 gen_text=generateInference(myinput, model, tokenizer, modelType,
44 -temperature=_temperature)
45
46 print("Generated text: with tempeature " + str(_temperature))
47
48 results = bertscore.compute(predictions=gen_text, references=
49 _labels, model_type=bertscore_model)
50
51
52 return results['f1'], gen_text
53
54
55 def getSimilarityResultsMultiStep(_references,_labels, sys_prompt,
56 model, tokenizer, bertscore, modelType="Llama8B", bertscore_model
57 ="distilbert-base-uncased"):
58
59     #perform inference and get similarity results through BERTscore
60
61
62     #SET UP PROMPTS
63
64     prompt1=sys_prompt[0]
65
66     prompt2=sys_prompt[1]
67
68
69     myinput=[]
70
71     for i in range (0,len(_references)):
72         myinput.append([{"role": "system", "content": prompt1}, {
73             "role": "user", "content": _references[i]}])
```

```

51     gen_text1=generateInference(myinput, model, tokenizer, modelType
52     )#subtext thought
53
54     gen_text1 = [re.sub("Thought:", "", i, 3) for idx, i in enumerate(
55         gen_text1)]
56
57
58
59     myinput=[]
60
61     for i in range(0, len(_references)):
62
63         myinput.append([{"role": "system", "content": prompt2}, {"role": "user", "content": _references[i] + "Thought: " +
64         gen_text1[i]}])
65
66
67     gen_text2=generateInference(myinput, model, tokenizer, modelType
68     )#final output (generated speaker sentence(s))
69
70     gen_text2 = [re.sub("Speaker [0-9]:" , "", i, 3) for idx, i in
71     enumerate(gen_text2)]
72
73
74     results = bertscore.compute(predictions=gen_text2, references=
75     _labels, model_type=bertscore_model)
76
77
78     return results['f1'], gen_text1, gen_text2
79
80
81
82 def getSimilarityResultsMultiStepRandomThought(_references, _labels,
83     sys_prompt, model, tokenizer, bertscore, modelType="Llama8B",
84     bertscore_model="distilbert-base-uncased", random_thought_list=
85     None):
86
87     #perform inference and get similarity results through BERTscore
88
89
90     #SET UP PROMPTS
91
92     prompt1=sys_prompt[0]
93
94     prompt2=sys_prompt[1]

```

```

73     myinput=[]
74
75     for i in range (0,len(_references)):
76         myinput.append([{"role": "system", "content": prompt1}, {"role": "user", "content": _references[i]}])
77
78     #gen_text1 = ["no actual thought"] * len(_references) # Create
79     # an empty array of the same length as _references
80
81     if random_thought_list is not None:
82
83         if len(random_thought_list) >= len(_references):
84
85             gen_text1 = random.sample(random_thought_list, k=len(
86             _references))
87
88         else:
89
90             gen_text1 = random.choices(random_thought_list, k=len(
91             _references))
92
93     myinput=[]
94
95     for i in range (0,len(_references)):
96
97         myinput.append([{"role": "system", "content": prompt2}, {"role": "user", "content": _references[i] + "Thought: " +
98             gen_text1[i] } ])
99
100
101     gen_text2=generateInference(myinput, model, tokenizer, modelType)
102     #final output (generated speaker sentence(s))
103
104     gen_text2 = [re.sub("Speaker [0-9]:" , "", i, 3) for idx, i in
105     enumerate(gen_text2)]
106
107
108     #gen_text2=gen_text1
109
110     results = bertscore.compute(predictions=gen_text2, references=
111     _labels, model_type=bertscore_model)
112
113
114     return results['f1'], gen_text1, gen_text2

```

```

96
97
98 def getSimilarityResultsVoteChoice(_references,_labels, sys_prompt,
99     subtexts, model, tokenizer, bertscore, modelType="Llama8B",
100    bertscore_model="distilbert-base-uncased"):
101
102    #perform inference and get similarity results through BERTscore
103    ##[[prompt1 (generate subtext)*10], prompt2 (voting), prompt3 (
104    generate next speaker sentence)]
105
106
107    #SET UP PROMPTS
108
109    prompt1=sys_prompt[0]  #list of subtext prompts
110    prompt2=sys_prompt[1] # voting prompt
111    prompt3=sys_prompt[2] # generate next speaker sentence
112
113
114    #generating subtext thoughts
115    #want to keep the references size, so generate along the other
116    #column and then reshape the array
117
118    gen_text_thoughts=[]
119
120    for i in range (0,len(prompt1)):
121
122        myinput=[]
123
124        for j in range (0,len(_references)):
125
126            myinput.append([{"role": "system", "content": prompt1[i]
127                }, {"role": "user", "content": _references[j]}])
128
129        gen_text1=generateInference(myinput, model, tokenizer,
130        modelType)#subtext thought
131
132        gen_text1 = [re.sub("Thought:", "",i,3) for idx, i in
133        enumerate(gen_text1)]
134
135        gen_text_thoughts.append(gen_text1)
136
137
138    # Transpose the list of lists
139
140    gen_text_thoughts = list(map(list, zip(*gen_text_thoughts)))
141
142    #output should be a list of lists (x is reference, y is different

```

```

        subtext_prompts)

120    print("Gen thoughts shape:")
121    print(len(gen_text_thoughts))
122    print(len(gen_text_thoughts[0]))
123    #voting
124
125    for i in range(len(gen_text_thoughts)):#this should be length
126        references
127            #create the list of sutexts to choose from
128            thoughtList=[]
129            thoughtPrompt = "\n".join([f"{idx+1}. {thought}" for idx,
130            thought in enumerate(subtexts)])
131
132            myinput=[]
133            for i in range (0,len(_references)):
134                myinput.append([{"role": "system", "content": prompt2},
135                {"role": "user", "content": "Conversation History: " +
136                _references[i] + "Thought options: " + thoughtPrompt } ])
137
138            gen_text2=generateInference(myinput, model, tokenizer, modelType
139            ,_temperature=0.0)#result of voting (should just be a single
140            number)
141
142            #need to select which one was voted for
143            voted_thought=[]
144            numErrors=0
145
146            for i in range (len(gen_text_thoughts)):
147                parsedString=None
148                match = re.search(r'Output:\s*(\d+)', gen_text2[i])
149                if match:
150                    parsedString = match.group(1)
151                else:

```

```

145         match = re.search(r'\d+', gen_text2[i])
146
147         if match:
148
149             parsedString = match.group(0)
150
151         else:
152
153             print("Error: no number found in the output: " +
154                 gen_text2[i])
155
156             parsedString = "1" # default to "1" if no match is
157             found
158
159             numErrors+=1
160
161
162
163             if int(parsedString) > len(gen_text_thoughts[i]):
164
165                 print("Error: number found is greater than the number of
166 thoughts. The gen_text was: " + gen_text2[i])
167
168                 parsedString = "1"
169
170                 if int(parsedString) < 1:
171
172                     print("Error: number found is less than 1, which is not
173 a valid option. The gen_text was: " + gen_text2[i])
174
175                     parsedString = "1"
176
177                     voted_thought.append(int(parsedString) - 1)
178
179
180
181             print("Number of errors in voting in this batch: " + str(
182                 numErrors))
183
184             #need to generate the next speaker sentence based on thought
185             like before
186
187
188             myinput=[]
189
190             for i in range (0,len(_references)):
191
192                 myinput.append([{"role": "system", "content": prompt3}, {""
193                 role": "user", "content": _references[i] + "Thought: " +
194                 gen_text_thoughts[i][voted_thought[i]] } ])
195
196
197             gen_text3=generateInference(myinput, model, tokenizer, modelType

```

```

) #final output (generated speaker sentence(s))

169     gen_text3 = [re.sub("Speaker [0-9]:" , "", i, 3) for idx, i in
170     enumerate(gen_text3)]
171
172     results = bertscore.compute(predictions=gen_text3, references=
173     _labels, model_type=bertscore_model)
174
175     return results['f1'], gen_text_thoughts, gen_text2, gen_text3,
176     voted_thought
177
178
179 def getSimilarityResultsVoteScoring(_references, _labels, sys_prompt,
180                                     subtexts, model, tokenizer, bertscore, modelType="Llama8B",
181                                     bertscore_model="distilbert-base-uncased"):
182
183     #perform inference and get similarity results through BERTscore
184     ## [[prompt1 (generate subtext)*10], prompt2 (voting), prompt3 (
185     generate next speaker sentence)]
186
187
188     #SET UP PROMPTS
189
190     prompt1=sys_prompt[0] #list of subtext prompts
191     prompt2=sys_prompt[1] # scoring prompt
192     prompt3=sys_prompt[2] # generate next speaker sentence
193
194
195     #generating subtext thoughts and then scoring them
196     #want to keep the references size, so generate along the other
197     #column and then reshape the array
198
199     gen_text_thoughts=[]
200
201     for i in range (0,len(prompt1)):
202
203         myinput=[]
204
205         for j in range (0,len(_references)):
206
207             myinput.append([{"role": "system", "content": prompt1[i]
208             }, {"role": "user", "content": _references[j]}])
209
210         gen_text1=generateInference(myinput, model, tokenizer,

```

```

modelType)#subtext thought

192     gen_text1 = [re.sub("Thought:", "", i, 3) for idx, i in
193         enumerate(gen_text1)]
194         gen_text_thoughts.append(gen_text1)
195         gen_text_thoughts = list(map(list, zip(*gen_text_thoughts))) # Transpose the list of lists, output should be a list of lists (x is reference, y is different subtext prompts)
196         print("Gen thoughts shape:")
197         print(len(gen_text_thoughts))
198         print(len(gen_text_thoughts[0]))
199
200
201         #voting - instead of inputting the whole thought options,
202         evaluate all the thought options separately with a probability/
203         score, parse the score and then choose the vote with the highest
204         score
205
206
207         #generating scores for each subtext
208         gen_text_scores = []
209         for i in range(0, len(prompt1)):
210             myinput = []
211             for j in range(0, len(_references)):
212                 myinput.append([{"role": "system", "content": prompt2},
213 {"role": "user", "content": "Conversation History: " +
214 _references[j] + "Subtext: " + subtexts[i]}])
215             gen_text2 = generateInference(myinput, model, tokenizer,
216 modelType, _temperature=0.0)
217             gen_text_scores.append(gen_text2)
218
219
220             gen_text_scores = list(map(list, zip(*gen_text_scores))) # Transpose the list of lists, output should be a list of lists (x is reference, y is different subtext prompts)
221             print("Gen scores shape:")

```

```

212     print(len(gen_text_scores))
213     print(len(gen_text_scores[0]))
214
215     #parsing the list of scores to get the best thought for each
216     #reference
217
218     all_subtext_scores=[] #contains all the scores for all references
219     for each_subtext
220
221     best_score_values=[]
222     best_scored_thoughts=[0]*len(gen_text_scores)
223     best_score_indexes=[]
224
225     numErrors=0
226     for i in range(len(gen_text_scores)):#this should be length
227         references
228             subtext_scores=[]
229
230             for j in range(len(gen_text_scores[i])):#this should be
231                 length subtexts
232
233                 parsedString=None
234
235                 match=None
236
237                 match = re.search(r'Output:\s*(\d+)', gen_text_scores[i]
238 [j])
239
240                 if match:
241
242                     parsedString = match.group(1)
243
244                 else:
245
246                     match = re.search(r'\d+', gen_text_scores[i][j])
247
248                     if match:
249
250                         parsedString = match.group(0)
251
252                     else:
253
254                         print("Error: no number found in the output: " +
255 gen_text_scores[i][j])
256
257                         parsedString = "1" # default to "1" if no match

```

```

    is found

238             numErrors+=1

239

240             subtext_scores.append(int(parsedString))

241

242

243             #get the max of the subtext scores

244             max_score=max(subtext_scores)

245             max_score_index=subtext_scores.index(max_score)

246

247             all_subtext_scores.append(subtext_scores)

248             best_score_values.append(max_score)

249             best_scored_thoughts[i]=gen_text_scores[i][max_score_index]

250             best_score_indexes.append(max_score_index)

251

252             print("Number of errors in parsing scores in this batch: " + str

253             (numErrors))

254

255             #need to generate the next speaker sentence on the thought that

256             #is evaluated to be the best one

257             myinput=[]

258             for i in range (0,len(_references)):

259                 myinput.append([{"role": "system", "content": prompt3}, {"role": "user", "content": "Conversation History: " + _references[i] + "Thought: " + best_scored_thoughts[i] }])

260             gen_text3=generateInference(myinput, model, tokenizer, modelType

261             )#final output (generated speaker sentence(s))

262             gen_text3 = [re.sub("Speaker [0-9]:" , "", i, 3) for idx, i in

263             enumerate(gen_text3)]

```

```

263     print(len(gen_text3))
264
265     results = bertscore.compute(predictions=gen_text3, references=
266     _labels, model_type=bertscore_model)
267
268     return results['f1'], gen_text_thoughts, gen_text_scores,
269     gen_text3, best_score_indexes, all_subtext_scores,
270     best_score_values
271
272
273 def getSimilarityResultsVoteChoiceRigged(_references, _labels,
274   sys_prompt, subtexts, model, tokenizer, bertscore, modelType="Llama8B",
275   bertscore_model="distilbert-base-uncased"):
276
277   #perform inference and get similarity results through BERTscore
278   ##[[prompt1 (generate subtext)*10], prompt2 (voting), prompt3 (
279   generate next speaker sentence)]
280
281
282   #SET UP PROMPTS
283
284   prompt1=sys_prompt[0]  #list of generation of thought prompts
285   inc. subtexts
286
287   prompt2=sys_prompt[1] # voting prompt (rigged)
288
289   prompt3=sys_prompt[2] # generate next speaker sentence
290
291
292   for i in range(len(_references)):#this should be length
293     references
294
295     #create the list of subtexts to choose from
296
297     thoughtList=[]
298
299     thoughtPrompt = "\n".join([f"{idx+1}. {thought}" for idx,
300     thought in enumerate(subtexts)])
301
302
303     myinput=[]
304
305     for i in range (0,len(_references)):
306
307       myinput.append([{"role": "system", "content": prompt2},
308       {"role": "user", "content": "Conversation History: " +

```

```

    _references[i] + "Last Speaker: " + _labels[i] + "Thought options
    : " + thoughtPrompt } ])
```

285

```

286     gen_text2=generateInference(myinput, model, tokenizer, modelType
287     ,_temperature=0.0)#result of voting (should just be a single
288     number)
```

287

```

288     #need to select which one was voted for
289     voted_thought=[]
290     numErrors=0
291     for i in range (len(_references)):
292         parsedString=None
293         match = re.search(r'Output:\s*(\d+)', gen_text2[i])
294         if match:
295             parsedString = match.group(1)
296         else:
297             match = re.search(r'\d+', gen_text2[i])
298             if match:
299                 parsedString = match.group(0)
300             else:
301                 print("Error: no number found in the output: " +
302                     gen_text2[i])
303                 parsedString = "1" # default to "1" if no match is
304                     found
305                     numErrors+=1
306
307                     if int(parsedString) > len(prompt1):
308                         print("Error: number found is greater than the number of
309 thoughts. The gen_text was: " + gen_text2[i])
310                         parsedString = "1"
311                         if int(parsedString) < 1:
312                             print("Error: number found is less than 1, which is not
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
559
560
561
562
563
564
565
566
567
568
569
569
570
571
572
573
574
575
576
577
578
579
579
580
581
582
583
584
585
586
587
588
589
589
590
591
592
593
594
595
596
597
598
599
599
600
601
602
603
604
605
606
607
608
609
609
610
611
612
613
614
615
616
617
618
619
619
620
621
622
623
624
625
626
627
628
629
629
630
631
632
633
634
635
636
637
638
639
639
640
641
642
643
644
645
646
647
648
649
649
650
651
652
653
654
655
656
657
658
659
659
660
661
662
663
664
665
666
667
668
669
669
670
671
672
673
674
675
676
677
678
679
679
680
681
682
683
684
685
686
687
688
689
689
690
691
692
693
694
695
696
697
698
699
699
700
701
702
703
704
705
706
707
708
709
709
710
711
712
713
714
715
716
717
718
719
719
720
721
722
723
724
725
726
727
728
729
729
730
731
732
733
734
735
736
737
738
739
739
740
741
742
743
744
745
746
747
748
749
749
750
751
752
753
754
755
756
757
758
759
759
760
761
762
763
764
765
766
767
768
769
769
770
771
772
773
774
775
776
777
778
779
779
780
781
782
783
784
785
786
787
787
788
789
789
790
791
792
793
794
795
796
797
798
799
799
800
801
802
803
804
805
806
807
808
809
809
810
811
812
813
814
815
816
817
818
819
819
820
821
822
823
824
825
826
827
828
829
829
830
831
832
833
834
835
836
837
838
839
839
840
841
842
843
844
845
846
847
848
849
849
850
851
852
853
854
855
856
857
858
859
859
860
861
862
863
864
865
866
867
868
869
869
870
871
872
873
874
875
876
877
878
879
879
880
881
882
883
884
885
886
887
888
889
889
890
891
892
893
894
895
896
897
898
899
899
900
901
902
903
904
905
906
907
908
909
909
910
911
912
913
914
915
916
917
918
919
919
920
921
922
923
924
925
926
927
928
929
929
930
931
932
933
934
935
936
937
938
939
939
940
941
942
943
944
945
946
947
948
949
949
950
951
952
953
954
955
956
957
958
959
959
960
961
962
963
964
965
966
967
968
969
969
970
971
972
973
974
975
976
977
978
979
979
980
981
982
983
984
985
986
987
988
989
989
990
991
992
993
994
995
996
997
998
999
999
1000
1000
1001
1002
1003
1004
1005
1006
1007
1008
1009
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1019
1020
1021
1022
1023
1024
1025
1026
1027
1028
1029
1029
1030
1031
1032
1033
1034
1035
1036
1037
1038
1039
1039
1040
1041
1042
1043
1044
1045
1046
1047
1048
1048
1049
1050
1051
1052
1053
1054
1055
1056
1057
1058
1059
1059
1060
1061
1062
1063
1064
1065
1066
1067
1068
1068
1069
1070
1071
1072
1073
1074
1075
1076
1077
1077
1078
1079
1080
1081
1082
1083
1084
1085
1086
1087
1087
1088
1089
1090
1091
1092
1093
1094
1095
1095
1096
1097
1098
1099
1099
1100
1101
1102
1103
1104
1105
1106
1107
1108
1109
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1128
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1158
1159
1160
1161
1162
1163
1164
1165
1166
1167
1167
1168
1169
1170
1171
1172
1173
1174
1175
1175
1176
1177
1178
1179
1180
1181
1182
1183
1183
1184
1185
1186
1187
1188
1189
1189
1190
1191
1192
1193
1194
1195
1195
1196
1197
1198
1199
1199
1200
1201
1202
1203
1204
1205
1206
1207
1208
1208
1209
1210
1211
1212
1213
1214
1215
1216
1216
1217
1218
1219
1220
1221
1222
1223
1224
1225
1226
1227
1228
1228
1229
1230
1231
1232
1233
1234
1235
1236
1237
1238
1238
1239
1240
1241
1242
1243
1244
1245
1246
1247
1247
1248
1249
1250
1251
1252
1253
1254
1255
1256
1257
1258
1258
1259
1260
1261
1262
1263
1264
1265
1266
1267
1267
1268
1269
1270
1271
1272
1273
1274
1275
1275
1276
1277
1278
1279
1280
1281
1282
1283
1283
1284
1285
1286
1287
1288
1289
1289
1290
1291
1292
1293
1294
1295
1295
1296
1297
1298
1299
1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1308
1309
1310
1311
1312
1313
1314
1315
1316
1317
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1347
1348
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1367
1368
1369
1370
1371
1372
1373
1374
1375
1375
1376
1377
1378
1379
1380
1381
1382
1383
1383
1384
1385
1386
1387
1388
1389
1389
1390
1391
1392
1393
1394
1395
1395
1396
1397
1398
1399
1399
1400
1401
1402
1403
1404
1405
1406
1407
1408
1408
1409
1410
1411
1412
1413
1414
1415
1416
1417
1417
1418
1419
1420
1421
1422
1423
1424
1425
1426
1427
1427
1428
1429
1430
1431
1432
1433
1434
1435
1436
1437
1437
1438
1439
1440
1441
1442
1443
1444
1445
1446
1447
1447
1448
1449
1450
1451
1452
1453
1454
1455
1456
1457
1458
1458
1459
1460
1461
1462
1463
1464
1465
1466
1467
1467
1468
1469
1470
1471
1472
1473
1474
1475
1475
1476
1477
1478
1479
1480
1481
1482
1483
1483
1484
1485
1486
1487
1488
1489
1489
1490
1491
1492
1493
1494
1495
1495
1496
1497
1498
1499
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1547
1548
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1558
1559
1560
1561
1562
1563
1564
1565
1566
1567
1567
1568
1569
1570
1571
1572
1573
1574
1575
1575
1576
1577
1578
1579
1580
1581
1582
1583
1583
1584
1585
1586
1587
1588
1589
1589
1590
1591
1592
1593
1594
1595
1595
1596
1597
1598
1599
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1617
1618
1619
1620
1621
1622
1623
1624
1625
1626
1627
1627
1628
1629
1630
1631
1632
1633
1634
1635
1636
1637
1637
1638
1639
1640
1641
1642
1643
1644
1645
1646
1647
1647
1648
1649
1650
1651
1652
1653
1654
1655
1656
1657
1658
1658
1659
1660
1661
1662
1663
1664
1665
1666
1667
1667
1668
1669
1670
1671
1672
1673
1674
1675
1675
1676
1677
1678
1679
1680
1681
1682
1683
1683
1684
1685
1686
1687
1688
1689
1689
1690
1691
1692
1693
1694
1695
1695
1696
1697
1698
1699
1699
1700
1701
1702
1703
1704
1705
1706
1707
1708
1708
1709
1710
1711
1712
1713
1714
1715
1716
1717
1717
1718
1719
1720
1721
1722
1723
1724
1725
1726
1727
1727
1728
1729
1730
1731
1732
1733
1734
1735
1736
1737
1737
1738
1739
1740
1741
1742
1743
1744
1745
1746
1747
1747
1748
1749
1750
1751
1752
1753
1754
1755
1756
1757
1758
1758
1759
1760
1761
1762
1763
1764
1765
1766
1767
1767
1768
1769
1770
1771
1772
1773
1774
1775
1775
1776
1777
1778
1779
1780
1781
1782
1783
1783
1784
1785
1786
1787
1788
1789
1789
1790
1791
1792
1793
1794
1795
1795
1796
1797
1798
1799
1799
1800
1801
1802
1803
1804
1805
1806
1807
1808
1808
1809
1810
1811
1812
1813
1814
1815
1816
1817
1817
1818
1819
1820
1821
1822
1823
1824
1825
1826
1827
1827
1828
1829
1830
1831
1832
1833
1834
1835
1836
1837
1837
1838
1839
1840
1841
1842
1843
1844
1845
1846
1847
1847
1848
1849
1850
1851
1852
1853
1854
1855
1856
1857
1858
1858
1859
1860
1861
1862
1863
1864
1865
1866
1867
1867
1868
1869
1870
1871
1872
1873
1874
1875
1875
1876
1877
1878
1879
1880
1881
1882
1883
1883
1884
1885
1886
1887
1888
1889
1889
1890
1891
1892
1893
1894
1895
1895
1896
1897
1898
1899
1899
1900
1901
1902
1903
1904
1905
1906
1907
1908
1908
1909
1910
1911
1912
1913
1914
1915
1916
1917
1917
1918
1919
1920
1921
1922
1923
1924
1925
1926
1927
1927
1928
1929
1930
1931
1932
1933
1934
1935
1936
1937
1937
1938
1939
1940
1941
1942
1943
1944
1945
1946
1947
1947
1948
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1958
1959
1960
1961
1962
1963
1964
1965
1966
1967
1967
1968
1969
1970
1971
1972
1973
1974
1975
1975
1976
1977
1978
1979
1980
1981
1982
1983
1983
1984
1985
1986
1987
1988
1989
1989
1990
1991
1992
1993
1994
1994
1995
1996
1997
1998
1999
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2058
2059
2060
2061
2062
2063
2064
2065
2066
2067
2067
2068
2069
2070
2071
2072
2073
2074
2075
2075
2076
2077
2078
2079
2080
2081
2082
2083
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2094
2095
2096
2097
2098
2099
2099
2100
2101
2102
2103
2104
2105
2106
2107
2108
2108
2109
2110
2111
2112
2113
2114
2115
2116
2117
2117
2118
2119
2120
2121
2122
2123
2124
2125
2126
2127
2127
2128
2129
2130
2131
2132
2133
2134
2135
2136
2137
2137
2138
2139
2140
2141
2142
2143
2144
2145
2146
2147
2147
2148
2149
2150
2151
2152
2153
2154
2155
2156
2157
2158
2158
2159
2160
2161
2162
2163
2164
2165
2166
2167
2167
2168
2169
2170
2171
2172
2173
2174
2175
2175
2176
2177
2178
2179
2180
2181
2182
2183
2183
2184
2185
2186
2187
2188
2189
2189
2190
2191
2192
2193
2194
2194
2195
2196
2197
2198
2199
2199
2200
2201
2202
2203
2204
2205
2206
2207
2208
2208
2209
2210
2211
2212
2213
2214
2215
2216
2217
2217
2218
2219
2220
2221
2222
2223
2224
2225
2226
2227
2227
2228
2229
2230
2231
2232
2233
2234
2235
2236
2237
2237
2238
2239
2240
2241
2242
2243
2244
2245
2246
2247
2247
2248
2249
2250
2251
2252
2253
2254
2255
2256
2257
2258
2258
2259
2260
2261
2262
2263
2264
2265
2266
2267
2267
2268
2269
2270
2271
2272
2273
2274
2275
2275
2276
2277
2278
2279
2280
2281
2282
2283
2283
2284
2285
2286
2287
2288
2289
2289
2290
2291
2292
2293
2294
2295
2295
2296
2297
2298
2299
2299
2300
2301
2302
2303
2304
2305
2306
2307
2308
2308
2309
2310
2311
2312
2313
2314
2315
2316
2317
2317
2318
2319
2320
2321
2322
2323
2324
2325
2326
2327
2327
2328
2329
2330
2331
2332
2333
2334
2335
2336
2337
2337
2338
2339
2340
2341
2342
2343
2344
2345
2346
2347
2347
2348
2349
2350
2351
2352
2353
2354
2355
2356
2357
2358
2358
2359
2360
2361
2362
2363
2364
2365
2366
2367
2367
2368
2369
2370
2371
2372
2373
2374
2375
2375
2376
2377
2378
2379
2380
2381
2382
2383
2383
2384
2385
2386
2387
2388
2389
2389
2390
2391
2392
2393
2394
2394
2395
2396
2397
2398
2399
2399
2400
2401
2402
2403
2404
2405
2406
2407
2408
2408
2409
2410
2411
2412
2413
2414
2415
2416
2417
2417
2418
2419
2420
2421
2422
2423
2424
2425
2426
2427
2427
2428
2429
2430
2431
2432
243
```

```

    a valid option. The gen_text was: " + gen_text2[i])

310     parsedString = "1"
311
312     voted_thought.append(int(parsedString) - 1)

313
314     print("Number of errors in voting in this batch: " + str(
315         numErrors))

316     #need to generate the next speaker sentence based on thought
317     # like before

318
319     # Generate thoughts based on voted_thought directly in batches
320     gen_text_thoughts = [[{"no thoughts": 1}]*len(sys_prompt[0]) for _ in
321     range(len(_references))]
322
323     batch_inputs = []
324
325     for i in range(len(_references)):
326
327         voted_thought_index = voted_thought[i]
328
329         batch_inputs.append([{"role": "system", "content": sys_prompt[0][voted_thought_index]},
330
331                         {"role": "user", "content": _references[i]}])
332
333
334     # Perform batch inference
335
336     batch_gen_texts = generateInference(batch_inputs, model,
337                                         tokenizer, modelType)
338
339
340     # Update gen_text_thoughts with the generated results
341
342     for i, gen_text in enumerate(batch_gen_texts):
343
344         #print(gen_text)
345
346         voted_thought_index = voted_thought[i]
347
348         gen_text_thoughts[i][voted_thought_index] = re.sub("Thought:",
349             "", gen_text, 3)
350
351     print("Generated thoughts based on voted_thought.")

```

```

334
335     myinput=[]
336
337     for i in range (0,len(_references)):
338         myinput.append([{"role": "system", "content": prompt3}, {"role": "user", "content": _references[i] + "Thought: " +
339         gen_text_thoughts[i][voted_thought[i]]}])
340
341
342     gen_text3=generateInference(myinput, model, tokenizer, modelType
343     )#final output (generated speaker sentence(s))
344     gen_text3 = [re.sub("Speaker [0-9]:" , "",i,3) for idx, i in
345     enumerate(gen_text3)]
346
347     results = bertscore.compute(predictions=gen_text3, references=
348     _labels, model_type=bertscore_model)
349
350
351     return results['f1'], gen_text_thoughts, gen_text2, gen_text3,
352     voted_thought
353
354
355
356 def getSimilarityResultsVoteScoringRigged(_references,_labels,
357     sys_prompt,subtexts, model, tokenizer, bertscore, modelType="Llama8B", bertscore_model="distilbert-base-uncased"):
358
359     #perform inference and get similarity results through BERTscore
360     ##[[prompt1 (generate subtext)*10], prompt2 (voting), prompt3 (generate next speaker sentence)]
361
362     #SET UP PROMPTS
363
364     prompt1=sys_prompt[0] #list of subtext prompts
365     prompt2=sys_prompt[1] # scoring prompt
366     prompt3=sys_prompt[2] # generate next speaker sentence

```

```

357     #voting - instead of inputting the whole thought options,
358     #evaluate all the thought options separately with a probability/
359     #score, parse the score and then choose the vote with the highest
360     #score
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377

```

`#voting - instead of inputting the whole thought options,
#evaluate all the thought options separately with a probability/
#score, parse the score and then choose the vote with the highest
#score

#generating scores for each subtext
gen_text_scores=[]
for i in range (0,len(prompt1)):
 myinput=[]
 for j in range (0,len(_references)):
 myinput.append([{"role": "system", "content": prompt2},
{"role": "user", "content": "Conversation History: " +
_references[j] + ". Last Speaker: " + _labels[j] + ". Subtext
Thought: " + subtexts[i] }])
 gen_text2=generateInference(myinput, model, tokenizer,
modelType,_temperature=0.0)
 gen_text_scores.append(gen_text2)

gen_text_scores = list(map(list, zip(*gen_text_scores))) #
Transpose the list of lists, output should be a list of lists (x
is reference, y is different subtext prompts)
print("Gen scores shape:")
print(len(gen_text_scores))
print(len(gen_text_scores[0]))

#parsing the list of scores to get the best thought for each
reference
all_subtext_scores=[]#contains all the scores for all references
for each subtext
best_score_values=[]
best_scored_thoughts=[0]*len(gen_text_scores)
best_score_indexes=[]`

```

378
379     numErrors=0
380
381     for i in range(len(gen_text_scores)):#this should be length
382         references
383
384         subtext_scores=[]
385
386         for j in range(len(gen_text_scores[i])):#this should be
387             length subtexts
388
389             parsedString=None
390
391             match=None
392
393             match = re.search(r'Output:\s*(\d+)', gen_text_scores[i]
394             ][j])
395
396             if match:
397
398                 parsedString = match.group(1)
399
400             else:
401
402                 match = re.search(r'\d+', gen_text_scores[i][j])
403
404                 if match:
405
406                     parsedString = match.group(0)
407
408                 else:
409
410                     print("Error: no number found in the output: " +
411             gen_text_scores[i][j])
412
413                     parsedString = "1" # default to "1" if no match
414             is found
415
416                     numErrors+=1
417
418
419                     subtext_scores.append(int(parsedString))
420
421
422                     #get the max of the subtext scores
423
424                     max_score=max(subtext_scores)
425
426                     max_score_index=subtext_scores.index(max_score)

```

```

405     all_subtext_scores.append(subtext_scores)
406
407     best_score_values.append(max_score)
408
409     best_scored_thoughts[i]=gen_text_scores[i][max_score_index]
410     best_score_indexes.append(max_score_index)
411
412
413     print("Number of errors in parsing scores in this batch: " + str
414     (numErrors))
415
416     # Generate thoughts based on voted_thought directly
417
418     # Generate thoughts based on voted_thought directly in batches
419     gen_text_thoughts = [["no thoughts"] * len(sys_prompt[0]) for _
420     in range(len(_references))]
421
422     batch_inputs = []
423
424     for i in range(len(_references)):
425
426         voted_thought_index = int(best_score_indexes[i]) # Ensure
427         it's an integer
428
429         batch_inputs.append([{"role": "system", "content": sys_prompt[0][voted_thought_index]},
430
431                         {"role": "user", "content": _references
432 [i]}])
433
434
435         # Perform batch inference
436
437         batch_gen_texts = generateInference(batch_inputs, model,
438         tokenizer, modelType)
439
440
441         # Update gen_text_thoughts with the generated results
442
443         for i, gen_text in enumerate(batch_gen_texts):
444
445             #print(gen_text)
446
447             voted_thought_index = best_score_indexes[i]
448
449             gen_text_thoughts[i][voted_thought_index] = re.sub("Thought:",
450             "", "", gen_text, 3)
451
452             print("Generated thoughts based on voted_thought.")

```

```
430
431
432
433     #need to generate the next speaker sentence on the thought that
434     #is evaluated to be the best one
435
436     myinput=[]
437
438     for i in range(0, len(_references)):
439         myinput.append([{"role": "system", "content": prompt3}, {"role": "user", "content": "Conversation History: " + _references[i] + "Thought: " + gen_text_thoughts[i][best_score_indexes[i]] }])
440
441
442     gen_text3=generateInference(myinput, model, tokenizer, modelType)
443
444     #final output (generated speaker sentence(s))
445
446     gen_text3 = [re.sub("Speaker [0-9]:" , "", i, 3) for idx, i in
447     enumerate(gen_text3)]
448
449     print(len(gen_text3))
450
451     results = bertscore.compute(predictions=gen_text3, references=_labels, model_type=bertscore_model)
452
453
454     return results['f1'], gen_text_thoughts, gen_text_scores,
455     gen_text3, best_score_indexes, all_subtext_scores,
456     best_score_values
457
458
459
460
461
462
463
464
465
466
467
468 #want this to be the same for all prompting types
469
470 def performAnalysisOnWholeData(references, labels, sys_prompt,
471     prompting_type, model_name, model, tokenizer, bertscore,
472     shuffled_indices, subtexts=None, _temperature=1.0,
473     random_thought_list=None):
```

```

450     batchSize=20
451     batchNum=math.ceil(len(references)/batchSize)
452
453
454     results_dict = {
455         "reference": references,
456         "label": labels,
457     }
458
459     if prompting_type != "multi_parallel":
460         results_dict["gen_speech"] = []
461         results_dict["bertscore"] = []
462
463     if prompting_type == "voting_score" or prompting_type == "voting_choice" or prompting_type == "voting_choice_rigged" or prompting_type == "voting_score_rigged":
464         results_dict["gens_thoughts"] = []
465         results_dict["gens_vote"] = []
466         results_dict["voted_thoughts_parsed"] = []
467         results_dict["voted_thoughts_true"] = []
468         results_dict["gen_for_voted_thought"] = []
469         results_dict["subtext_for_voted_thoughts"] = []
470
471     for i in range(batchNum):
472         start=batchSize*i
473         end = len(references) if i == batchNum - 1 else batchSize * (i + 1)
474
475         if prompting_type == "single":
476             batch_scores, batch_gens_speech = (
477                 getSimilarityResults(references[start:end], labels[start:end], sys_prompt, model, tokenizer, bertscore, model_name,
478                 _temperature=_temperature)

```

```

477
478     )
479
480     if prompting_type == "multi_parallel":
481
482         for p in range(len(sys_prompt[0])):
483
484             print(f"Running for subtext {p}")
485
486             batch_prompt=[sys_prompt[0][p],sys_prompt[1]]
487
488             if random_thought_list is None:
489
490                 batch_scores, batch_gens_thought,
491
492                 batch_gens_speech = (
493
494                     getSimilarityResultsMultiStep(references[
495                         start:end], labels[start:end], batch_prompt, model, tokenizer,
496                         bertscore, model_name)
497
498                 )
499
500             else:
501
502                 batch_scores, batch_gens_thought,
503
504                 batch_gens_speech = (
505
506                     getSimilarityResultsMultiStepRandomThought(
507                         references[start:end], labels[start:end], batch_prompt, model,
508                         tokenizer, bertscore, model_name, random_thought_list=
509                         random_thought_list)
510
511                 )
512
513                 results_dict.setdefault(f"subtext{p}", []).extend([
514                     subtexts[p]]*(end-start))#may need to expand this by batch size
515
516                 results_dict.setdefault(f"gen_thought{p}", []).extend(batch_gens_thought)
517
518                 results_dict.setdefault(f"gen_speech{p}", []).extend(
519                     batch_gens_speech)
520
521                 results_dict.setdefault(f"bertscore{p}", []).extend(
522                     batch_scores)
523
524
525             elif prompting_type == "voting_score":
526
527                 batch_scores, batch_gens_thought, batch_gens_vote,

```

```

batch_gens_speech, batch_voted_thought, batch_all_subtext_scores,
batch_best_score_values = (
    getSimilarityResultsVoteScoring(references[start:end],
                                    labels[start:end], sys_prompt, subtexts, model, tokenizer,
                                    bertscore, model_name)
)

500     elif prompting_type == "voting_choice":
501         batch_scores, batch_gens_thought, batch_gens_vote,
batch_gens_speech, batch_voted_thought = (
    getSimilarityResultsVoteChoice(references[start:end],
                                    labels[start:end], sys_prompt, model, tokenizer, bertscore,
model_name)
)

504     elif prompting_type == "voting_choice_rigged":#
505         batch_scores, batch_gens_thought, batch_gens_vote,
batch_gens_speech, batch_voted_thought = (
    getSimilarityResultsVoteChoiceRigged(references[
start:end], labels[start:end], sys_prompt, subtexts, model,
tokenizer, bertscore, model_name)
)

507     elif prompting_type == "voting_score_rigged":
508         batch_scores, batch_gens_thought, batch_gens_vote,
batch_gens_speech, batch_voted_thought, batch_all_subtext_scores,
batch_best_score_values = (
    getSimilarityResultsVoteScoringRigged(references[
start:end], labels[start:end], sys_prompt, subtexts, model,
tokenizer, bertscore, model_name)
)

511     )
512
513     if prompting_type == "voting_choice" or prompting_type == "voting_score" or prompting_type == "voting_choice_rigged" or
prompting_type == "voting_score_rigged":

```

```

514         voted_thoughts_strings = [batch_gens_thought[i][
515             voted_thought] for i, voted_thought in enumerate(
516                 batch_voted_thought)]
517
518         voted_thoughts_subtext = [subtexts[i] for i in
519             batch_voted_thought]
520
521
522         results_dict.setdefault("gens_thoughts", []).extend(
523             batch_gens_thought)
524
525         results_dict.setdefault("gens_vote", []).extend(
526             batch_gens_vote)#for scoring choice this is the list of generated
527             scores
528
529         if prompting_type=="voting_score" or prompting_type=="voting_score_rigged":
530
531             results_dict.setdefault("gens_vote_unshuffled", []).extend(
532                 unshuffleArray(batch_gens_vote, shuffled_indices))
533
534             results_dict.setdefault("voted_thoughts_parsed", []).extend(
535                 batch_voted_thought)
536
537
538             # Undo the shuffling for subtexts, gens_thoughts,
539             voted_thoughts
540
541             reverse_mapping = shuffled_indices #map from
542             order in shuffled array to order in original array
543
544             # Apply the reverse mapping to restore the original
545             order
546
547             test_array=[shuffled_indices for _ in range(len(
548                 batch_gens_thought))]
549
550             unshuffled_subtexts=unshuffleArray(test_array,
551                 shuffled_indices)
552
553             unshuffled_gens_thoughts=unshuffleArray(
554                 batch_gens_thought, shuffled_indices)
555
556             voted_thoughts_true = [reverse_mapping[
557                 batch_voted_thought[i]] for i in range(len(batch_voted_thought))]
```

```

530
531
532         results_dict.setdefault("voted_thoughts_true", []).extend(voted_thoughts_true)
533         results_dict.setdefault("gen_for_voted_thought", []).extend(voted_thoughts_strings)
534         results_dict.setdefault("subtext_for_voted_thoughts", []).extend(voted_thoughts_subtext)
535         results_dict.setdefault("gens_thoughts_unshuffled", []).extend(unshuffled_gens_thoughts)
536         results_dict.setdefault("unshuffled_subtexts_indices", []).extend(unshuffled_subtexts)
537         if prompting_type == "voting_score" or prompting_type == "voting_score_rigged":#the part that's unique to voting_scores
538             results_dict.setdefault("all_subtext_scores_unshuffled", []).extend(unshuffleArray(
539                 batch_all_subtext_scores, shuffled_indices))
540             results_dict.setdefault("best_score_values", []).extend(batch_best_score_values)
541             if prompting_type != "multi_parallel":
542                 results_dict["bertscore"].extend(batch_scores)
543                 results_dict["gen_speech"].extend(batch_gens_speech)
544
545             for key, value in results_dict.items():
546                 print(f"{key}: {len(value)}")
547
548             if "results" in results_dict:
549                 results = results_dict["results"]
550                 print("Mean result: " + str(np.mean(results)))
551                 print("Standard deviation results: " + str(np.std(results)))
552                 print("Number of samples: " + str(len(results)))

```

```

553     print("Std deviation of mean: " + str(np.std(results) / np.
554     sqrt(len(results))))
555
556
557
558 def unshuffleArray(gens_thoughts, shuffled_indices):
559     temp_array1=[]
560     for i in range(len(gens_thoughts)):
561         temp_array2=[]
562         for j in range(len(shuffled_indices)):
563             index_in_shuffled = shuffled_indices.index(j)
564             temp_array2.append(gens_thoughts[i][index_in_shuffled])
565         temp_array1.append(temp_array2)
566     gens_thoughts=temp_array1
567
568
569
570 def main():
571     parser = argparse.ArgumentParser(description="Run analysis with
572     specified model and data.")
573     parser.add_argument('--model_name', type=str, required=True,
574     help='Name of the LM to use')
575     parser.add_argument('--data_path', type=str, required=True, help
576     ='Path to the data file')
577     parser.add_argument('--prompt_path', type=str, required=True,
578     help='Path to the prompt file')
579     parser.add_argument('--num_shuffles', type=int, required=True,
580     help='Number of random shuffles to do (first set is always not
581     shuffled)')
582     parser.add_argument('--prompting_type', type=str, required=True,
583     choices=["single","multi_parallel","voting_choice","voting_score"]

```

```

    ", "voting_choice_rigged", "voting_score_rigged"], help='Prompting
    type: standard or voting')
577     parser.add_argument('--temperature', type=float, required=False,
578                         help='Temperature of generation sampling', default=1.0)
579     parser.add_argument('--random_subtext', type=str, required=False
580                         , help='Whether a set of random subtexts are used', default="
581                         False")
582
583     parser.add_argument('--random_thought', type=str, required=False
584                         , help='Whether a set of random thoughts are used', default="
585                         False")
586
587     args = parser.parse_args()
588
589     model_name = args.model_name
590     data_path = Path(args.data_path)
591     prompt_path = Path(args.prompt_path)
592     num_shuffles = args.num_shuffles
593     random_shuffle = num_shuffles > 1 #if num_shuffles is 1, don't
594     shuffle
595
596     prompting_type = args.prompting_type
597     _temperature = args.temperature
598     random_subtext = str2bool(args.random_subtext)
599     random_thought = str2bool(args.random_thought)

596     device = torch.device("cuda" if torch.cuda.is_available() else "
597                           cpu")
598
599     print(f"Using device: {device}")

```

```

600     #LOADING AND SETTING UP THE MODEL
601
602     #set up model id based on model type
603
604     if (model_name=="Llama1B"):
605
606         model_id = "meta-llama/Llama-3.2-1B-Instruct"
607
608         model = AutoModelForCausalLM.from_pretrained(model_id,
609             torch_dtype=torch.bfloat16, device_map="auto").to(device)
610
611     elif (model_name=="Llama8B"):
612
613         model_id = "meta-llama/Llama-3.1-8B-Instruct"
614
615         model = AutoModelForCausalLM.from_pretrained(model_id,
616             torch_dtype=torch.bfloat16, device_map="auto").to(device)
617
618     elif (model_name=="Llama70B"):
619
620         model_id ="hugging-quant/Meta-Llama-3.1-70B-Instruct-GPTQ-
INT4"
621
622         #model_id = "unsloth/Meta-Llama-3.1-70B-Instruct-bnb-4bit"
623
624         model = AutoModelForCausalLM.from_pretrained(model_id,
625             torch_dtype=torch.float16).to(device)
626
627         # model = AutoModelForCausalLM.from_pretrained(model_id,
628             device_map="auto", cache_dir='/scratch/gpfs/ij9216/cached_models
').to(device)
629
630         # model = AutoModelForCausalLM.from_pretrained(model_id,
631             torch_dtype=torch.float16, device_map="auto", cache_dir='/scratch
/gpfs/ij9216/cached_models')
632
633
634     print("TESTING AA1")
635
636     tokenizer = AutoTokenizer.from_pretrained(model_id, padding_side
= "left")
637
638     tokenizer.pad_token_id = tokenizer.eos_token_id
639
640
641     print("TESTING AA2")
642
643     #print model size
644
645     module_sizes = accelerate.utils.compute_module_sizes(model)

```

```

623     print(f"The model size is {module_sizes[''] * 1e-9} GB")
624
625     if torch.cuda.is_available():
626         gpu_id = 0 # or whichever GPU you're using
627         total_memory = torch.cuda.get_device_properties(gpu_id).
628         total_memory
629
630         print(f"Total GPU memory: {total_memory / 1024**3:.2f} GB")
631
632         with open(prompt_path, 'r') as file:
633             promptsDict = json.load(file)
634
635         with open(data_path, 'r') as file:
636             data = json.load(file)
637
638         #get the first X datapoints
639         numSamplesTotal = len(data["references"])
640
641
642         subtexts=None
643
644         if prompting_type=="single":
645             prompts=promptsDict["prompt_singleStep"]
646             subtexts=None
647
648         elif prompting_type=="multi_parallel":
649             if not random_subtext:
650                 prompts=promptsDict["prompts_subtext_multiStep"]
651                 subtexts=promptsDict["subtexts"]
652
653             else:
654                 prompts=promptsDict["prompts_subtext_random"]
655                 subtexts = promptsDict["randomSubtextList"]

```

```

654     elif prompting_type == "voting_choice":
655         prompts = promptsDict["prompts_subtext_voting1"]
656         subtexts = promptsDict["subtexts"]
657     elif prompting_type == "voting_score":
658         prompts = promptsDict["prompts_subtext_voting2"]
659         subtexts = promptsDict["subtexts"]
660     elif prompting_type == "voting_choice_rigged":
661         prompts = promptsDict["prompts_subtext_voting3"]
662         subtexts = promptsDict["subtexts"]
663     elif prompting_type == "voting_score_rigged":
664         prompts = promptsDict["prompts_subtext_voting4"]
665         subtexts = promptsDict["subtexts"]
666     if random_thought:
667         random_thought_list = pd.read_csv("data/
RandomGens20250409_Llama8B.csv")["GeneratedText"].tolist()
668     else:
669         random_thought_list = None
670
671     print(f"Using model: {model_name}")
672     print(f"Data path: {data_path}")
673     print(f"Prompt path: {prompt_path}")
674     print(f"Prompting type: {prompting_type}")
675
676     all_res_dfs = []
677
678     for shuffle_id in range(1, num_shuffles + 1):
679         original_indices = list(range(len(prompts[0])))
680         shuffled_indices = original_indices[:]
681         if random_shuffle:
682             if shuffle_id != 1: # don't shuffle the first time
683                 random.shuffle(shuffled_indices)
684             if prompting_type == "single":

```

```

685         pass
686
687     elif prompting_type == "multi_parallel":
688
689         prompts[0] = [prompts[0][i] for i in
690         shuffled_indices]
691
692     elif prompting_type == "voting_score":
693
694         prompts[0] = [prompts[0][i] for i in
695         shuffled_indices]
696
697     elif prompting_type == "voting_choice":
698
699         prompts[0] = [prompts[0][i] for i in
700         shuffled_indices]
701
702
703     elif prompting_type == "voting_choice_rigged":
704
705         prompts[0] = [prompts[0][i] for i in
706         shuffled_indices]
707
708
709     if subtexts is not None:
710
711         subtexts = [subtexts[i] for i in shuffled_indices]
712
713
714     print(prompts)
715
716
717     today_date = datetime.now().strftime("%Y%m%d")
718
719     start_time = datetime.now()
720
721
722     results_dict = performAnalysisOnWholeData(references, labels,
723
724         prompts, prompting_type, model_name, model, tokenizer, bertscore,
725
726         shuffled_indices, subtexts, _temperature=_temperature,
727
728         random_thought_list=random_thought_list)
729
730
731     res_df=pd.DataFrame.from_dict(results_dict)
732
733
734
735     #almost exactly the same, just need to get reference/label
736     from csv instead of json file like before

```

```

709         #and then also add another column which says if voted
710         subtext index matches real subtext index, (i.e. whether
711         generation was correct
712
713         #also want to do it with multi parallel realistically
714
715         #can probably handle almost everything within the
716         performAnalysisOnWholeData function, only the input needs to
717         change asw
718
719         res_df.insert(0, 'sample_id', range(1, numSamplesTotal + 1))
720         res_df.insert(1, 'shuffle_id', shuffle_id)
721
722         res_df['shuffle_indices'] = [shuffled_indices] *
723         numSamplesTotal
724         res_df['subtexts'] = [subtexts] * numSamplesTotal
725
726         all_res_dfs.append(res_df)
727
728         final_res_df = pd.concat(all_res_dfs, ignore_index=True)
729         print(final_res_df.head())
730
731         # Calculate the duration
732         duration = datetime.now() - start_time
733         print(f"Time taken: {duration}")
734
735         # save the df to a csv
736
737         # filename: file_path _ prompt_file _ prompt index _ model name
738         # _ date
739
740         res_f_name = f"noramlrand{today_date}_{data_path.stem}_{prompt_path.stem}_{model_name}_{prompting_type}_shuffle{num_shuffles}_temp{_temperature}_random{str(random_subtext)}"
```

```
733  
734     # Check and create the directory  
735     if not os.path.exists("results"):  
736         os.makedirs("results")  
737     final_res_df.to_csv(f"results/{res_f_name}")  
738  
739     print(f"Results saved to results/{res_f_name}")  
740  
741  
742 if __name__ == "__main__":  
743     main()
```

```
1
2     import torch
3
4 import time, subprocess, argparse, json, re, math
5 import accelerate.utils
6 from transformers import AutoModelForCausalLM, AutoTokenizer,
7     AutoModel
8 from evaluate import load
9
10 #default values from generation_config.json      "temperature": 0.6, "
11 #                                            "top_p": 0.9,
12
13 def generateInference(inputPrompts, model, tokenizer, modelType="Llama8B", debugTime=False, device='cuda', _temperature=0.6):
14
15     #tokenizing inputs
16
17     texts = tokenizer.apply_chat_template(inputPrompts,
18                                         add_generation_prompt=True, tokenize=False)
19
20     inputs = tokenizer(texts, padding=True, return_tensors="pt").to(
21         device)
```

```
16     temp_texts=tokenizer.batch_decode(inputs["input_ids"] ,  
17                                         skip_special_tokens=True) #to see how tokenizer structured input  
18  
19     if debugTime:  
20         #generate inferred next sentence  
21         print("Starting generation")  
22         start_time = time.time()  
23  
24     if (modelType=="Llama1B"):  
25         terminators = [  
26             tokenizer.eos_token_id ,  
27             tokenizer.convert_tokens_to_ids("<| eot_id|>")  
28         ]  
29         if _temperature==0.0:  
30             gen_tokens = model.generate(  
31                 **inputs ,  
32                 max_new_tokens=50 ,  
33                 pad_token_id=tokenizer.eos_token_id ,eos_token_id=  
34                 terminators ,  
35                 do_sample=False ,  
36             )  
37         else:  
38             gen_tokens = model.generate(  
39                 **inputs ,  
40                 max_new_tokens=50 ,  
41                 pad_token_id=tokenizer.eos_token_id ,eos_token_id=  
42                 terminators ,  
43                 temperature=_temperature ,  
44             )  
45     else:  
46         if _temperature==0.0:  
47             gen_tokens = model.generate(
```

```

45         **inputs ,
46
47         max_new_tokens=50 ,
48
49         pad_token_id=tokenizer.eos_token_id ,
50
51         do_sample=False ,
52
53     )
54
55     else:
56
57         gen_tokens = model.generate(
58
59             **inputs ,
60
61             max_new_tokens=50 ,
62
63             pad_token_id=tokenizer.eos_token_id ,
64
65             temperature=_temperature ,
66
67         )
68
69     if debugTime:
70
71         print(f"Generation finished, Time: {time.time()-start_time}")
72
73
74     #decode generated text and remove the prompt
75     gen_text = tokenizer.batch_decode(gen_tokens ,
76
77         skip_special_tokens=True)
78
79     gen_text = [i[len(temp_texts[idx]):] for idx, i in enumerate(
80
81         gen_text)]
82
83     return gen_text

```

```

1 import torch
2 import time, subprocess, argparse, json, re, math
3 import accelerate.utils
4 import numpy as np
5 import pandas as pd
6 from transformers import AutoModelForCausalLM, AutoTokenizer,
7     AutoModel
8 from evaluate import load
9 from pathlib import Path

```

```

9  from datetime import datetime
10 import os
11 import math
12 import random
13 from inference import generateInference
14 from distutils.util import strtobool
15
16
17
18 def generateRandomText(_references, sys_prompt, model, tokenizer,
19   modelType="Llama8B", bertscore_model="distilbert-base-uncased",
20   _temperature=1.0):
21
22   myinput=[]
23
24   for i in range (0,len(_references)):
25
26     myinput.append([{"role": "system", "content": sys_prompt}, {
27       "role": "user", "content": _references[i]]] )
28
29   #getting next sentence and outputting similarity results
30   gen_text=generateInference(myinput, model, tokenizer, modelType,
31   _temperature=_temperature)
32
33   print("Generated text: with tempeature " + str(_temperature))
34
35
36   return gen_text
37
38
39
40 def main():
41
42   parser = argparse.ArgumentParser(description="Run analysis with
43   specified model and data.")
44
45   parser.add_argument('--model_name', type=str, required=True,
46   help='Name of the LM to use')
47
48   args = parser.parse_args()
49
50
51   model_name = args.model_name

```

```

35     #set up device to GPU if available
36
37     device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
38
39     print(f"Using device: {device}")
40
41     #LOADING AND SETTING UP THE MODEL
42
43     #set up model id based on model type
44
45
46     if (model_name=="Llama1B"):
47
48         model_id = "meta-llama/Llama-3.2-1B-Instruct"
49
50         model = AutoModelForCausalLM.from_pretrained(model_id,
51             torch_dtype=torch.bfloat16, device_map="auto").to(device)
52
53     elif (model_name=="Llama8B"):
54
55         model_id = "meta-llama/Llama-3.1-8B-Instruct"
56
57         model = AutoModelForCausalLM.from_pretrained(model_id,
58             torch_dtype=torch.bfloat16, device_map="auto").to(device)
59
60     elif (model_name=="Llama70B"):
61
62         model_id = "hugging-quants/Meta-Llama-3.1-70B-Instruct-GPTQ-
63             INT4"
64
65
66         #model_id = "unsloth/Meta-Llama-3.1-70B-Instruct-bnb-4bit"
67
68         model = AutoModelForCausalLM.from_pretrained(model_id,
69             torch_dtype=torch.float16).to(device)
70
71
72         # model = AutoModelForCausalLM.from_pretrained(model_id,
73             device_map="auto", cache_dir='/scratch/gpfs/ij9216/cached_models
74             ').to(device)
75
76
77         # model = AutoModelForCausalLM.from_pretrained(model_id,
78             torch_dtype=torch.float16, device_map="auto", cache_dir='/scratch
79             /gpfs/ij9216/cached_models')
80
81
82     print("TESTING AA1")
83
84     tokenizer = AutoTokenizer.from_pretrained(model_id, padding_side
85         = "left")
86
87     tokenizer.pad_token_id = tokenizer.eos_token_id

```

```

57
58     print("TESTING AA2")
59
60     #print model size
61     module_sizes = accelerate.utils.compute_module_sizes(model)
62     print(f"The model size is {module_sizes[''] * 1e-9} GB")
63
64
65     if torch.cuda.is_available():
66         gpu_id = 0 # or whichever GPU you're using
67         total_memory = torch.cuda.get_device_properties(gpu_id).
68         total_memory
69
70         print(f"Total GPU memory: {total_memory / 1024**3:.2f} GB")
71
72
73     numGens=8000
74
75     batchSize=50
76
77     references=["Generate some random text" for _ in range(numGens)]
78
79
80     numBatches=math.ceil(len(references)/batchSize)
81
82
83     allGeneratedText=[]
84
85     for i in range(numBatches):
86
87         startIndex=i*batchSize
88
89         endIndex=min((i+1)*batchSize ,len(references))
90
91         batchReferences=references[startIndex:endIndex]
92
93         #generate random text
94
95         generatedText=generateRandomText(batchReferences , "Generate
96         some random text" , model , tokenizer , modelType="Llama8B",
97         bertscore_model="distilbert-base-uncased" , temperature=2.0)
98
99         allGeneratedText.extend(generatedText)
100
101
102     results_dict = {"GeneratedText": allGeneratedText}
103
104     final_res_df = pd.DataFrame.from_dict(results_dict)

```

```
86
87
88     today_date = datetime.now().strftime("%Y%m%d")
89
90
91     # save the df to a csv
92
93     # filename: file_path - prompt_file - prompt index - model name
94     # - date
95
96     res_f_name = f"RandomGens{today_date}_{model_name}.csv"
97
98
99     # Check and create the directory
100
101    if not os.path.exists("results"):
102
103        os.makedirs("results")
104
105    final_res_df.to_csv(f"results/{res_f_name}")
106
107
108    print(f"Results saved to results/{res_f_name}")
109
110
111
112
113 if __name__ == "__main__":
114
115     main()
```

A.5 Language model hyperparameters

Hyperparameters reterived from "generation_on_fq.json". It is the same for LLaMa3.11B, LLaMa3.28B.

```
1 {  
2     "bos_token_id": 128000,  
3     "do_sample": true,  
4     "eos_token_id": [  
5         128001,  
6         128008,  
7         128009  
8     ],
```

```

9     "temperature": 0.6,
10    "top_p": 0.9,
11    "transformers_version": "4.46.2"
12 }

```

A.6 Banded Ridge Regression Code

dataframe_setup.py

```

1 import pandas as pd
2
3
4 #load transcript and extract sentences
5 # Extract sentences/sentence indices and add them to the transcript
6 # DataFrame.
7
8 punctTranscript_path = "monkey_util-whisperx_transcript.csv" # Change
9 # as needed
10
11 df = pd.read_csv(punctTranscript_path)
12 df.dropna(subset=[ 'start' ], inplace=True)
13 df.sort_values("start", inplace=True)
14 df.word.tail(10)
15
16 df.index = list(range(len(df.word)))# fix word index numbering
17 sentenceEnds = df.index[df.word.str.contains('[.?!]')]
18 df.loc[sentenceEnds]
19 sentenceBegs = sentenceEnds + 1
20 sentenceBegs = sentenceBegs.insert(0, 0)
21 sentenceBegs = sentenceBegs.delete(-1)
22
23 sentenceIdx = [0] * len(df.word);
24 sentence = ['] * len(df.word);

```

```

22 for i, j, k in zip(sentenceBegs, sentenceEnds, range(len(
23     sentenceBegs))):
24     fullSentence = (df.word[i:(j+1)]).str.cat(sep=' ')#concatenate
25     words in sentence
26
27
28 sentenceIdx[i:(j+1)] = [k] * len(range(i, (j+1)))#set sentence
29     index for each word in sentence to be the same
30 sentence[i:(j+1)] = [fullSentence] * len(range(i, (j+1)))
31
32 sentenceIdx = pd.DataFrame({'sentence_idx': sentenceIdx})
33 sentence = pd.DataFrame({'sentence': sentence})
34 df = df.join(sentenceIdx)
35 df = df.join(sentence)
36 df.insert(0, "word_idx", df.index.values)
37
38 df.to_csv('transcript_inc_sentence.csv', index=False)

```

create_embeddings.py

```

1 import h5py
2 import torch
3 import numpy as np
4 import pandas as pd
5
6 from himalaya.backend import set_backend, get_backend
7
8 from accelerate import Accelerator, find_executable_batch_size
9 from transformers import AutoModelForCausalLM, AutoTokenizer
10 from sentence_transformers import SentenceTransformer
11 import gensim.downloader
12 import re
13
14 if torch.cuda.is_available():
15     set_backend("torch_cuda")

```

```

16     print("Using cuda!")
17
18
19 #load transcript and extract sentences
20 # Extract sentences/sentence indices and add them to the transcript
21 # DataFrame.
22
23 transcript_path = "transcript_inc_sentence.csv" # Change as needed
24 df = pd.read_csv(transcript_path)
25
26 # Download 300-dimensional word2vec embeddings
27 model_name = 'word2vec-google-news-300'
28 n_features = 300
29
30
31 #creating static word embeddings
32 transcript_w2v = df.copy()
33
34 # Convert words to lowercase
35 transcript_w2v['word'] = transcript_w2v.word.str.lower()
36 # Filter words
37 transcript_w2v['filtered_word'] = transcript_w2v.word.apply(lambda
38     word: re.sub(r"[^\w\s]", "", word))
39
40 # Function to extract embeddings if available
41 def get_vector(word):
42     #word = ''.join(char for char in word if char.isalnum() and char
43     #not in [",", ",", ".", "?", "!", "-"])
44     if word in model.key_to_index:
45         return model.get_vector(word, norm=True).astype(np.float32)
46     return np.nan

```

```

45
46 # Extract embedding for each word
47 transcript_w2v['embedding'] = transcript_w2v.filtered_word.apply(
48     get_vector) # word2vec embeddings are 300-dimensional
49 # Insert zero vectors for entries with no embeddings
50 embedding_length = n_features # Length of the embeddings (300 for
51     word2vec)
52
53
54
55 # Print out words not found in vocabulary
56 print(f'{(transcript_w2v.embedding.isna()).sum()} words not found:')
57 print(transcript_w2v.filtered_word[transcript_w2v.embedding.isna()].value_counts())
58
59
60 #creating contextual word embeddings
61 #model to tokenize the data
62 modelname = "gpt2"
63 context_len = 32
64 device = torch.device("cpu")
65 if torch.cuda.is_available():
66     device = torch.device("cuda", 0)
67     print("Using cuda!")
68
69 model_sentence = SentenceTransformer("intfloat/e5-mistral-7b-
70     instruct")
71 # In case you want to reduce the maximum sequence length:
72 model_sentence.max_seq_length = 1000

```

```

72
73 # Load model for tokenization
74 tokenizer = AutoTokenizer.from_pretrained(modelname)
75
76 #explode dataframe using tokens
77 df["hftoken"] = df.word.apply(lambda x: tokenizer.tokenize(" " + x))
78
79 df = df.explode("hftoken", ignore_index=True)
80 df["token_id"] = df.hftoken.apply(tokenizer.convert_tokens_to_ids)
81
82 df.head(10)
83
84 #model to get token embeddings
85 print("Loading model...")
86 model = AutoModelForCausalLM.from_pretrained(modelname)
87
88 print(
89     f"Model : {modelname}"
90     f"\nLayers: {model.config.num_hidden_layers}"
91     f"\nEmbDim: {model.config.hidden_size}"
92     f"\nConfig: {model.config}"
93 )
94 model = model.eval()
95 model = model.to(device)
96
97 # setting up data for getting embeddings (using previous 32 tokens
98 # as context)
99 # could change this to use e.g. last sentence as context
100 token_ids = df.token_id.tolist()
101 print(f"Token id length: {len(token_ids)}")
102

```

```

103 fill_value = 0
104 if tokenizer.pad_token_id is not None:
105     fill_value = tokenizer.pad_token_id
106
107 #data ends up being an array of length token_ids containing arrays
108 # of length context_len + 1
109 data = torch.full((len(token_ids), context_len + 1), fill_value,
110                   dtype=torch.long)
111 for i in range(len(token_ids)):
112     example_tokens = token_ids[max(0, i - context_len) : i + 1]
113     data[i, -len(example_tokens) :] = torch.tensor(example_tokens)
114
115
116
117 #extract embeddings of all of the tokens
118 #use accekerator to make extracting features more efficient
119 accelerator = Accelerator()
120 @find_executable_batch_size(starting_batch_size=32)
121 def inference_loop(batch_size=32):
122     # nonlocal accelerator # Ensure they can be used in our context
123     accelerator.free_memory() # Free all lingering references
124
125     data_dl = torch.utils.data.DataLoader(
126         data, batch_size=batch_size, shuffle=False
127     )
128
129     top_guesses = []
130     ranks = []
131     true_probs = []

```

```

132     entropies = []
133     embeddings = []
134
135     #batch size is number of entries in each batch (i.e 32,32... (something less than 32 at end))
136
137     with torch.no_grad():
138         i=0
139         for batch in data_dl:
140             # Get output from model
141             output = model(batch.to(device), output_hidden_states=True)
142             logits = output.logits
143             states = output.hidden_states
144
145             true_ids = batch[:, -1]
146             brange = list(range(len(true_ids)))
147             logits_order = logits[:, -2, :].argsort(descending=True)
148             batch_top_guesses = logits_order[:, 0]
149             batch_ranks = torch.eq(logits_order, true_ids.reshape(-1, 1).to(device)).nonzero()[:, 1]
150             batch_probs = torch.softmax(logits[:, -2, :], dim=-1)
151             batch_true_probs = batch_probs[brange, true_ids]
152             batch_entropy = torch.distributions.Categorical(probs=batch_probs).entropy()
153
154             #hidden states is a list of tensors where each tensor has shape (batch_size, sequence length, hidden_size)
155             #hidden size is dimensionality of the model's hidden representations
156
157             #extracts the hidden state for the last token of each

```

```
sequence in the batch, giving a list of 13 tensors of shape (
batch size, hidden size) i.e (32, 768)

batch_embeddings = [state[:, -1, :].numpy(force=True)

for state in states]

top_guesses.append(batch_top_guesses.numpy(force=True))

ranks.append(batch_ranks.numpy(force=True))

true_probs.append(batch_true_probs.numpy(force=True))

entropies.append(batch_entropy.numpy(force=True))

embeddings.append(batch_embeddings)

i+=1

print(f"Number of batches is {i}")

return top_guesses, ranks, true_probs, entropies, embeddings

top_guesses, ranks, true_probs, entropies, embeddings =
inference_loop()

#embeddings is a list of 180 lists, each list has 13 tensors of
shape (32, 768)

#GPT 2 has 13 layers of embeddings, each word embedding is 768
dimensions long

print(f"There are {len(embeddings[0])} layers of embeddings")

print(f"Each word embedding is {embeddings[0][0].shape[1]}"
      dimensions long)

embeddingsA=np.hstack(embeddings)#this should be len 180, containing
      a bunch of ndarrays of shape (32,768)

print(embeddingsA.shape)

embeddingsB=embeddingsA[12,:,:,:]

print(embeddingsB.shape)# Shape: (num tokens, embedding size)
```

```

182
183 print("List shape")
184 print(len(embeddingsB.tolist()))
185 print(len(embeddingsB.tolist()[0]))
186
187 df["rank"] = np.concatenate(ranks)
188 df["true_prob"] = np.concatenate(true_probs)
189 df["top_pred"] = np.concatenate(top_guesses)
190 df["entropy"] = np.concatenate(entropies)
191 df["tkn_embedding"] = embeddingsB.tolist()
192 df.head(10)
193
194
195 #average embeddings over tokens in each word, and add to dataframe
196 aligned_embeddings = []
197 for _, group in df.groupby("word_idx"): # group by word index
198     indices = group.index.to_numpy()
199     average_emb = embeddingsB[indices].mean(0) # average features
200     aligned_embeddings.append(average_emb)
201 aligned_embeddings = np.stack(aligned_embeddings)
202 print(f"LLM embeddings matrix has shape: {aligned_embeddings.shape}")
203
204 print(len(aligned_embeddings.tolist()))
205
206 #get list of sentences in order
207 # Get the indices of the first occurrence of each unique sentence
208 unique_indices = df["sentence_idx"].drop_duplicates(keep="first").
    index.tolist()
209 # Access each sentence in order
210 sentences = df.loc[unique_indices, "sentence"].tolist()
211
```

```

212 #create new list of embeddings for each sentence
213 sentence_embeddings=[]
214 sentence_embeddings = model_sentence.encode(sentences)
215 dataLength=df.shape[0]
216
217 # explode the sentence list to line up with words(for each word add
# its sentence embedding to array)
218 unique_word_indices = df[ "word_idx" ].drop_duplicates(keep="first") .
    index.tolist()
219 expanded_sentence_embeddings=[]
220 for idx in unique_word_indices:
221     curSentence=df[ "sentence_idx" ].iloc[idx]
222     expanded_sentence_embeddings.append(sentence_embeddings [
        curSentence])
223
224 print(len(expanded_sentence_embeddings))
225
226
227 df_sentence_embeddings = pd.DataFrame({ 'sentence_embedding' :
    expanded_sentence_embeddings})
228 df_word_embeddings = pd.DataFrame({ 'word_embedding' :
    aligned_embeddings.tolist()})
229 df_static_word_embeddings =pd.DataFrame({ 'static_word_embedding' :
    transcript_w2v.embedding.tolist()})
230
231 df_word_embeddings = df_word_embeddings.join(df_sentence_embeddings)
232 df_word_embeddings = df_word_embeddings.join(
    df_static_word_embeddings)
233
234
235
236 print(df_word_embeddings.head())

```

```
237 df_word_embeddings.to_hdf('embeddings.h5', key='df', mode='w')
```

joint_encoding.py

```
1 import mne
2 import h5py
3 import torch
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 from nilearn.plotting import plot_markers
8
9 from himalaya.backend import set_backend, get_backend
10 from himalaya.ridge import RidgeCV
11 from himalaya.scoring import correlation_score
12 from himalaya.scoring import correlation_score_split
13
14 from sklearn.model_selection import KFold
15 from sklearn.pipeline import make_pipeline
16 from sklearn.preprocessing import StandardScaler
17
18 from himalaya.kernel_ridge import Kernelizer, ColumnKernelizer,
   MultipleKernelRidgeCV
19 import joblib
20
21 if torch.cuda.is_available():
22     set_backend("torch_cuda")
23     print("Using cuda!")
24
25 #span of the epochs
26 tMin=-6.0
27 tMax=6.0
28
29 #sampling rate
```

```

30 sampling_rate = 512
31 step_size=16
32
33 #load transcript and load embeddings
34 transcript_path = "transcript_inc_sentence.csv"
35 df = pd.read_csv(transcript_path)
36
37 df_embeddings= pd.read_hdf('embeddings.h5', 'df')
38 #print(df_embeddings.head())
39 aligned_embeddings = np.vstack(df_embeddings["word_embedding"].
40                                values)
40 sentence_embeddings = np.vstack(df_embeddings["sentence_embedding"].
41                                values)
41 static_word_embeddings = np.vstack(df_embeddings["
42                                static_word_embedding"].values)
42
43
44 print(aligned_embeddings.shape)
45 print(sentence_embeddings.shape)
46 print(static_word_embeddings.shape)
47
48 #construct separate dataframe containins words with their start and
49     end times
50
51 df_word = df.groupby("word_idx").agg(dict(word="first", start="first",
52                                             , end="last"))
51
52 #load the preprocessed brain data
53 file_path = "sub-03_task-podcast_desc-highgamma_ieeg.fif"
54 raw = mne.io.read_raw_fif(file_path, verbose=False)
55 picks = mne.pick_channels_regreg(raw.ch_names, "LG[AB]*")
55 raw = raw.pick(picks)
56

```

```

57 print(raw)

58

59 #map seconds to samples by multiplying to sampling rate
60 events = np.zeros((len(df_word), 3), dtype=int)
61 events[:, 0] = (df_word.start * raw.info['sfreq']).astype(int)
62 print(events.shape)

63

64 #create epochs and downsample to 32 Hz
65 epochs = mne.Epochs(raw, events, tmin=tMin, tmax=tMax, baseline=None,
66                      proj=False, event_id=None, preload=True, event_repeated="merge")
67 print(f"Epochs object has a shape of: {epochs._data.shape}")
68 epochs = epochs.resample(sfreq=32, npad='auto', method='fft', window
69                      ='hamming')
70
71 print(f"Epochs object has a shape of: {epochs._data.shape}")

72 #create Y matrices
73 #reshape target matrix Y by horizontally stacking electrodes and lags
74 #along the second dimension
75 epochs_data = epochs.get_data(copy=True)
76 epochs_data = epochs_data.reshape(len(epochs), -1)
77 print(f"ECOG data matrix shape: {epochs_data.shape}")

78 selected_df = df_word.iloc[epochs.selection]
79 averaged_embeddings = aligned_embeddings[epochs.selection]
80 averaged_sentence_embeddings=sentence_embeddings[epochs.selection]

81

82

83 #Change the float precision to float32 for all data to take
84 #advantage of the GPU memory and computational speed.

```

```

84 X_word_ctx = averaged_embeddings
85 X_word_static = static_word_embeddings[epochs.selection]
86 X_sentence=averaged_sentence_embeddings
87 X_word_used=X_word_ctx
88
89 X_joint = np.hstack([X_word_used, X_sentence]) # Horizontal-stack
          both embeddings to create joint model
90 Y = epochs_data
91
92 print(f"Joint predictor matrix shape: {X_joint.shape}")
93 print(f"X_word_ctx shape: {X_word_ctx.shape}")
94 print(f"X_word_static shape: {X_word_static.shape}")
95 print(f"X_sentence shape: {X_sentence.shape}")
96 print(f"Y shape: {Y.shape}")
97
98 #change float precision to float32 for all data to take advantage of
          GPU memory/computational speed
99 if "torch" in get_backend().__name__:
100     X_joint = X_joint.astype(np.float32)
101     X_word_ctx=X_word_ctx.astype(np.float32)
102     X_word_static=X_word_static.astype(np.float32)
103     Y = Y.astype(np.float32)
104
105 #do the banded ridge regression
106
107 # Cross validation is done with an outer part split into 2 and inner
          part split into 5
108 inner_cv = KFold(n_splits=5)
109
110 # Make pipeline with kernelizer for each feature space
111 column_pipeline = make_pipeline(
112     StandardScaler(with_mean=True, with_std=True),

```

```

113     Kernelizer(kernel="linear"),
114 )
115
116 #set up slices
117 width_w = X_word_used.shape[1]
118 width_s = X_sentence.shape[1]
119
120 slice_w = slice(0, width_w)
121 slice_s = slice(width_w, width_w + width_s)
122 print(f"Word slice: {slice_w}")
123 print(f"Sentence slice: {slice_s}")
124
125 # Compile joint column kernelizer
126 column_kernelizer = ColumnKernelizer(
127     [(
128         'word', column_pipeline, slice_w),
129         ('sentence', column_pipeline, slice_s)])
130
131 # Ridge regression with alpha grid and nested CV
132 solver = 'random_search'
133 n_iter = 20
134 alphas = np.logspace(1, 10, 10)
135 solver_params = dict(n_iter=n_iter, alphas=alphas)
136
137 # Banded ridge regression with column kernelizer
138 banded_ridge = MultipleKernelRidgeCV(kernels="precomputed", solver=
139             solver, solver_params=solver_params, cv=inner_cv)
140
141 # Chain transforms and estimator into pipeline
142 pipeline = make_pipeline(column_kernelizer, banded_ridge)
143
144 def train_joint_encoding(X, Y, _split=False, _epochs_shape=None):

```

```

144
145     corrs = [] # empty array to store correlation results
146
147     kfold = KFold(2, shuffle=False) # outer 2-fold cross-validation
148
149     setup
150
151     for train_index, test_index in kfold.split(X): # loop through
152         folds
153
154
155         # Split train and test datasets
156
157         X1_train, X1_test = X[train_index], X[test_index]
158
159         Y_train, Y_test = Y[train_index], Y[test_index]
160
161
162         # Standardize Y
163
164         scaler = StandardScaler()
165
166         Y_train = scaler.fit_transform(Y_train)
167
168         Y_test = scaler.transform(Y_test)
169
170
171         pipeline.fit(X1_train, Y_train) # Fit pipeline with
172         transforms and ridge estimator
173
174
175         if _split:
176
177             Y_preds = pipeline.predict(X1_test, split=True) # Use
178             trained model to predict on test set
179
180             corr = correlation_score_split(Y_test, Y_preds) #
181             Compute correlation score
182
183         else:
184
185             Y_preds = pipeline.predict(X1_test)
186
187             corr = correlation_score(Y_test, Y_preds).reshape(
188                 epochs_shape) # Compute correlation score
189
190
191             if "torch" in get_backend().__name__: # if using gpu,
192                 transform tensor back to numpy
193
194             corr = corr.numpy(force=True)

```

```

169
170     corrs.append(corr) # append fold correlation results to
171     final_results
172
173
174 alphas = np.logspace(1, 10, 10) # specify alpha values
175 inner_cv_single = KFold(n_splits=5, shuffle=False) # inner 5-fold
176     cross-validation setup
177 single_model = make_pipeline(
178     StandardScaler(), RidgeCV(alphas, fit_intercept=True, cv=
179     inner_cv_single) # pipeline
180 )
181
182
183 def train_single_encoding(X, Y,_epochs_shape):
184
185     corrs = [] # empty array to store correlation results
186     kfold = KFold(2, shuffle=False) # outer 2-fold cross-validation
187     setup
188
189     for train_index, test_index in kfold.split(X): # loop through
190         folds
191
192         # Split train and test datasets
193         X1_train, X1_test = X[train_index], X[test_index]
194         Y_train, Y_test = Y[train_index], Y[test_index]
195
196         # Standardize Y
197
198         scaler = StandardScaler()
199
200         Y_train = scaler.fit_transform(Y_train)
201
202         Y_test = scaler.transform(Y_test)

```

```

196     single_model.fit(X1_train, Y_train) # Fit pipeline with
197     transforms and ridge estimator
198
199     Y_preds = single_model.predict(X1_test) # Use trained model
200     to predict on test set
201
202     corr = correlation_score(Y_test, Y_preds).reshape(
203         _epochs_shape) # Compute correlation score
204
205
206
207     if "torch" in get_backend().__name__: # if using gpu,
208     transform tensor back to numpy
209     corr = corr.numpy(force=True)
210
211
212     corrs.append(corr) # append fold correlation results to
213     final results
214
215     return np.stack(corrs)
216
217
218
219 epochs_shape = epochs._data.shape[1:] # number of electrodes *
220     number of lags
221
222 corrs_embedding = train_joint_encoding(X_joint, Y, _split=False,
223     _epochs_shape=epochs_shape)
224
225 print(f"Encoding performance correlating matrix shape: {
226     corrs_embedding.shape}")
227
228 # Encoding performance correlating matrix shape: (2, 127, 128)
229
230
231 corrs_embedding_split = train_joint_encoding(X_joint, Y, True)
232 corrs_embedding_split = corrs_embedding_split.reshape(2, 2, *
233     epochs_shape)
234
235 print(f"Encoding performance correlating matrix shape: {
236     corrs_embedding_split.shape}")
237
238 # Encoding performance correlating matrix shape: (2, 2, 127, 128)
239
240 corrs_embedding_word_ctx = train_single_encoding(X_word_ctx, Y,
241     _epochs_shape=epochs_shape) # predictions with just contextual

```

```

    embeddings

217 print(f"Encoding performance correlating matrix shape: {
        corrs_embedding_word_ctx.shape}")

218 # Encoding performance correlating matrix shape: (2, 127, 128)

219 corrs_embedding_word_static = train_single_encoding(X_word_static, Y
            , _epochs_shape=epochs_shape) # predictions with just contextual
            embeddings

220 print(f"Encoding performance correlating matrix shape: {
        corrs_embedding_word_ctx.shape}")

221

222 #127 electrodes, 128 lags

223

224 joblib.dump(corrs_embedding, 'corrs_embedding_joint_ctx_sentence.pkl'
            )

225 joblib.dump(corrs_embedding_split, '
            corrs_embedding_split_ctx_sentence.pkl')

226 joblib.dump(corrs_embedding_word_ctx, 'corrs_embedding_word_ctx.pkl'
            )

227 joblib.dump(corrs_embedding_word_static, '
            corrs_embedding_word_static.pkl')

228

229 #Note:
230 #      - The dimensions of the following arrays are:
231 #          - corrs_embedding: (2, number_of_electrodes, number_of_lags
232 #          )
233 #          - corrs_embedding_split: (2, 2, number_of_electrodes,
234 #          number_of_lags)
235 #          - corrs_embedding_word_ctx: (2, number_of_electrodes,
236 #          number_of_lags)

237

238

239 print("Done!")

```

Appendix B

Engineering Standards

This independent project adheres to the following engineering and industrial standards:

- **Programming Languages:** Python,
- **Open Source Software:** NumPy, PyTorch, HuggingFace Transformers Library

Standards Compliance

All tools, programming languages, and methodologies employed in the development of this thesis conform to well-established academic and industrial standards. This adherence ensures that both the theoretical insights and the associated codebases are robust, scalable, and easily adaptable to a variety of practical applications in the engineering domain. Additionally, the thesis consistently utilizes the International System of Units (SI units), ensuring standardization and precision in the presentation of technical data.

Bibliography

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention is all you need,” June 2017.
- [2] F. Xu, Q. Hao, Z. Zong, J. Wang, Y. Zhang, J. Wang, X. Lan, J. Gong, T. Ouyang, F. Meng, C. Shao, Y. Yan, Q. Yang, Y. Song, S. Ren, X. Hu, Y. Li, J. Feng, C. Gao, and Y. Li, “Towards large reasoning models: A survey of reinforced reasoning with large language models,” 2025.
- [3] A. Goldstein, Z. Zada, E. Buchnik, M. Schain, A. Price, B. Aubrey, S. A. Nas-tase, A. Feder, D. Emanuel, A. Cohen, A. Jansen, H. Gazula, G. Choe, A. Rao, C. Kim, C. Casto, L. Fanda, W. Doyle, D. Friedman, P. Dugan, L. Melloni, R. Reichart, S. Devore, A. Flinker, L. Hasenfratz, O. Levy, A. Hassidim, M. Brenner, Y. Matias, K. A. Norman, O. Devinsky, and U. Hasson, “Shared computational principles for language processing in humans and deep language models,” *Nat. Neurosci.*, vol. 25, pp. 369–380, Mar. 2022.
- [4] Z. Wu, L. Qiu, A. Ross, E. Akyürek, B. Chen, B. Wang, N. Kim, J. Andreas, and Y. Kim, “Reasoning or reciting? exploring the capabilities and limitations of language models through counterfactual tasks,” 2023.
- [5] S. Bubeck, V. Chandrasekaran, R. Eldan, J. Gehrke, E. Horvitz, E. Kamar, P. Lee, Y. T. Lee, Y. Li, S. Lundberg, H. Nori, H. Palangi, M. T. Ribeiro,

- and Y. Zhang, “Sparks of artificial general intelligence: Early experiments with GPT-4,” 2023.
- [6] J. Gauthier, J. Hu, E. Wilcox, P. Qian, and R. Levy, “SyntaxGym: An online platform for targeted evaluation of language models,” in *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, (Stroudsburg, PA, USA), Association for Computational Linguistics, 2020.
 - [7] A. Warstadt and S. R. Bowman, “What artificial neural networks can tell us about human language acquisition,” 2022.
 - [8] K. Mahowald, A. A. Ivanova, I. A. Blank, N. Kanwisher, J. B. Tenenbaum, and E. Fedorenko, “Dissociating language and thought in large language models,” 2023.
 - [9] A. M. Paunov, I. A. Blank, O. Jouravlev, Z. Mineroff, J. Gallée, and E. Fedorenko, “Differential tracking of linguistic vs. mental state content in naturalistic stimuli by language and theory of mind (ToM) brain networks,” *Neurobiol. Lang. (Camb.)*, vol. 3, pp. 413–440, July 2022.
 - [10] E. Fedorenko, A. A. Ivanova, and T. I. Regev, “The language network as a natural kind within the broader landscape of the human brain,” *Nat. Rev. Neurosci.*, vol. 25, pp. 289–312, May 2024.
 - [11] E. Fedorenko and R. Varley, “Language and thought are not the same thing: evidence from neuroimaging and neurological patients,” *Ann. N. Y. Acad. Sci.*, vol. 1369, pp. 132–153, Apr. 2016.
 - [12] G. Granato, A. M. Borghi, and G. Baldassarre, “A computational model of language functions in flexible goal-directed behaviour,” *Sci. Rep.*, vol. 10, p. 21623, Dec. 2020.

- [13] T. Ruffman, L. Slade, and E. Crowe, “The relation between children’s and mothers’ mental state language and theory-of-mind understanding,” *Child Dev.*, vol. 73, pp. 734–751, May 2002.
- [14] O. et al., “OpenAI o1 system card,” Dec. 2024.
- [15] D. et al., “DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning,” 2025.
- [16] Q. Team, “Qwq: Reflect deeply on the boundaries of the unknown,” Nov 2024.
- [17] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. Chi, Q. Le, and D. Zhou, “Chain-of-thought prompting elicits reasoning in large language models,” 2022.
- [18] T. Kojima, S. S. Gu, M. Reid, Y. Matsuo, and Y. Iwasawa, “Large language models are Zero-Shot reasoners,” 2022.
- [19] Y. Du, S. Li, A. Torralba, J. B. Tenenbaum, and I. Mordatch, “Improving factuality and reasoning in language models through multiagent debate,” 2023.
- [20] S. Chen, B. Li, and D. Niu, “Boosting of thoughts: Trial-and-error problem solving with large language models,” 2024.
- [21] N. Shinn, F. Cassano, E. Berman, A. Gopinath, K. Narasimhan, and S. Yao, “Reflexion: Language agents with verbal reinforcement learning,” 2023.
- [22] L. Yang, Z. Yu, T. Zhang, S. Cao, M. Xu, W. Zhang, J. E. Gonzalez, and B. Cui, “Buffer of thoughts: Thought-augmented reasoning with large language models,” 2024.
- [23] S. Parashar, B. Olson, S. Khurana, E. Li, H. Ling, J. Caverlee, and S. Ji, “Inference-time computations for LLM reasoning and planning: A benchmark and insights,” 2025.

- [24] S. Yao, D. Yu, J. Zhao, I. Shafran, T. L. Griffiths, Y. Cao, and K. Narasimhan, “Tree of thoughts: Deliberate problem solving with large language models,” 2023.
- [25] L. Chen, J. Q. Davis, B. Hanin, P. Bailis, I. Stoica, M. Zaharia, and J. Zou, “Are more LLM calls all you need? towards scaling laws of compound inference systems,” 2024.
- [26] M. K. Ho, R. Saxe, and F. Cushman, “Planning with theory of mind,” *Trends Cogn. Sci.*, vol. 26, pp. 959–971, Nov. 2022.
- [27] T. Ullman, “Large language models fail on trivial alterations to Theory-of-Mind tasks,” 2023.
- [28] N. Shapira, M. Levy, S. H. Alavi, X. Zhou, Y. Choi, Y. Goldberg, M. Sap, and V. Shwartz, “Clever hans or neural theory of mind? stress testing social reasoning in large language models,” in *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)* (Y. Graham and M. Purver, eds.), (St. Julian’s, Malta), pp. 2257–2273, Association for Computational Linguistics, Mar. 2024.
- [29] H. Kim, M. Sclar, X. Zhou, R. Bras, G. Kim, Y. Choi, and M. Sap, “FAN-ToM: A benchmark for stress-testing machine theory of mind in interactions,” in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing* (H. Bouamor, J. Pino, and K. Bali, eds.), (Singapore), pp. 14397–14413, Association for Computational Linguistics, Dec. 2023.
- [30] K. Gandhi, J.-P. Fränken, T. Gerstenberg, and N. Goodman, “Understanding social reasoning in language models with language models,” in *Thirty-seventh Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2023.

- [31] C. Jin, Y. Wu, J. Cao, J. Xiang, Y.-L. Kuo, Z. Hu, T. Ullman, A. Torralba, J. Tenenbaum, and T. Shu, “MMToM-QA: Multimodal theory of mind question answering,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (L.-W. Ku, A. Martins, and V. Srikumar, eds.), (Bangkok, Thailand), pp. 16077–16102, Association for Computational Linguistics, Aug. 2024.
- [32] Y. Wu, Y. He, Y. Jia, R. Mihalcea, Y. Chen, and N. Deng, “Hi-ToM: A benchmark for evaluating higher-order theory of mind reasoning in large language models,” in *Findings of the Association for Computational Linguistics: EMNLP 2023* (H. Bouamor, J. Pino, and K. Bali, eds.), (Singapore), pp. 10691–10706, Association for Computational Linguistics, Dec. 2023.
- [33] M. Sclar, S. Kumar, P. West, A. Suhr, Y. Choi, and Y. Tsvetkov, “Minding language models’ (lack of) theory of mind: A plug-and-play multi-character belief tracker,” in *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (A. Rogers, J. Boyd-Graber, and N. Okazaki, eds.), (Toronto, Canada), pp. 13960–13980, Association for Computational Linguistics, July 2023.
- [34] Y. Gu, O. Tafjord, H. Kim, J. Moore, R. L. Bras, P. Clark, and Y. Choi, “SimpleToM: Exposing the gap between explicit ToM inference and implicit ToM application in LLMs,” 2024.
- [35] C. Jung, D. Kim, J. Jin, J. Kim, Y. Seonwoo, Y. Choi, A. Oh, and H. Kim, “Perceptions to beliefs: Exploring precursory inferences for theory of mind in large language models,” 2024.
- [36] J. Wu, Z. Chen, J. Deng, S. Sabour, H. Meng, and M. Huang, “COKE: A cognitive knowledge graph for machine theory of mind,” May 2023.

- [37] H. Kim, M. Sclar, T. Zhi-Xuan, L. Ying, S. Levine, Y. Liu, J. B. Tenenbaum, and Y. Choi, “Hypothesis-driven theory-of-mind reasoning for large language models,” Feb. 2025.
- [38] A. Afrasiyabi, E. Busch, R. Singh, D. Bhaskar, L. Caplette, N. Turk-Browne, and S. Krishnaswamy, “Looking through the mind’s eye via multimodal encoder-decoder networks,” 2024.
- [39] G. Tuckute, N. Kanwisher, and E. Fedorenko, “Language in brains, minds, and machines,” *Annu. Rev. Neurosci.*, vol. 47, pp. 277–301, Aug. 2024.
- [40] C. Caucheteux and J.-R. King, “Brains and algorithms partially converge in natural language processing,” *Commun. Biol.*, vol. 5, p. 134, Feb. 2022.
- [41] A. Goldstein, Z. Zada, E. Buchnik, M. Schain, A. Price, B. Aubrey, S. A. Nastase, A. Feder, D. Emanuel, A. Cohen, A. Jansen, H. Gazula, G. Choe, A. Rao, C. Kim, C. Casto, L. Fanda, W. Doyle, D. Friedman, P. Dugan, L. Melloni, R. Reichart, S. Devore, A. Flinker, L. Hasenfratz, O. Levy, A. Hassidim, M. Brenner, Y. Matias, K. A. Norman, O. Devinsky, and U. Hasson, “Shared computational principles for language processing in humans and deep language models,” *Nat. Neurosci.*, vol. 25, pp. 369–380, Mar. 2022.
- [42] A. Goldstein, E. Ham, S. A. Nastase, Z. Zada, A. Dabush, B. Bobbi Aubrey, M. Schain, H. Gazula, A. Feder, W. Doyle, S. Devore, P. Dugan, D. Friedman, M. Brenner, A. Hassidim, O. Devinsky, A. Flinker, O. Levy, and U. Hasson, “Correspondence between the layered structure of deep language models and temporal structure of natural language processing in the human brain.” July 2022.

- [43] G. Tuckute, A. Sathe, S. Srikant, M. Taliaferro, M. Wang, M. Schrimpf, K. Kay, and E. Fedorenko, “Driving and suppressing the human language network using large language models,” *Nat. Hum. Behav.*, vol. 8, pp. 544–561, Mar. 2024.
- [44] F. Pereira, B. Lou, B. Pritchett, S. Ritter, S. J. Gershman, N. Kanwisher, M. Botvinick, and E. Fedorenko, “Toward a universal decoder of linguistic meaning from brain activation,” *Nat. Commun.*, vol. 9, Mar. 2018.
- [45] C. Kauf, G. Tuckute, R. Levy, J. Andreas, and E. Fedorenko, “Lexical-semantic content, not syntactic structure, is the main contributor to ANN-brain similarity of fMRI responses in the language network,” *Neurobiol. Lang. (Camb.)*, vol. 5, pp. 7–42, Apr. 2024.
- [46] D. Jurafsky and J. H. Martin, *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition with Language Models*. 2024. Online manuscript released August 20, 2024.
- [47] E. Smith, O. Hsu, R. Qian, S. Roller, Y.-L. Boureau, and J. Weston, “Human evaluation of conversations is an open problem: comparing the sensitivity of various methods for evaluating dialogue agents,” in *Proceedings of the 4th Workshop on NLP for Conversational AI*, (Stroudsburg, PA, USA), Association for Computational Linguistics, 2022.
- [48] C.-W. Liu, R. Lowe, I. V. Serban, M. Noseworthy, L. Charlin, and J. Pineau, “How NOT to evaluate your dialogue system: An empirical study of unsupervised evaluation metrics for dialogue response generation,” 2016.
- [49] A. B. Sai, A. K. Mohankumar, and M. M. Khapra, “A survey of evaluation metrics used for NLG systems,” 2020.
- [50] S. Mehri, J. Choi, L. F. D’Haro, J. Deriu, M. Eskenazi, M. Gasic, K. Georgila, D. Hakkani-Tur, Z. Li, V. Rieser, S. Shaikh, D. Traum, Y.-T. Yeh, Z. Yu,

Y. Zhang, and C. Zhang, “Report from the NSF future directions workshop on automatic evaluation of dialog: Research directions and challenges,” 2022.

- [51] Y.-T. Yeh, M. Eskenazi, and S. Mehri, “A comprehensive assessment of dialog evaluation metrics,” 2021.
- [52] J. I. Choi, M. Collins, E. Agichtein, O. Rokhlenko, and S. Malmasi, “Combining multiple metrics for evaluating retrieval-augmented conversations,” in *Proceedings of the Third Workshop on Bridging Human–Computer Interaction and Natural Language Processing*, (Stroudsburg, PA, USA), Association for Computational Linguistics, 2024.
- [53] D. Khashabi, G. Stanovsky, J. Bragg, N. Lourie, J. Kasai, Y. Choi, N. A. Smith, and D. S. Weld, “GENIE: Toward reproducible and standardized human evaluation for text generation,” 2021.
- [54] T. Hashimoto, H. Zhang, and P. Liang, “Unifying human and statistical evaluation for natural language generation,” in *Proceedings of the 2019 Conference of the North*, (Stroudsburg, PA, USA), Association for Computational Linguistics, 2019.
- [55] P. Gupta, S. Mehri, T. Zhao, A. Pavel, M. Eskenazi, and J. Bigham, “Investigating evaluation of open-domain dialogue systems with human generated multiple references,” in *Proceedings of the 20th Annual SIGdial Meeting on Discourse and Dialogue*, (Stroudsburg, PA, USA), Association for Computational Linguistics, 2019.
- [56] R. Lowe, M. Noseworthy, I. V. Serban, N. Angelard-Gontier, Y. Bengio, and J. Pineau, “Towards an automatic Turing test: Learning to evaluate dialogue responses,” in *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (R. Barzilay and M.-Y.

Kan, eds.), (Vancouver, Canada), pp. 1116–1126, Association for Computational Linguistics, July 2017.

- [57] A. B. Sai, M. D. Gupta, M. M. Khapra, and M. Srinivasan, “Re-evaluating ADEM: A deeper look at scoring dialogue responses,” *Proc. Conf. AAAI Artif. Intell.*, vol. 33, pp. 6220–6227, July 2019.
- [58] C. Callison-Burch, M. Osborne, and P. Koehn, “Re-evaluating the role of Bleu in machine translation research,” in *11th Conference of the European Chapter of the Association for Computational Linguistics* (D. McCarthy and S. Wintner, eds.), (Trento, Italy), pp. 249–256, Association for Computational Linguistics, Apr. 2006.
- [59] A. Smith, C. Hardmeier, and J. Tiedemann, “Climbing mont BLEU: The strange world of reachable high-BLEU translations,” in *Proceedings of the 19th Annual Conference of the European Association for Machine Translation*, pp. 269–281, 2016.
- [60] E. Sulem, O. Abend, and A. Rappoport, “BLEU is not suitable for the evaluation of text simplification,” in *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing* (E. Riloff, D. Chiang, J. Hockenmaier, and J. Tsujii, eds.), (Brussels, Belgium), pp. 738–744, Association for Computational Linguistics, Oct.-Nov. 2018.
- [61] J. Deriu, A. Rodrigo, A. Otegi, G. Echegoyen, S. Rosset, E. Agirre, and M. Cieliebak, “Survey on evaluation methods for dialogue systems,” *arXiv [cs.CL]*, 2019.
- [62] F. Jelinek, R. L. Mercer, L. R. Bahl, and J. K. Baker, “Perplexity—a measure of the difficulty of speech recognition tasks,” *J. Acoust. Soc. Am.*, vol. 62, pp. S63–S63, Dec. 1977.

- [63] S. F. Chen, D. Beeferman, and R. Rosenfeld, “Evaluation metrics for language models,” 1998.
- [64] X. Gao, Y. Zhang, M. Galley, C. Brockett, and B. Dolan, “Dialogue response ranking training with large-scale human feedback data,” in *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)* (B. Webber, T. Cohn, Y. He, and Y. Liu, eds.), (Online), pp. 386–395, Association for Computational Linguistics, Nov. 2020.
- [65] A. K. Vijayakumar, M. Cogswell, R. R. Selvaraju, Q. Sun, S. Lee, D. Crandall, and D. Batra, “Diverse beam search: Decoding diverse solutions from neural sequence models,” 2016.
- [66] J. Li, M. Galley, C. Brockett, J. Gao, and B. Dolan, “A diversity-promoting objective function for neural conversation models,” in *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies* (K. Knight, A. Nenkova, and O. Rambow, eds.), (San Diego, California), pp. 110–119, Association for Computational Linguistics, June 2016.
- [67] Y. Zhang, M. Galley, J. Gao, Z. Gan, X. Li, C. Brockett, and B. Dolan, “Generating informative and diverse conversational responses via adversarial information maximization,” in *Advances in Neural Information Processing Systems* (S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, eds.), vol. 31, Curran Associates, Inc., 2018.
- [68] W. Xu, S. Pakhomov, P. Heagerty, E. Horvitz, E. R. Bradley, A. Campbell, J. D. Woolley, A. Cohen, D. Ben-Zeev, and T. Cohen, “Perplexity and proximity: Large language model perplexity complements semantic distance metrics for the detection of incoherent speech.” 2025.

- [69] T. Zhang, V. Kishore, F. Wu, K. Q. Weinberger, and Y. Artzi, “BERTScore: Evaluating text generation with BERT,” 2019.
- [70] G. Datta, N. Joshi, and K. Gupta, “Analysis of automatic evaluation metric on low-resourced language: Bertscore vs bleu score,” in *Speech and Computer* (S. R. M. Prasanna, A. Karpov, K. Samudravijaya, and S. S. Agrawal, eds.), (Cham), pp. 155–162, Springer International Publishing, 2022.
- [71] S. Xu, W. Xie, L. Zhao, and P. He, “Chain of draft: Thinking faster by writing less,” 2025.
- [72] C. Beaudoin, É. Leblanc, C. Gagner, and M. H. Beauchamp, “Systematic review and inventory of theory of mind measures for young children,” *Front. Psychol.*, vol. 10, p. 2905, 2019.
- [73] A. Grattafiori, A. Dubey, A. Jauhri, A. Pandey, A. Kadian, A. Al-Dahle, A. Letman, A. Mathur, A. Schelten, A. Vaughan, A. Yang, A. Fan, A. Goyal, A. Hartshorn, A. Yang, A. Mitra, A. Sravankumar, A. Korenev, A. Hinsvark, A. Rao, A. Zhang, A. Rodriguez, A. Gregerson, A. Spataru, B. Roziere, B. Biron, B. Tang, B. Chern, C. Caucheteux, C. Nayak, C. Bi, C. Marra, C. McConnell, C. Keller, C. Touret, C. Wu, C. Wong, C. C. Ferrer, C. Nikolaidis, D. Allonsius, D. Song, D. Pintz, D. Livshits, D. Wyatt, D. Esiobu, D. Choudhary, D. Mahajan, D. Garcia-Olano, D. Perino, D. Hupkes, E. Lakomkin, E. AlBadawy, E. Lobanova, E. Dinan, E. M. Smith, F. Radenovic, F. Guzmán, F. Zhang, G. Synnaeve, G. Lee, G. L. Anderson, G. Thattai, G. Nail, G. Mialon, G. Pang, G. Cucurell, H. Nguyen, H. Korevaar, H. Xu, H. Touvron, I. Zarov, I. A. Ibarra, I. Kloumann, I. Misra, I. Evtimov, J. Zhang, J. Copet, J. Lee, J. Geffert, J. Vranes, J. Park, J. Mahadeokar, and J. a. Shah, “The llama 3 herd of models,”

- [74] R. Stureborg, D. Alikaniotis, and Y. Suhara, “Large language models are inconsistent and biased evaluators,” May 2024.
- [75] J. E. Eicher and R. F. Irgolić, “Reducing selection bias in large language models,” 2024.
- [76] M. Sap, H. Rashkin, D. Chen, R. LeBras, and Y. Choi, “SocialIQA: Commonsense reasoning about social interactions,” Apr. 2019.
- [77] T. Lanham, A. Chen, A. Radhakrishnan, B. Steiner, C. Denison, D. Hernandez, D. Li, E. Durmus, E. Hubinger, J. Kernion, K. Lukošiūtė, K. Nguyen, N. Cheng, N. Joseph, N. Schiefer, O. Rausch, R. Larson, S. McCandlish, S. Kundu, S. Kadavath, S. Yang, T. Henighan, T. Maxwell, T. Telleen-Lawton, T. Hume, Z. Hatfield-Dodds, J. Kaplan, J. Brauner, S. R. Bowman, and E. Perez, “Measuring faithfulness in Chain-of-Thought reasoning,” 2023.
- [78] M. Turpin, J. Michael, E. Perez, and S. R. Bowman, “Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting,” 2023.
- [79] S. H. Tanneru, D. Ley, C. Agarwal, and H. Lakkaraju, “On the hardness of faithful Chain-of-Thought reasoning in large language models,” 2024.
- [80] C. Agarwal, S. H. Tanneru, and H. Lakkaraju, “Faithfulness vs. plausibility: On the (un)reliability of explanations from large language models,” 2024.
- [81] A. Welivita, Y. Xie, and P. Pu, “A large-scale dataset for empathetic response generation,” in *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing* (M.-F. Moens, X. Huang, L. Specia, and S. W.-t. Yih, eds.), (Online and Punta Cana, Dominican Republic), pp. 1251–1264, Association for Computational Linguistics, Nov. 2021.

- [82] M. J. Traxler and M. A. Gernsbacher, eds., *Handbook of Psycholinguistics*. San Diego, CA: Academic Press, 2 ed., Nov. 2006.
- [83] M. K. Tanenhaus, M. J. Spivey-Knowlton, K. M. Eberhard, and J. C. Sedivy, “Integration of visual and linguistic information in spoken language comprehension,” *Science*, vol. 268, pp. 1632–1634, June 1995.
- [84] J. E. Hanna, M. K. Tanenhaus, and J. C. Trueswell, “The effects of common ground and perspective on domains of referential interpretation,” *J. Mem. Lang.*, vol. 49, pp. 43–61, July 2003.
- [85] C. G. Chambers, M. K. Tanenhaus, and J. S. Magnuson, “Actions and affordances in syntactic ambiguity resolution,” *J. Exp. Psychol. Learn. Mem. Cogn.*, vol. 30, pp. 687–696, May 2004.
- [86] S. E. Brennan, “Invited talk: Processes that shape conversation and their implications for computational linguistics,” in *Proceedings of the 38th Annual Meeting of the Association for Computational Linguistics*, (Hong Kong), pp. 1–11, Association for Computational Linguistics, Oct. 2000.
- [87] G. Aist, J. Allen, E. Campana, L. Galescu, C. A. G. Gallo, S. C. Stoness, M. Swift, and M. Tanenhaus, “Software architectures for incremental understanding of human speech,” in *Interspeech 2006*, (ISCA), pp. aper 1869–Wed2FoP.5–0, ISCA, Sept. 2006.
- [88] Z. Zada, S. A. Nastase, B. Aubrey, I. Jalon, S. Michelmann, H. Wang, L. Hasenfratz, W. Doyle, D. Friedman, P. Dugan, L. Melloni, S. Devore, A. Flinker, O. Devinsky, A. Goldstein, and U. Hasson, “The “podcast” ECoG dataset for modeling neural activity during natural language comprehension.” Feb. 2025.
- [89] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos,

- D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay, “Scikit-learn: Machine learning in python,” *J. Mach. Learn. Res.*, vol. 12, p. 2825–2830, Nov. 2011.
- [90] T. Dupré la Tour, M. Visconti di Oleggio Castello, and J. L. Gallant, “The voxelwise modeling framework: a tutorial introduction to fitting encoding models to fMRI data.” Apr. 2024.
- [91] T. Dupré la Tour, M. Eickenberg, A. O. Nunez-Elizalde, and J. L. Gallant, “Feature-space selection with banded ridge regression,” *NeuroImage*, vol. 264, p. 119728, 2022.
- [92] D. Barrett, *Supernormal Stimuli: How Primal Urges Overran Their Evolutionary Purpose*. New York: W. W. Norton & Company, 2010.
- [93] C. Fernyhough and A. M. Borghi, “Inner speech as language process and cognitive tool,” *Trends Cogn. Sci.*, vol. 27, pp. 1180–1193, Dec. 2023.